

# **CS 152 Computer Architecture and Engineering**

# **CS252A Graduate Computer Architecture**

## **Lecture 3 - Pipelining**

Krste Asanovic  
Electrical Engineering and Computer Sciences  
University of California at Berkeley

**<http://www.eecs.berkeley.edu/~krste>**  
**<http://inst.eecs.berkeley.edu/~cs152>**

## Q&A in lecture

- Please post on Piazza thread so we can capture Q&A for others, and stay muted during lecture
- I'll answer some questions live in break and at end of lecture with recording paused to avoid recording your voice.

## Last Time in Lecture 2

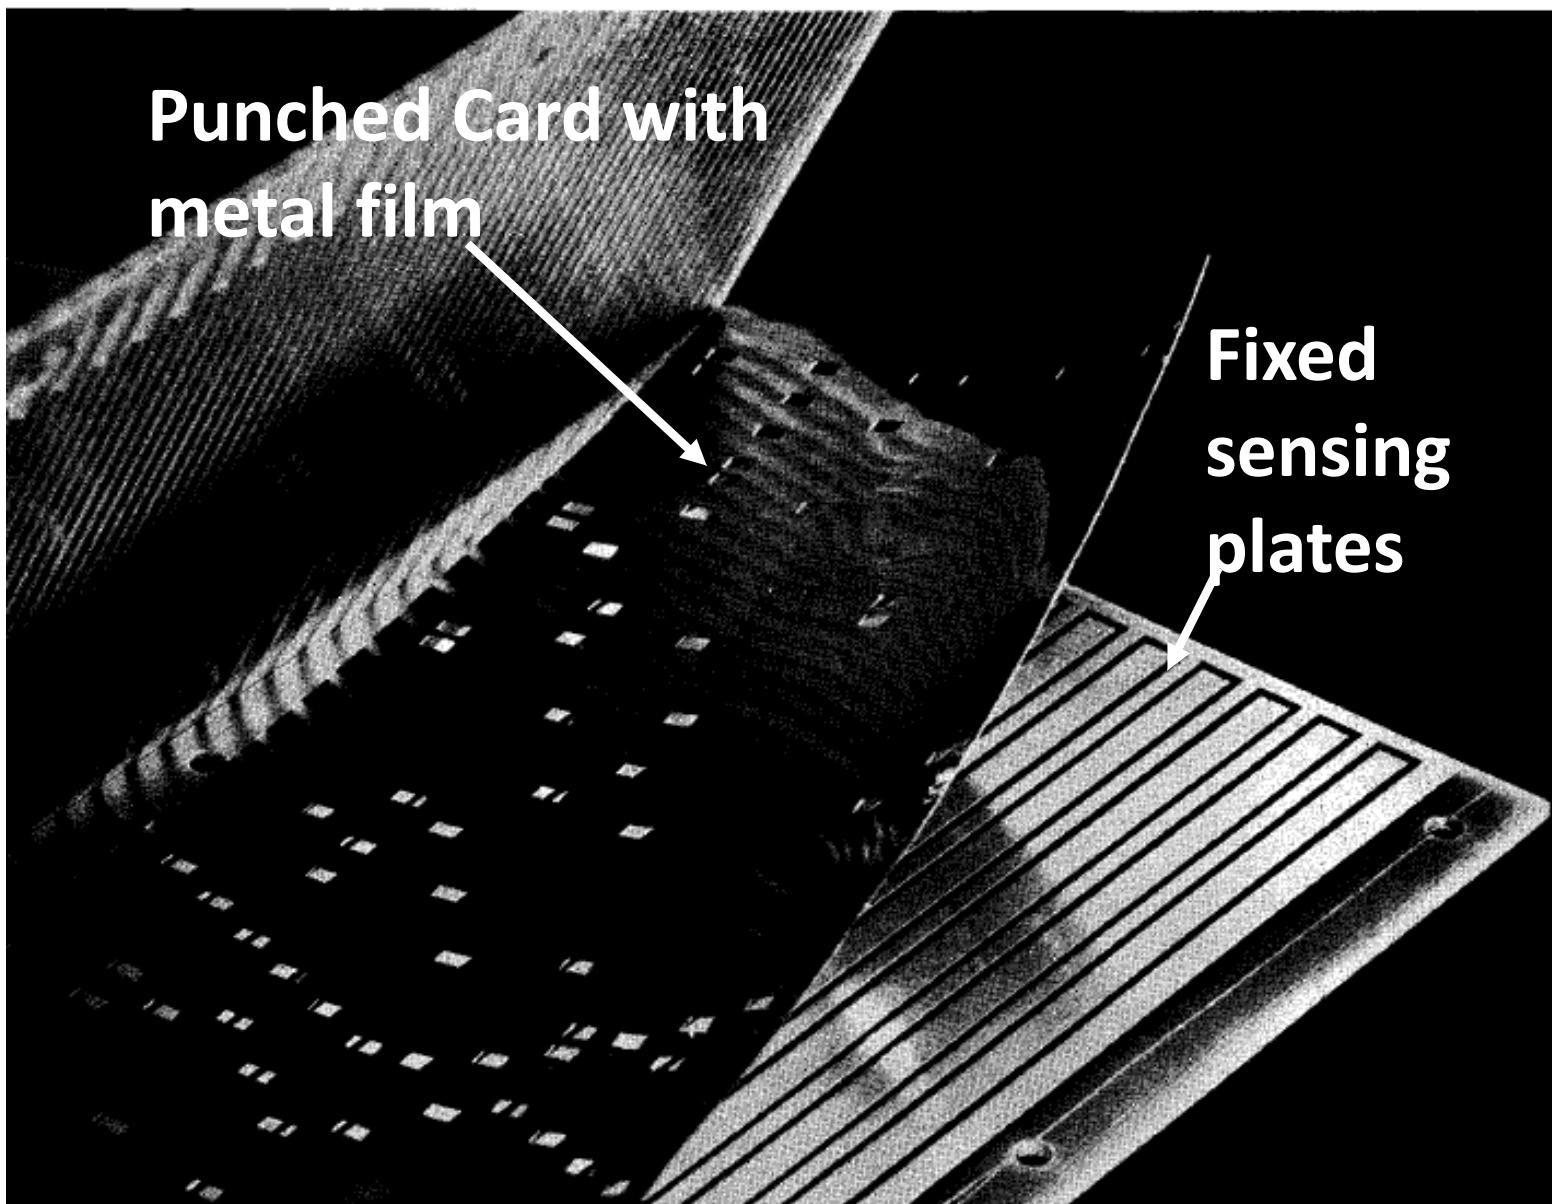
- Microcoding, an effective technique to manage control unit complexity, invented in era when logic (tubes), main memory (magnetic core), and ROM (diodes) used different technologies
- Difference between ROM and RAM speed motivated additional complex instructions
- Technology advances leading to fast SRAM made technology assumptions invalid
- Complex instructions sets impede parallel and pipelined implementations

# Recap: Microprogramming in IBM 360

	M30	M40	M50	M65
Datapath width (bits)	8	16	32	64
$\mu$ inst width (bits)	50	52	85	87
$\mu$ code size (K $\mu$ insts)	4	4	2.75	2.75
$\mu$ store technology	CCROS	TCROS	BCROS	BCROS
$\mu$ store cycle (ns)	750	625	500	200
memory cycle (ns)	1500	2500	2000	750
Rental fee (\$K/month)	4	7	15	35

- Only the fastest models (75 and 95) were hardwired

# IBM Card-Capacitor Read-Only Storage



[ IBM Journal, January 1961]

# Microcode Emulation

- IBM initially miscalculated the importance of software compatibility with earlier models when introducing the 360 series
- Honeywell stole some IBM 1401 customers by offering translation software (“Liberator”) for Honeywell H200 series machine
- IBM retaliated with optional additional microcode for 360 series that could emulate IBM 1401 ISA, later extended for IBM 7000 series
  - one popular program on 1401 was a 650 simulator, so some customers ran many 650 programs on emulated 1401s
  - i.e., 650 simulated on 1401 emulated on 360

# Microprogramming thrived in '60s and '70s

- Significantly faster ROMs than DRAMs were available
- For complex instruction sets, datapath and controller were cheaper and simpler
- New instructions , e.g., floating point, could be supported without datapath modifications
- Fixing bugs in the controller was easier
- ISA compatibility across various models could be achieved easily and cheaply

*Except for the cheapest and fastest machines, all computers were microprogrammed*

# Microprogramming: early 1980s

- Evolution bred more complex micro-machines
  - Complex instruction sets led to need for subroutine and call stacks in μcode
  - Need for fixing bugs in control programs was in conflict with read-only nature of μROM
  - → Writable Control Store (WCS) (B1700, QMachine, Intel i432, ...)
- With the advent of VLSI technology assumptions about ROM & RAM speed became invalid → more complexity
- Better compilers made complex instructions less important.
- Use of numerous micro-architectural innovations, e.g., pipelining, caches and buffers, made multiple-cycle execution of reg-reg instructions unattractive

# VAX 11-780 Microcode

; P1WFUD,1 [600,1205]  
; CALL2 .MIC [600,1205]

MICRO2 1F(12)  
Procedure call

26-May-81 14:58:1  
: CALLG, CALLS

VAX11/780 Microcode : PCS 01, FPLA 0D, WCS122

Page 771

```

;29744 ;HERE FOR CALLG OR CALLS, AFTER PROBING THE EXTENT OF THE STACK
;29745
;29746 =0 ;-----;CALL SITE FOR MPUSH
;29747 CALL.7: D_Q,AND,RC[T2], ;STRIP MASK TO BITS 11-0
;29748 CALL,J/MPUSH ;PUSH REGISTERS
;-----;RETURN FROM MPUSH
;29750 CACHE_D[LONG], ;PUSH PC
;29751 LAB_R[SP] ; BY SP
;-----;UPDATE SP FOR PUSH OF PC &
;29752 CACHE_D[LONG], ;PUSH PC
;29753 LAB_R[SP] ; BY SP
;-----;READY TO PUSH FRAME POINTER
;29754 CALL.8: R[SP]&VA_LA-K[,8]
;29755 ;-----;UPDATE SP FOR PUSH OF PC &
;29756 D_R[FPP]
;29757 ;-----;READY TO PUSH FRAME POINTER
;29758 =0 ;-----;CALL SITE FOR PSHSP
;29759 CACHE_D[LONG], ;STORE FP,
;29760 LAB_R[SP], ; GET SP AGAIN
;29761 SC_K[.FFF0], ;16 TO SC
;29762 CALL,J/PSHSP
;29763 ;-----;READY TO PUSH AP
;29764 D_R[AP], ;READY TO PUSH AP
;29765 Q_ID[PSL] ; AND GET PSW FOR COMBINATIO
;29766 ;-----;STORE OLD AP
;29767 CACHE_D[LONG], ;CLEAR PSW<T,N,Z,V,C>
;29768 Q_Q,ANDNOT,K[.1F], ;GET SP INTO LATCHES AGAIN
;29769 LAB_R[SP]
;-----;LOAD NEW PC AND CLEAR OUT
;29770 PC&VA_RC[T1], FLUSH,IB
;29771 ;-----;PSW TO D<31:16>
;29772 D_DAL,SC, ;RECOVER MASK
;29773 Q_RC[T2], ;PUT -13 IN SC
;29774 SC_SC+K[.3], ;START FETCHING SUBROUTINE I
;29775 LOAD,IB, PC_PC+1
;-----;MASK AND PSW IN D<31:03>
;29776 D_DAL,SC, ;GET LOW BITS OF OLD SP TO Q<1:0>
;29777 Q_RC[T4], ;PUT -3 IN SC
;29778 SC_SC+K[.A]
;-----;
```

# Writable Control Store (WCS)

- Implement control store in RAM not ROM
  - MOS SRAM memories now almost as fast as control store (core memories/DRAMs were 2-10x slower)
  - Bug-free microprograms difficult to write
- User-WCS provided as option on several minicomputers
  - Allowed users to change microcode for each processor
- User-WCS failed
  - Little or no programming tools support
  - Difficult to fit software into small space
  - Microcode control tailored to original ISA, less useful for others
  - Large WCS part of processor state - expensive context switches
  - Protection difficult if user can change microcode
  - Virtual memory required restartable microcode

# Microprogramming is far from extinct

- Played a crucial role in micros of the Eighties
  - DEC uVAX, Motorola 68K series, Intel 286/386
- Plays an assisting role in most modern micros
  - e.g., AMD Zen, Intel Sky Lake, Intel Atom, IBM PowerPC, ...
  - Most instructions executed directly, i.e., with hard-wired control
  - Infrequently-used and/or complicated instructions invoke microcode
- Patchable microcode common for post-fabrication bug fixes, e.g. Intel processors load µcode patches at bootup
  - Intel had to scramble to resurrect microcode tools and find original microcode engineers to patch Meltdown/Spectre security vulnerabilites

# Analyzing Microcoded Machines

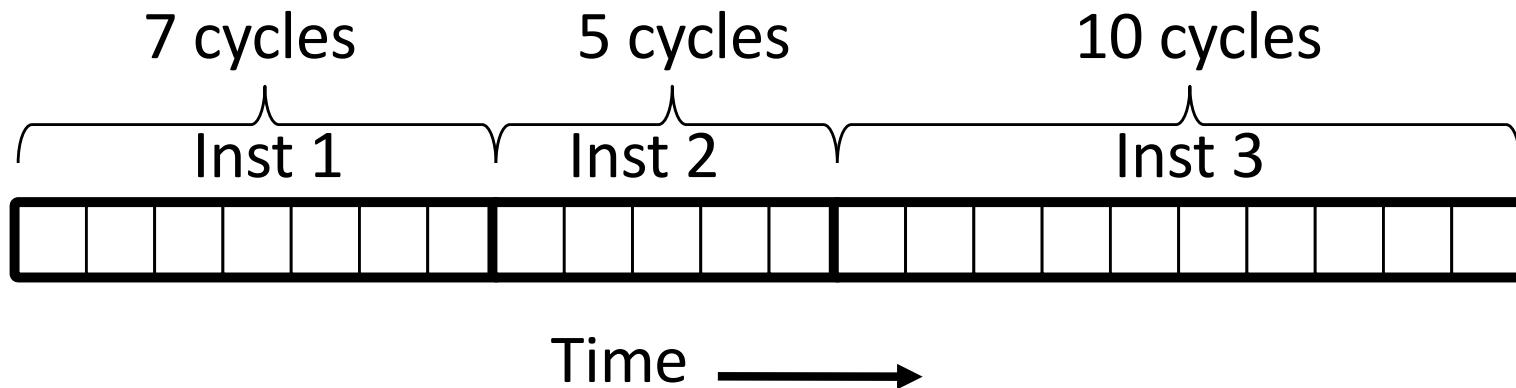
- John Cocke and group at IBM
  - Working on a simple pipelined processor, 801, and advanced compilers inside IBM
  - Ported experimental PL.8 compiler to IBM 370, and only used simple register-register and load/store instructions similar to 801
  - Code ran faster than other existing compilers that used all 370 instructions! (up to 6MIPS whereas 2MIPS considered good before)
- Emer, Clark, at DEC
  - Measured VAX-11/780 using external hardware
  - Found it was actually a 0.5MIPS machine, although usually assumed to be a 1MIPS machine
  - Found 20% of VAX instructions responsible for 60% of microcode, but only account for 0.2% of execution time!
- VAX8800
  - Control Store: 16K\*147b RAM, Unified Cache: 64K\*8b RAM
  - 4.5x more microstore RAM than cache RAM!

# “Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

- Instructions per program depends on source code, compiler technology, and ISA
- Cycles per instructions (CPI) depends on ISA and μarchitecture
- Time per cycle depends upon the μarchitecture and base technology

# CPI for Microcoded Machine



$$\text{Total clock cycles} = 7+5+10 = 22$$

$$\text{Total instructions} = 3$$

$$\text{CPI} = 22/3 = 7.33$$

*CPI is always an average over a large number of instructions.*

# IC Technology Changes Tradeoffs

- Logic, RAM, ROM all implemented using MOS transistors
- Semiconductor RAM ~ same speed as ROM

# Reconsidering Microcode Machine (microcoded 68000 example)

RISC!

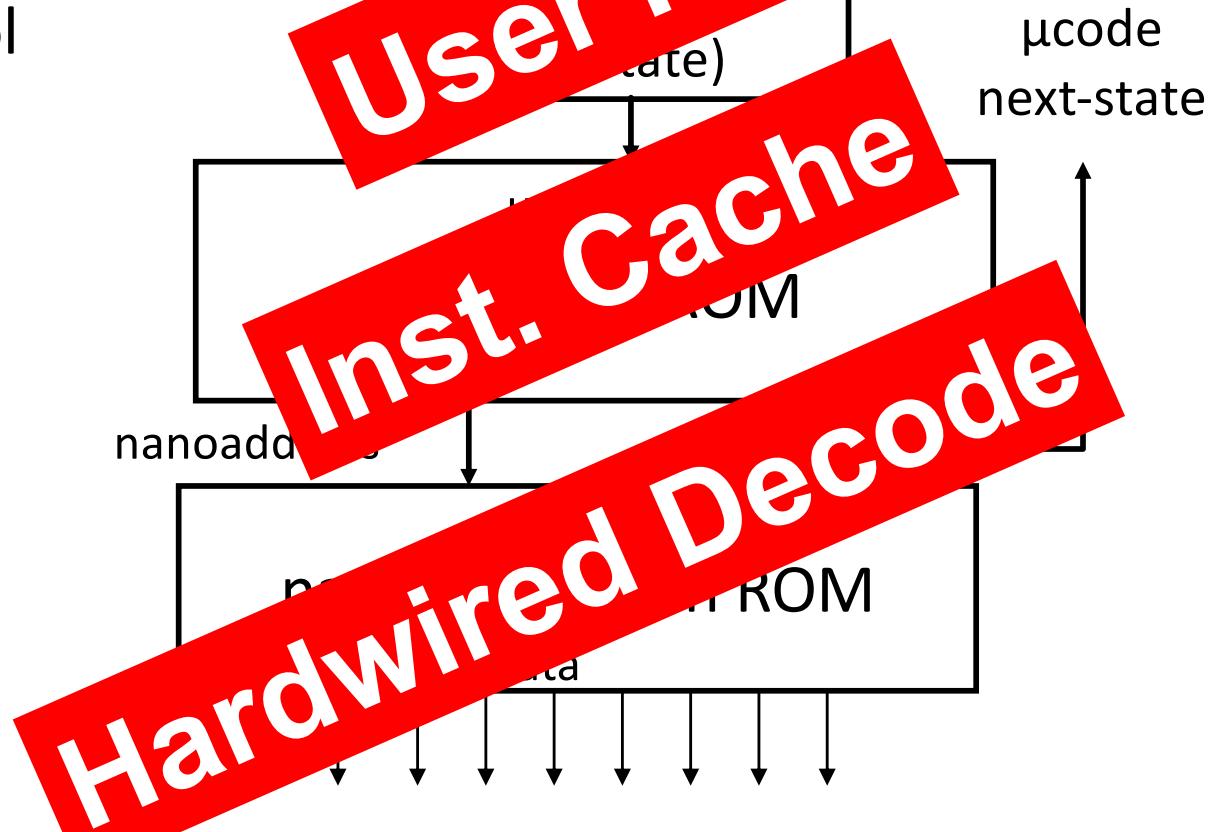
Exploits recurring control  
signal patterns in μcode,  
e.g.,

ALU0 A ← Reg[rs1]

...

ALUI0 A ← Reg[rs1]

...



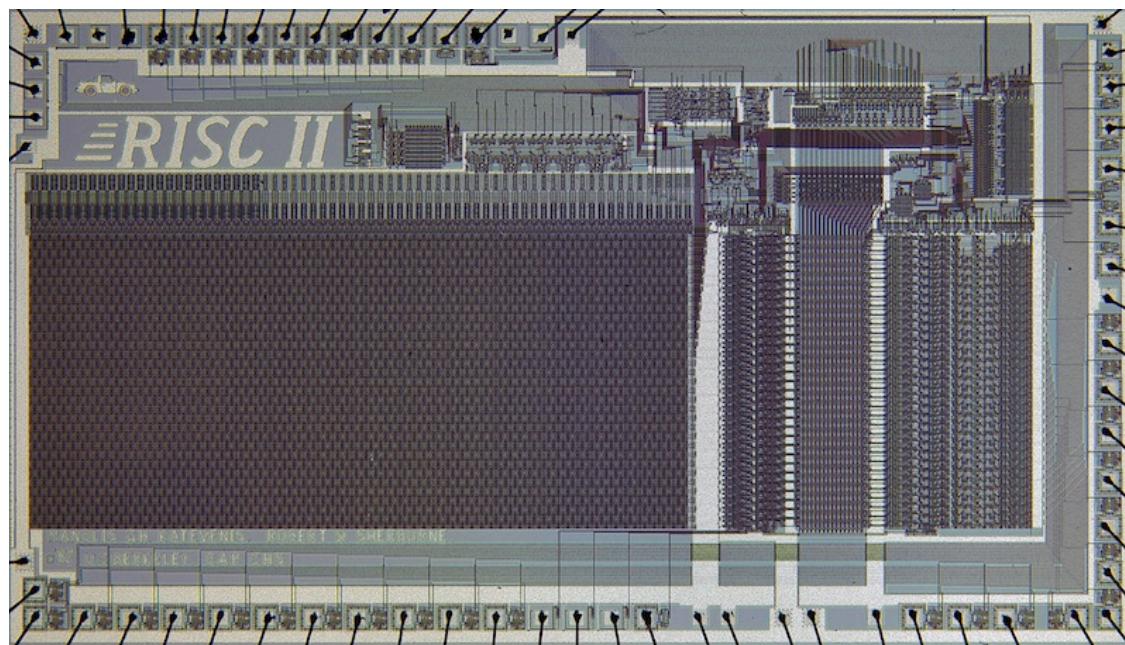
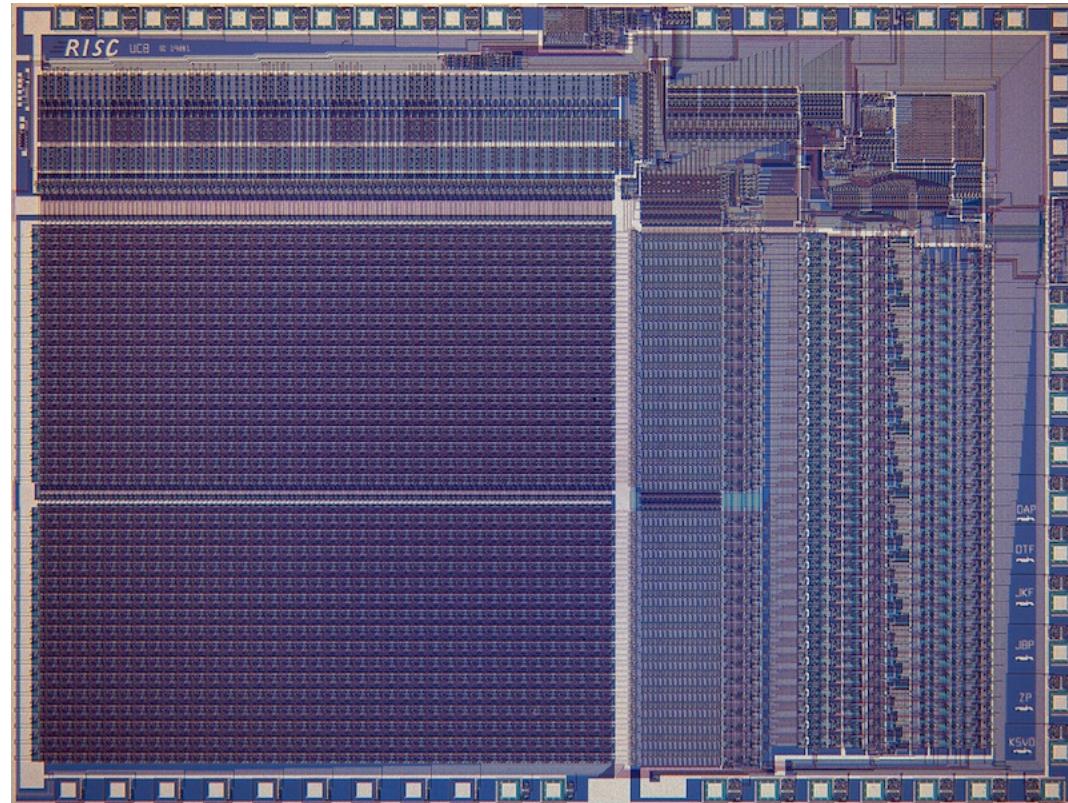
- Motorola 68000 had 17-bit μcode containing either 10-bit μjump or 9-bit nanoinstruction pointer
  - Nanoinstructions were 68 bits wide, decoded to give 196 control signals

# From CISC to RISC

- Use fast RAM to build fast instruction *cache* of user-visible instructions, not fixed hardware microroutines
  - Contents of fast instruction memory change to fit application needs
- Use simple ISA to enable hardwired pipelined implementation
  - Most compiled code only used few CISC instructions
  - Simpler encoding allowed pipelined implementations
  - RISC ISA comparable to vertical microcode
- Further benefit with integration
  - In early ‘80s, finally fit 32-bit datapath + small caches on single chip
  - No chip crossings in common case allows faster operation

# Berkeley RISC Chips

RISC-I (1982) Contains 44,420 transistors, fabbed in 5  $\mu\text{m}$  NMOS, with a die area of 77  $\text{mm}^2$ , ran at 1 MHz. This chip is probably the first VLSI RISC.



RISC-II (1983) contains 40,760 transistors, was fabbed in 3  $\mu\text{m}$  NMOS, ran at 3 MHz, and the size is 60  $\text{mm}^2$ .

**Stanford** built some too...

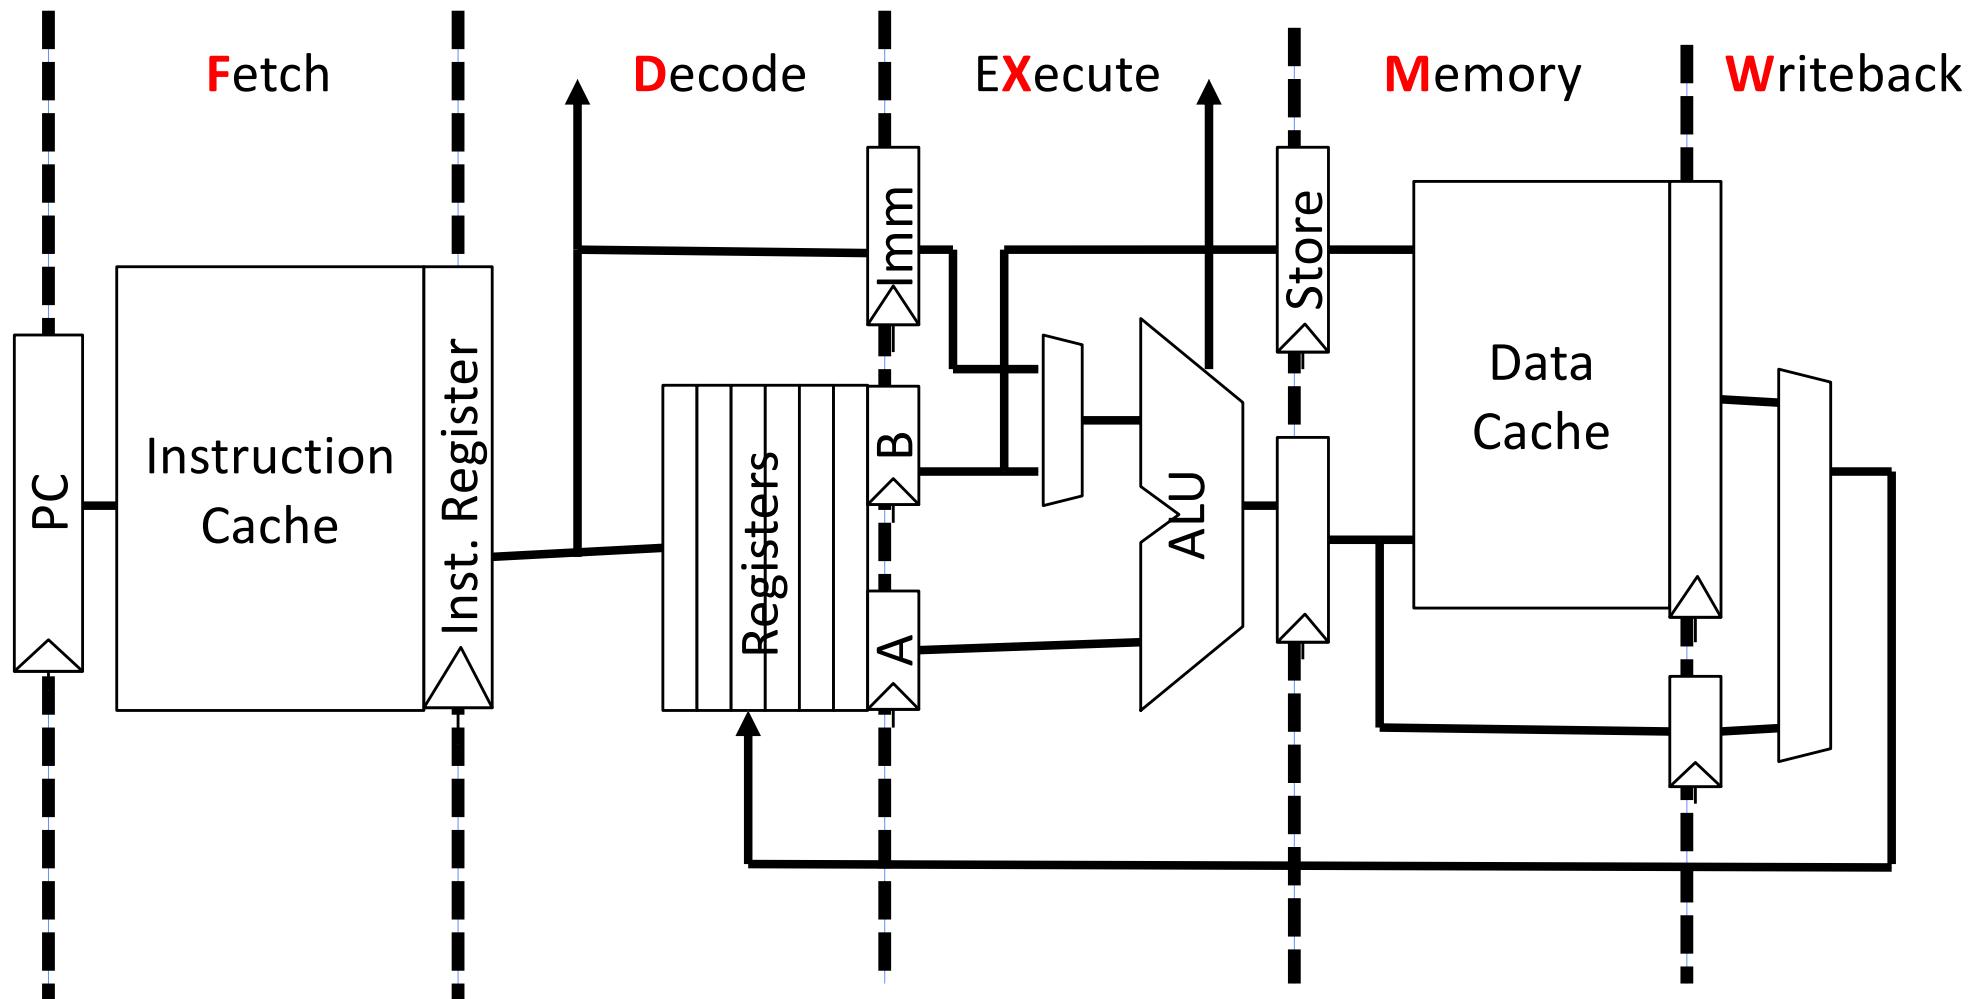
# “Iron Law” of Processor Performance

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

- Instructions per program depends on source code, compiler technology, and ISA
- Cycles per instructions (CPI) depends on ISA and μarchitecture
- Time per cycle depends upon the μarchitecture and base technology

Microarchitecture	CPI	cycle time
Microcoded	>1	short
Single-cycle unpipelined	1	long
Pipelined	1	short

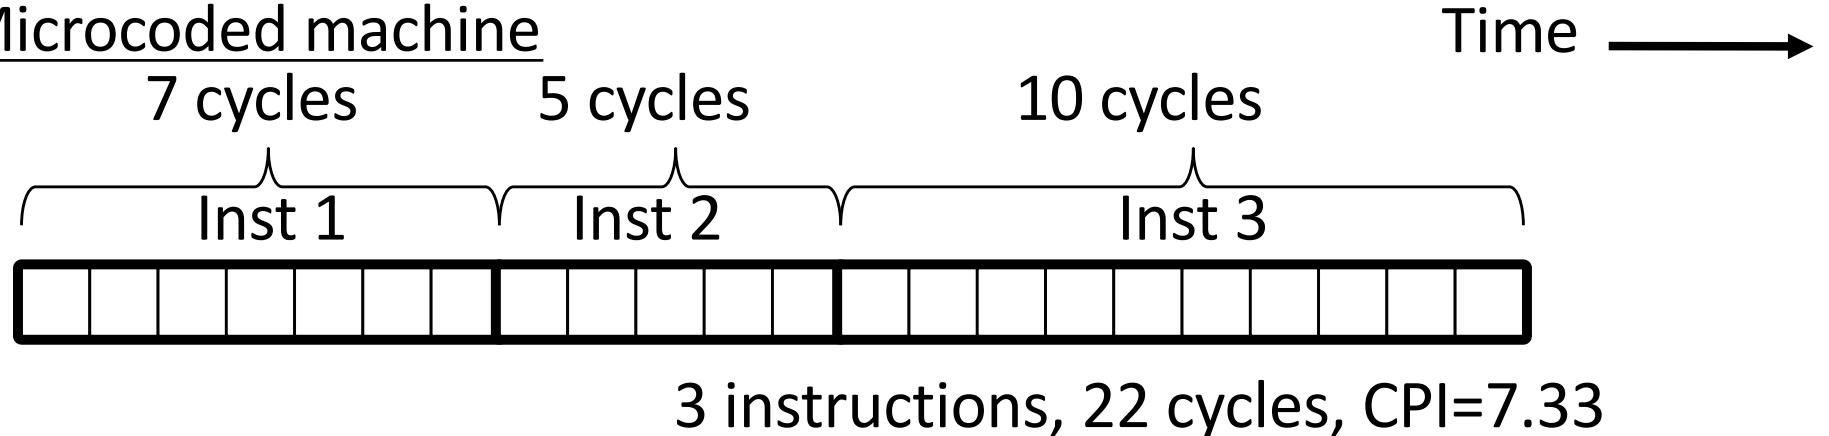
# Classic 5-Stage RISC Pipeline



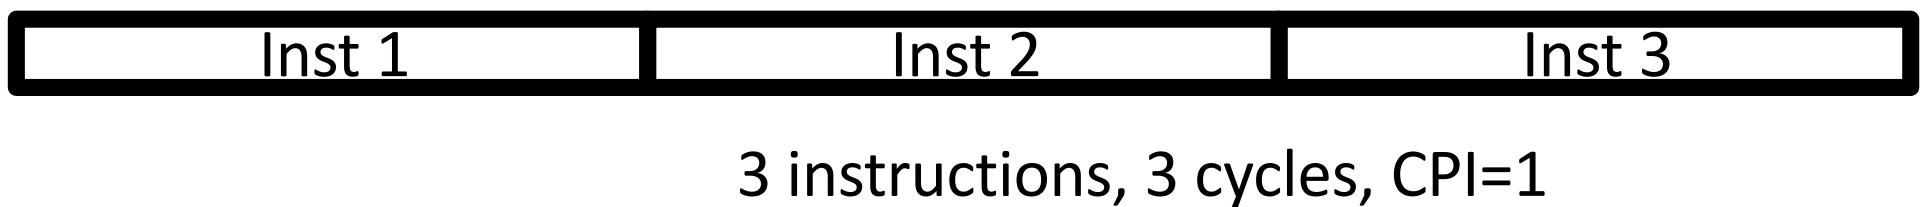
*This version designed for regfiles/memories  
with synchronous reads and writes.*

# CPI Examples

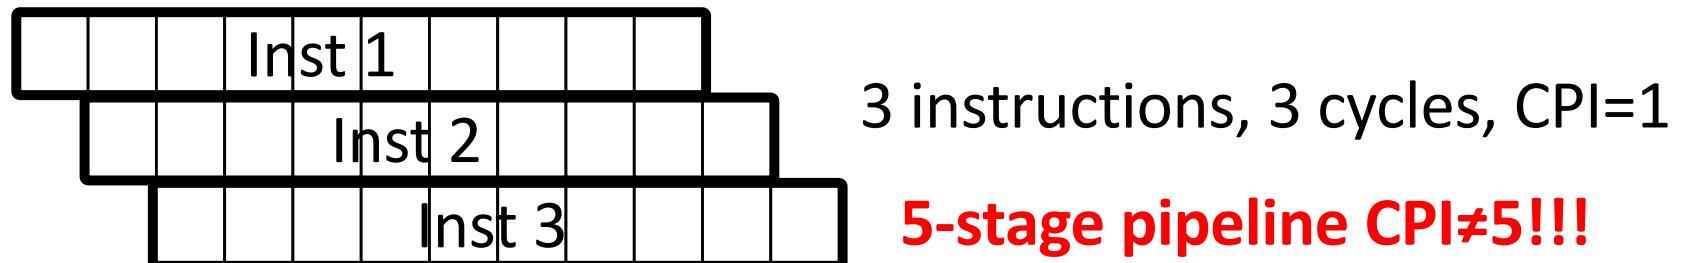
## Microcoded machine



## Unpipelined machine



## Pipelined machine



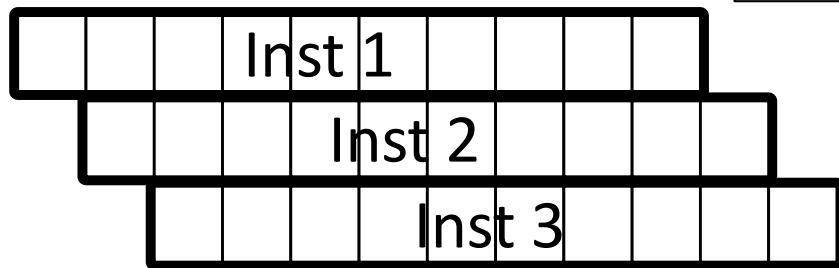
# Instructions interact with each other in pipeline

- An instruction in the pipeline may need a resource being used by another instruction in the pipeline → *structural hazard*
- An instruction may depend on something produced by an earlier instruction
  - Dependence may be for a data value  
→ *data hazard*
  - Dependence may be for the next instruction's address  
→ *control hazard (branches, exceptions)*
- Handling hazards generally introduces bubbles into pipeline and reduces ideal CPI > 1

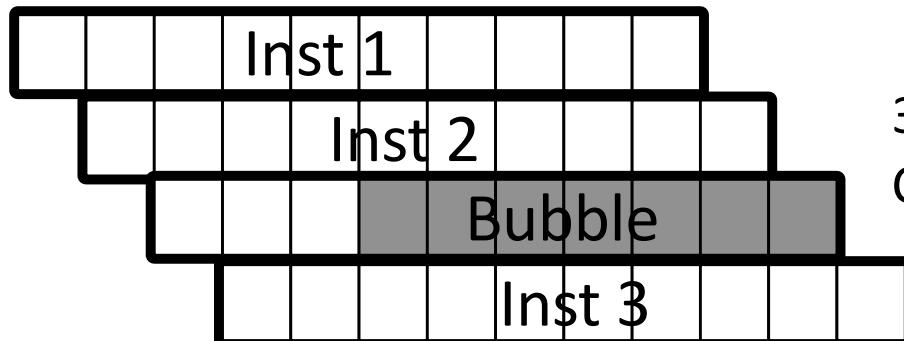
# Pipeline CPI Examples

Time →

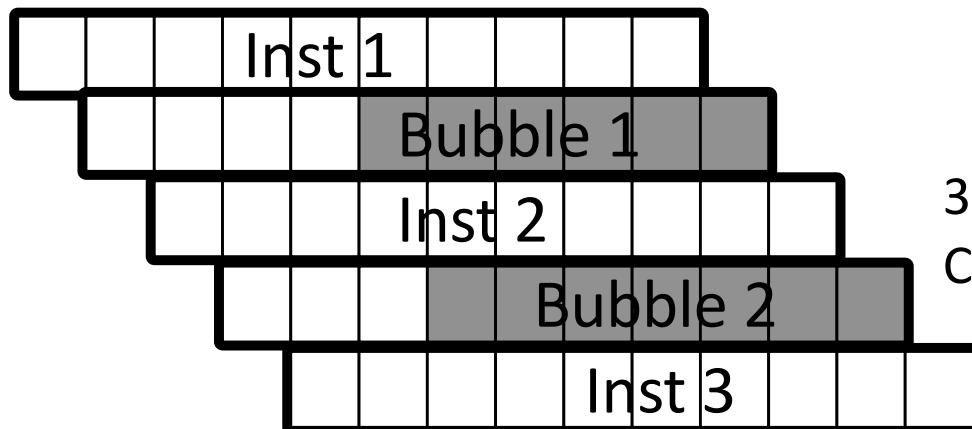
*Measure from when first instruction finishes to when last instruction in sequence finishes.*



3 instructions finish in 3 cycles  
CPI =  $3/3 = 1$



3 instructions finish in 4 cycles  
CPI =  $4/3 = 1.33$



3 instructions finish in 5cycles  
CPI =  $5/3 = 1.67$

# Resolving Structural Hazards

- Structural hazard occurs when two instructions need same hardware resource at same time
  - Can resolve in hardware by stalling newer instruction till older instruction finished with resource
- A structural hazard can always be avoided by adding more hardware to design
  - E.g., if two instructions both need a port to memory at same time, could avoid hazard by adding second port to memory
- Classic RISC 5-stage integer pipeline has no structural hazards by design
  - Many RISC implementations have structural hazards on multi-cycle units such as multipliers, dividers, floating-point units, etc., and can have on register writeback ports

# Types of Data Hazards

Consider executing a sequence of register-register instructions of type:

$$r_k \leftarrow r_i \text{ op } r_j$$

Data-dependence

$$\begin{aligned} r_3 &\leftarrow r_1 \text{ op } r_2 \\ r_5 &\leftarrow r_3 \text{ op } r_4 \end{aligned}$$

Read-after-Write  
(RAW) hazard

Anti-dependence

$$\begin{aligned} r_3 &\leftarrow r_1 \text{ op } r_2 \\ r_1 &\leftarrow r_4 \text{ op } r_5 \end{aligned}$$

Write-after-Read  
(WAR) hazard

Output-dependence

$$\begin{aligned} r_3 &\leftarrow r_1 \text{ op } r_2 \\ r_3 &\leftarrow r_6 \text{ op } r_7 \end{aligned}$$

Write-after-Write  
(WAW) hazard

# Three Strategies for Data Hazards

- Interlock

- Wait for hazard to clear by holding dependent instruction in issue stage

- Bypass

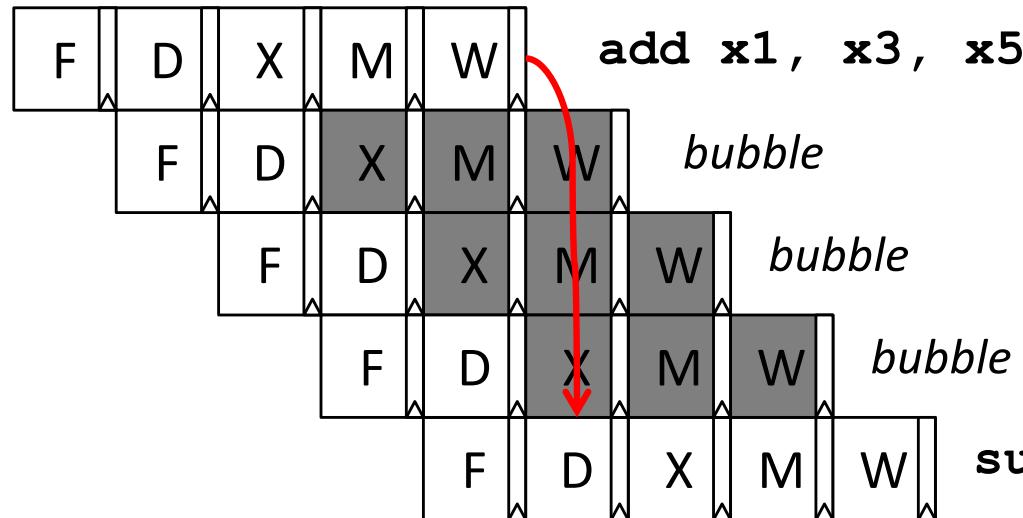
- Resolve hazard earlier by bypassing value as soon as available

- Speculate

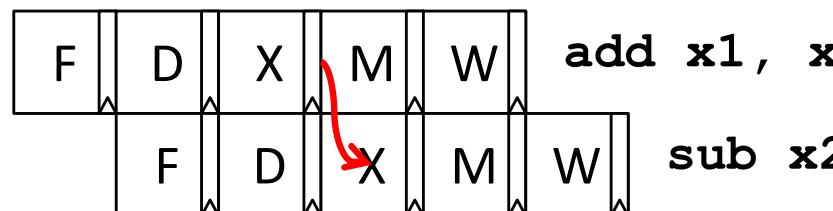
- Guess on value, correct if wrong

# Interlocking Versus Bypassing

add x1, x3, x5  
sub x2, **x1, x4**

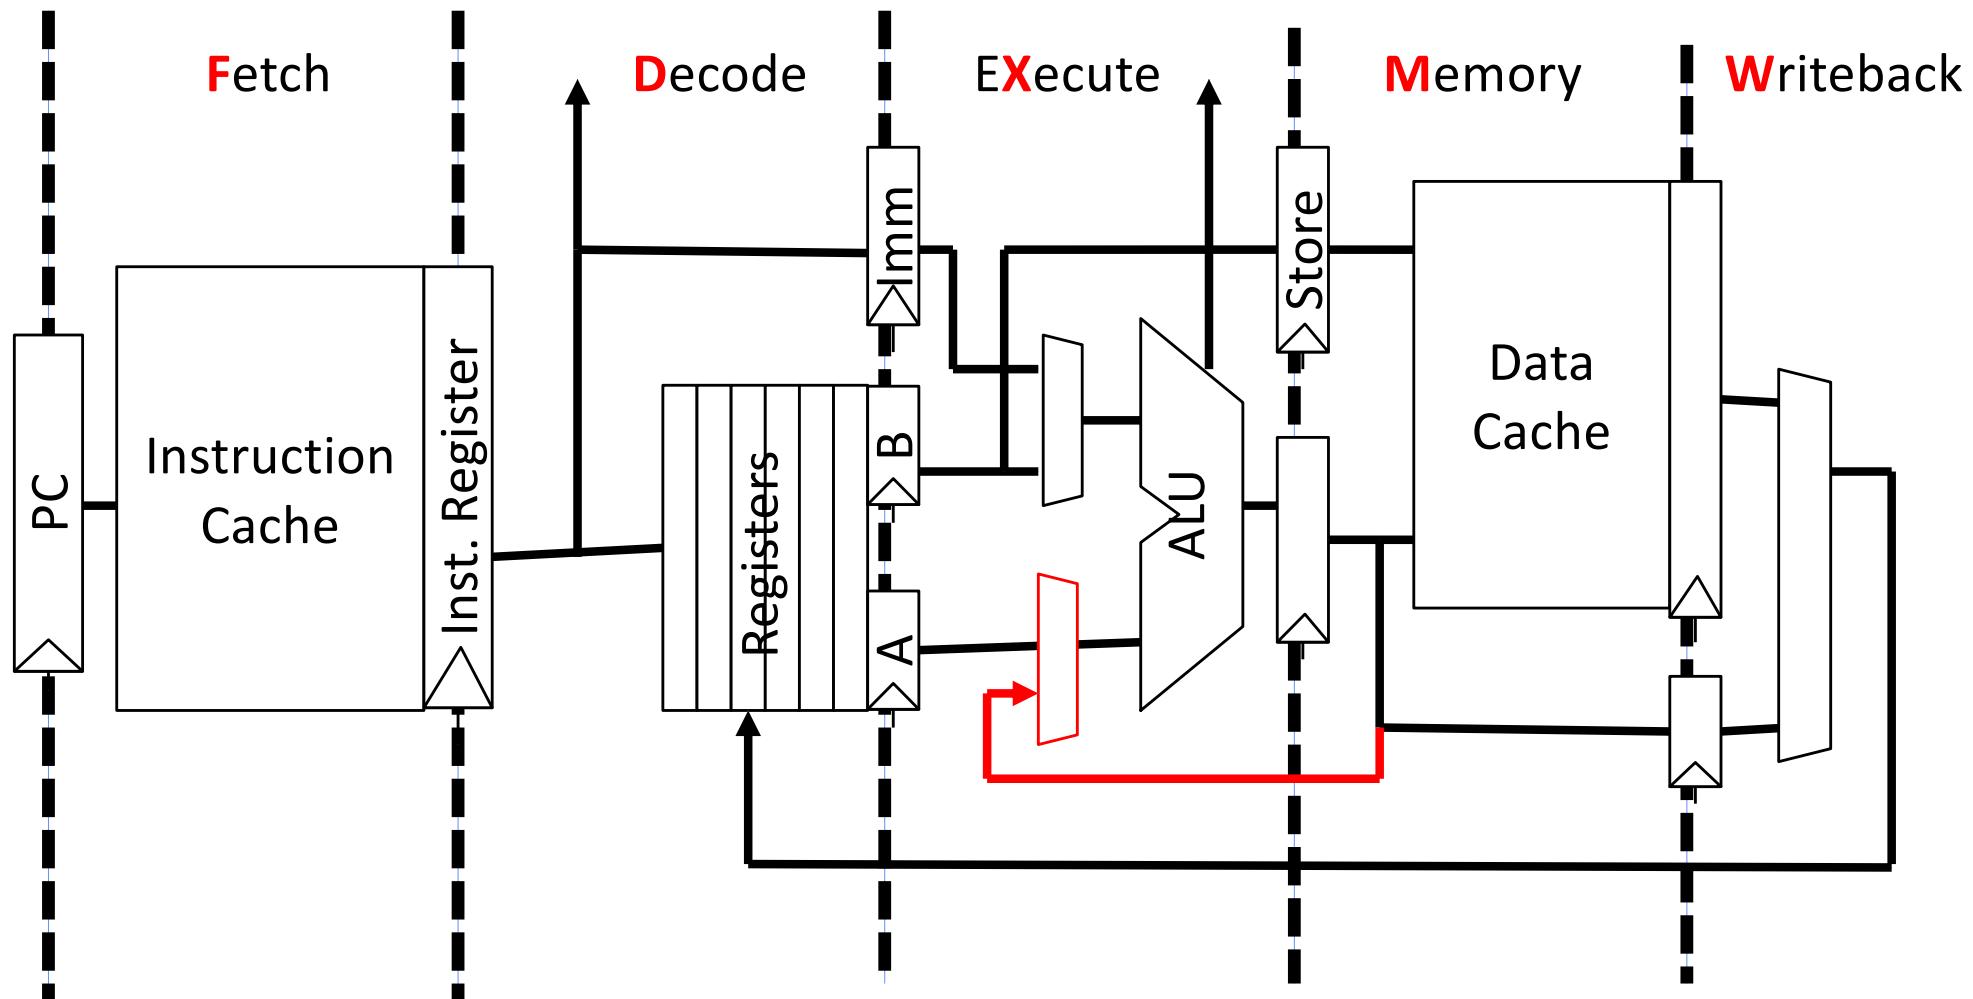


Instruction interlocked  
in decode stage

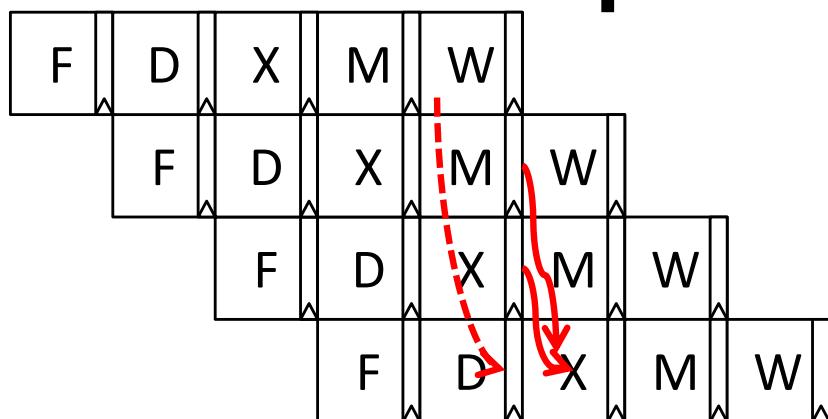
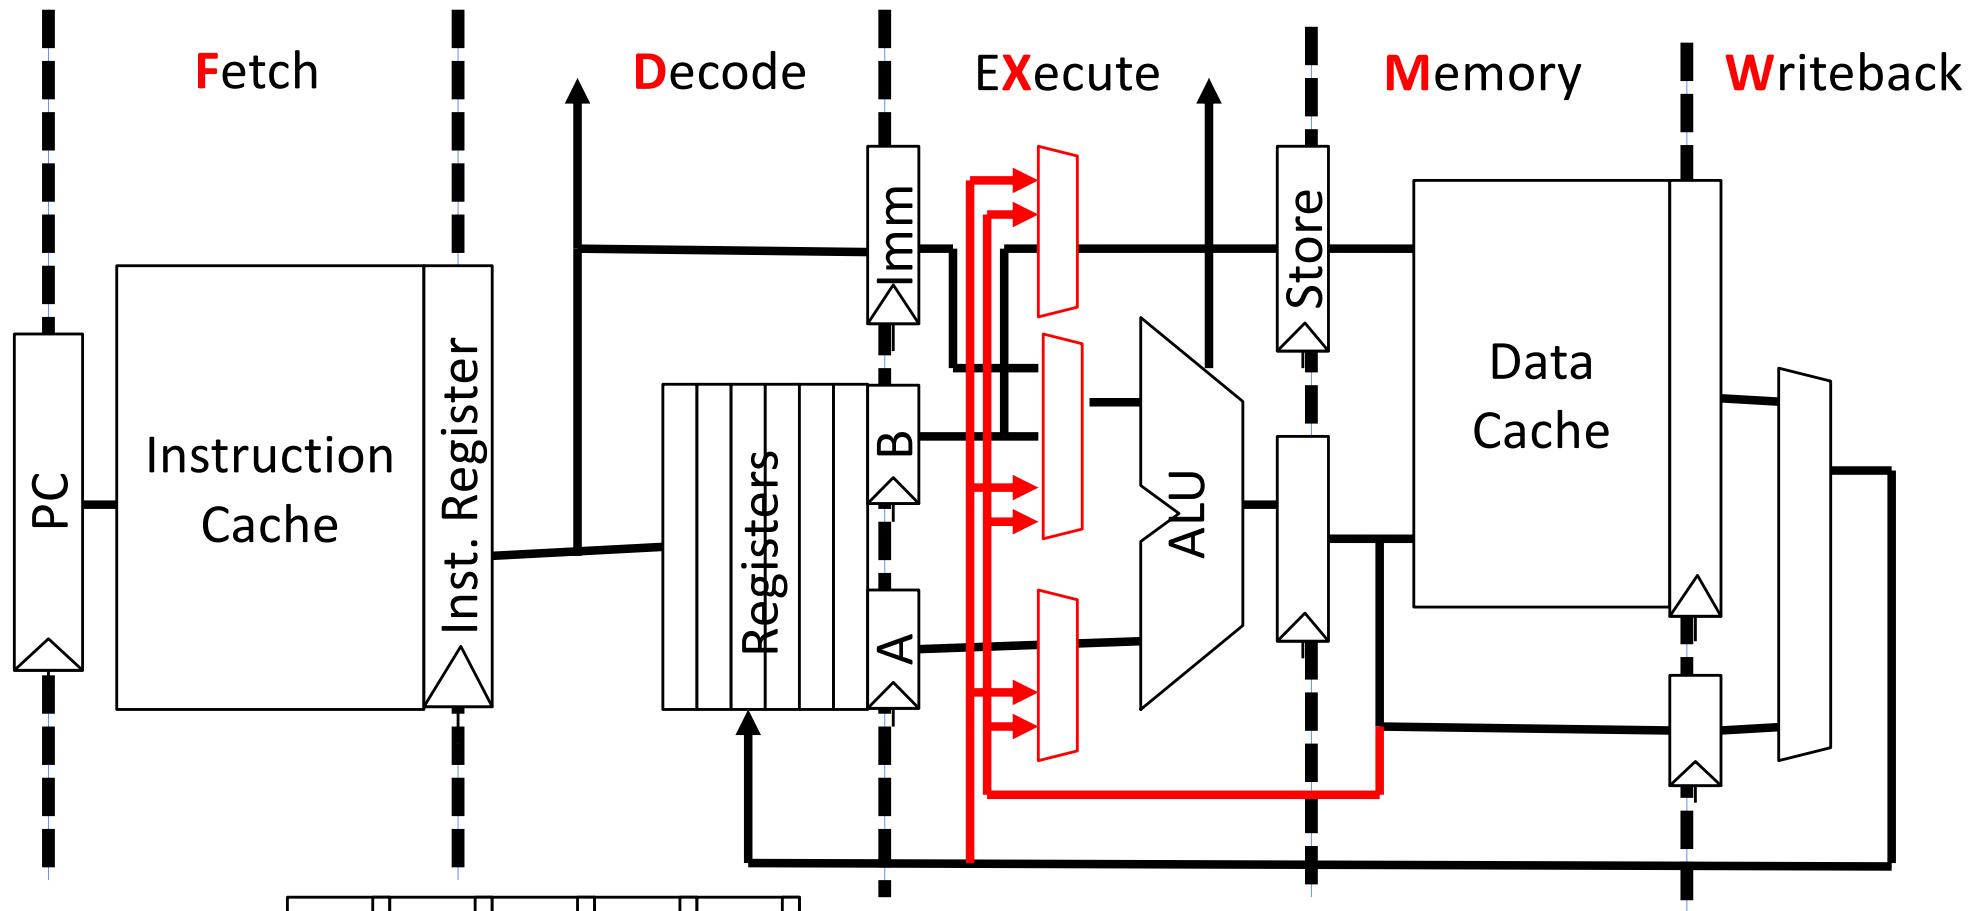


Bypass around ALU  
with no bubbles

# Example Bypass Path



# Fully Bypassed Data Path



[ Assumes data written to registers in a W cycle is readable in parallel D cycle (dotted line). Extra write data register and bypass paths required if this is not possible. ]

# Value Speculation for RAW Data Hazards

- Rather than wait for value, can guess value!
- So far, only effective in certain limited cases:
  - Branch prediction
  - Stack pointer updates
  - Memory address disambiguation

# CS152 Administrivia

- PS 1 is posted, due 11:59PM on Monday Feb 8
- Lab 1 out today, due 11:59PM Wed Feb 17
- Lab 1 overview in Sections on Friday
  - There will be live demo walking through lab install process
- Midterm and final exams will use remote proctoring
  - Details TBD, awaiting further instruction from campus

# CS252A Administrivia

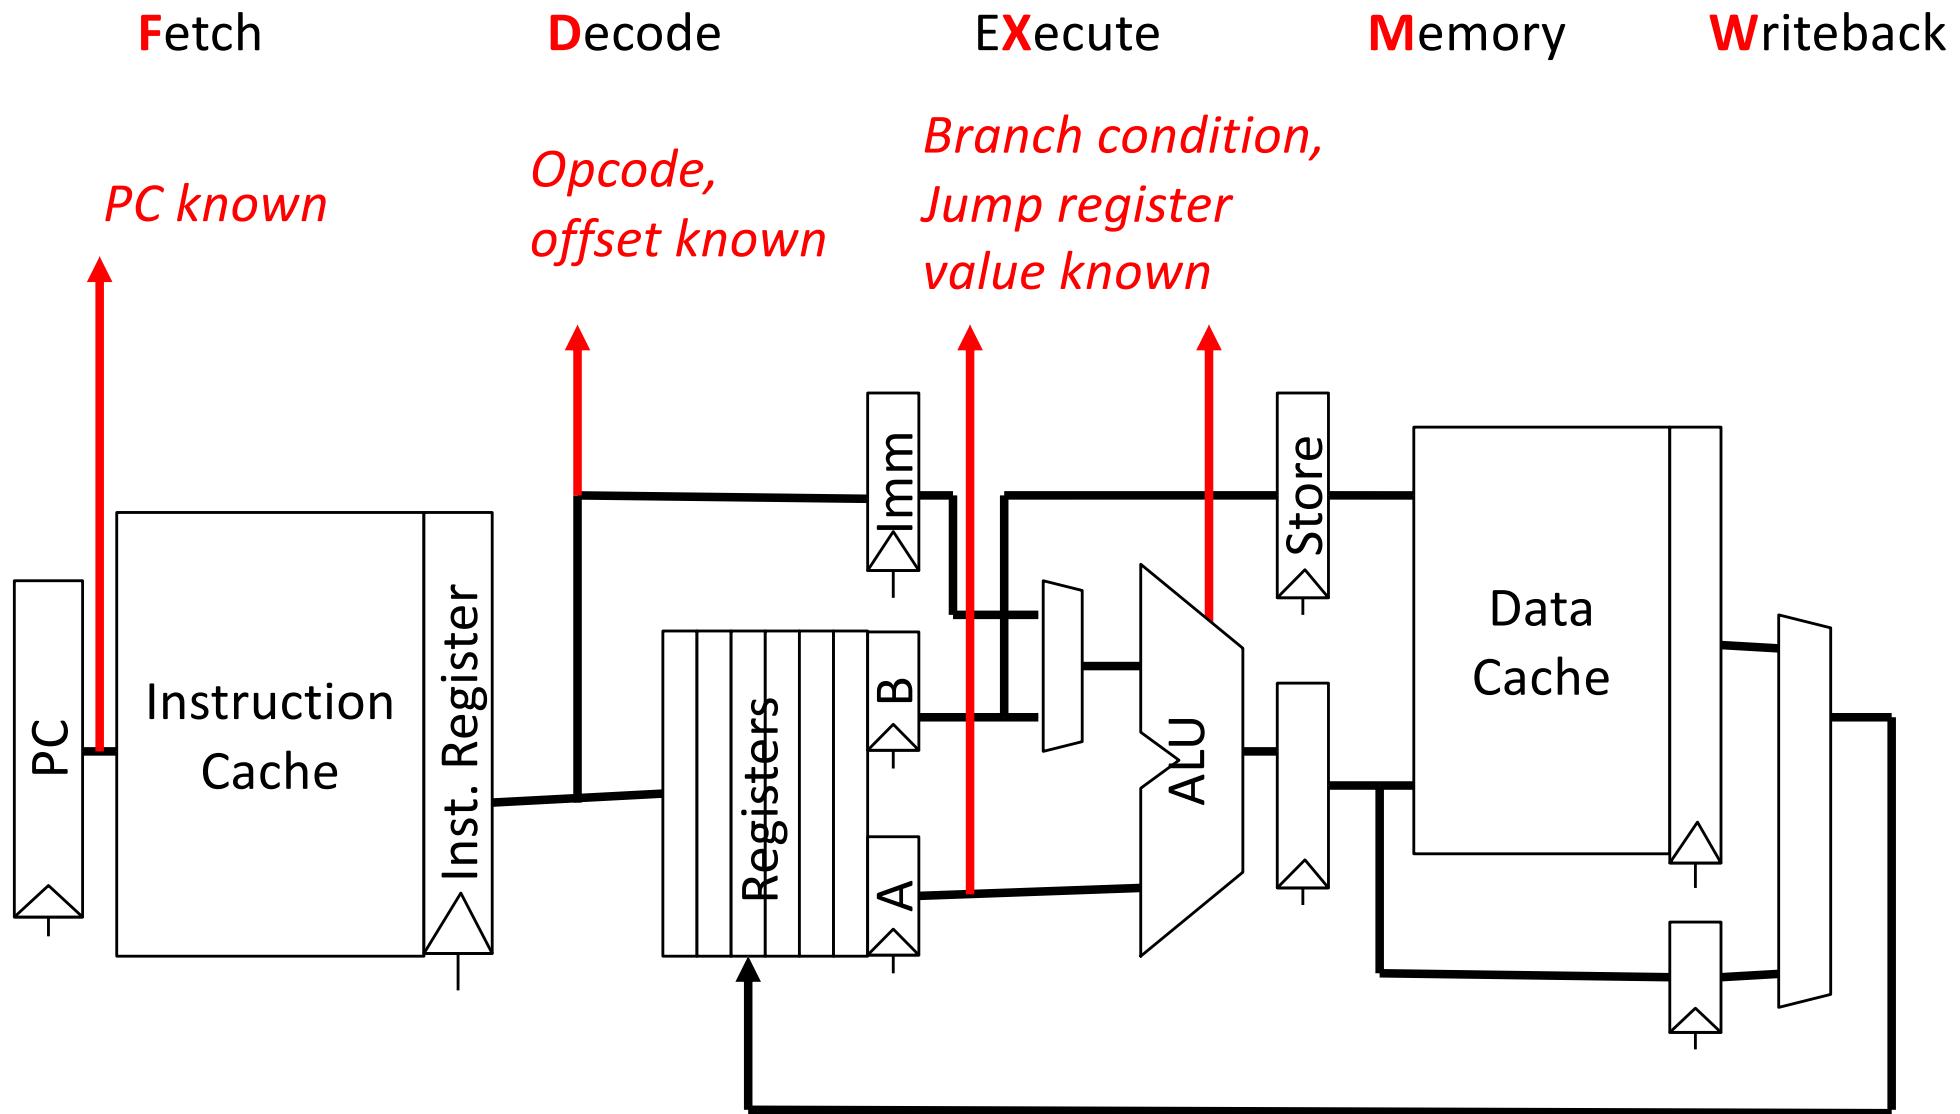
- CS252 discussions grading policy
  - We'll ignore your two lowest scores in grading, which includes absences
  - Send in summary even if you can't attend discussion
- CS252 Piazza class has been created
  - Sign up for this as well as CS152 Piazza
- Each CS252 paper has dedicated thread
  - Post your response as private note to instructors
  - Due Wednesday evening before Thursday discussion section

# Control Hazards

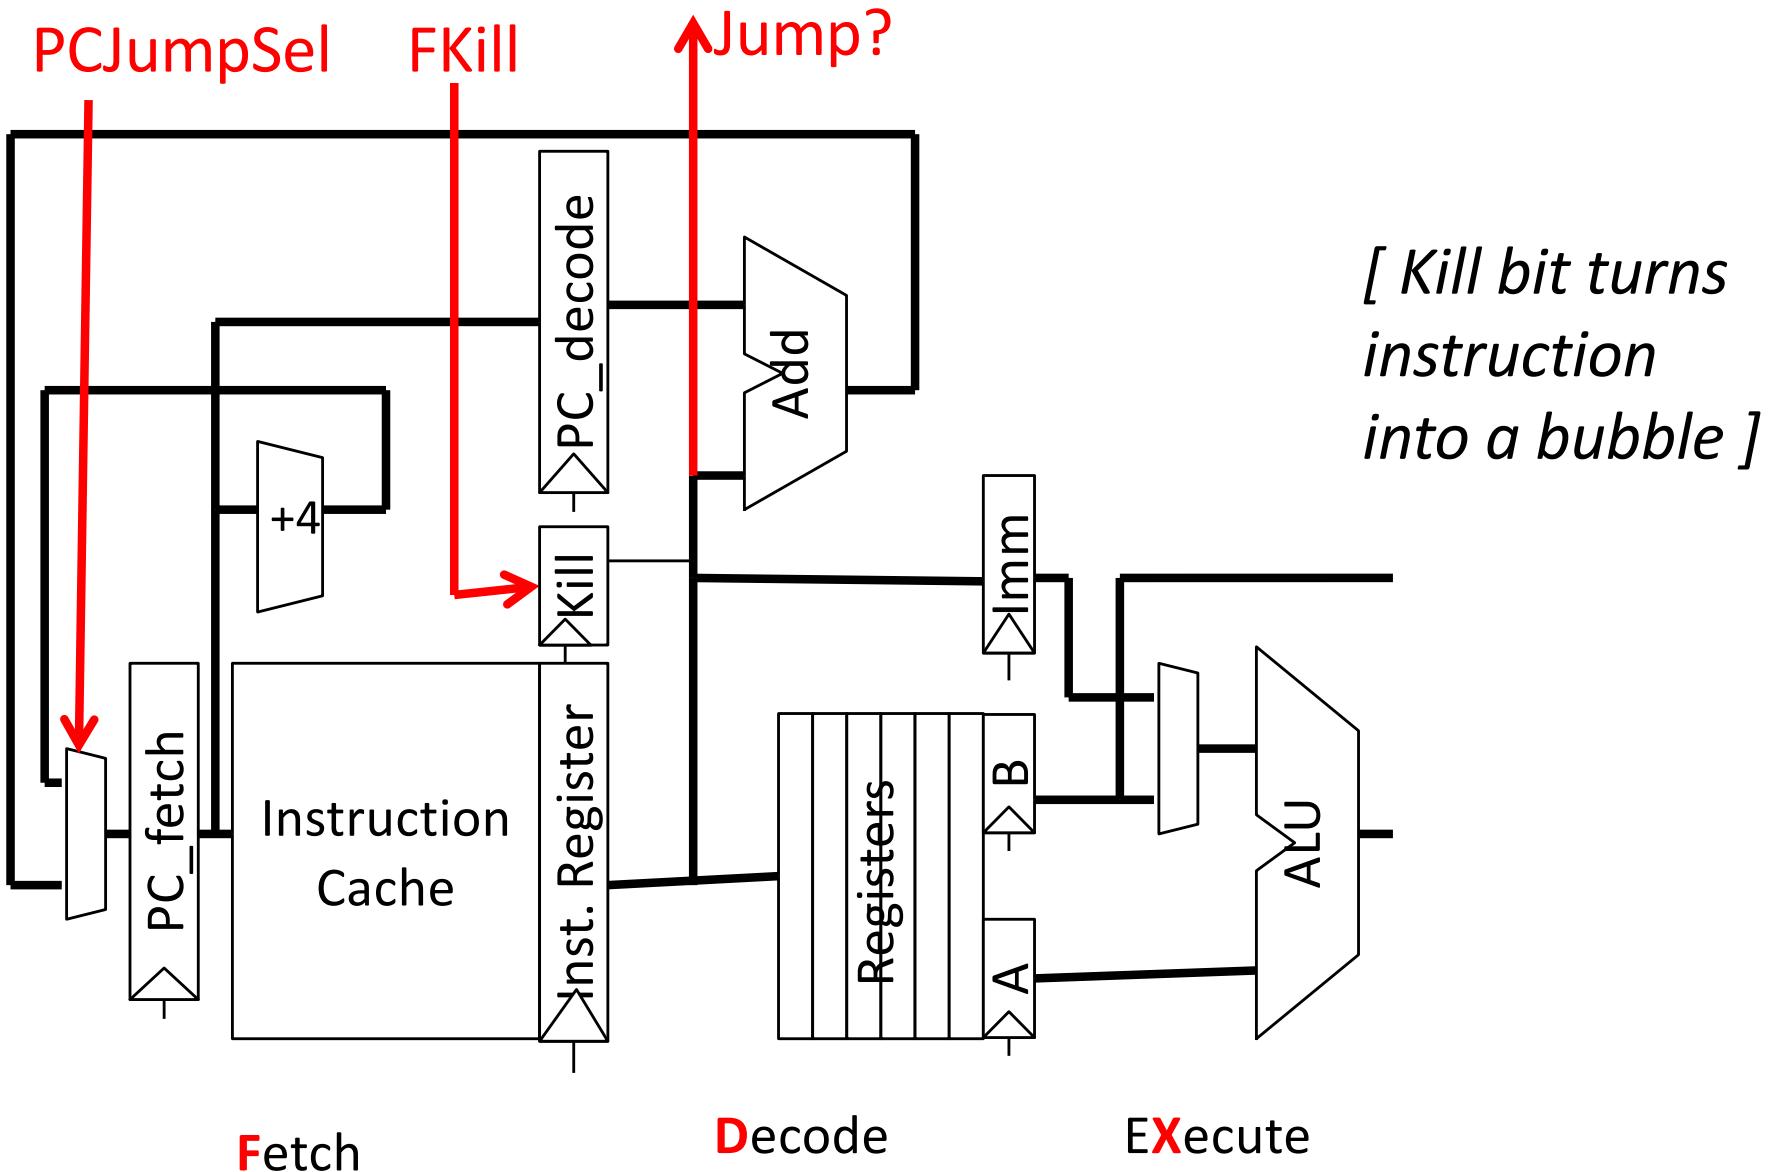
What do we need to calculate next PC?

- For Unconditional Jumps
  - Opcode, PC, and offset
- For Jump Register
  - Opcode, Register value, and offset
- For Conditional Branches
  - Opcode, Register (for condition), PC and offset
- For all other instructions
  - Opcode and PC ( and have to know it's not one of above )

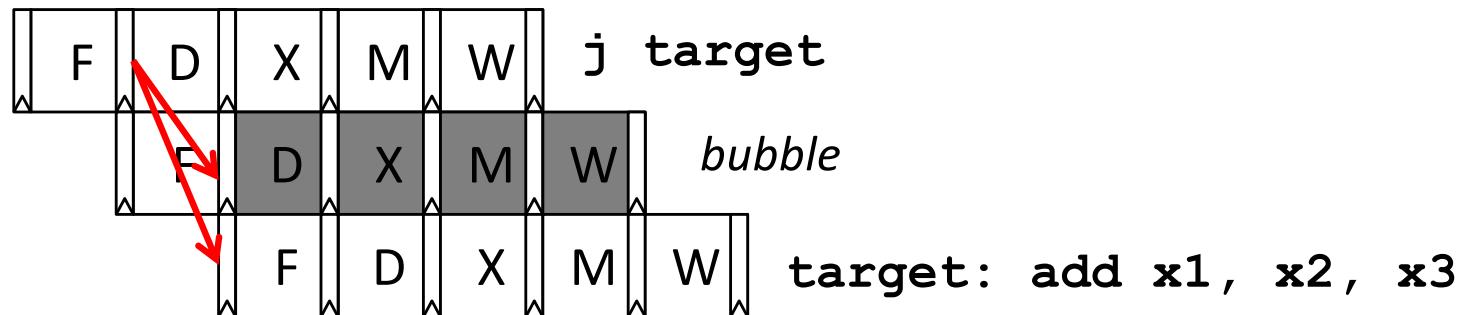
# Control flow information in pipeline



# RISC-V Unconditional PC-Relative Jumps



# Pipelining for Unconditional PC-Relative Jumps



# Branch Delay Slots

- Early RISCs adopted idea from pipelined microcode engines, and changed ISA semantics so instruction *after* branch/jump is always executed before control flow change occurs:

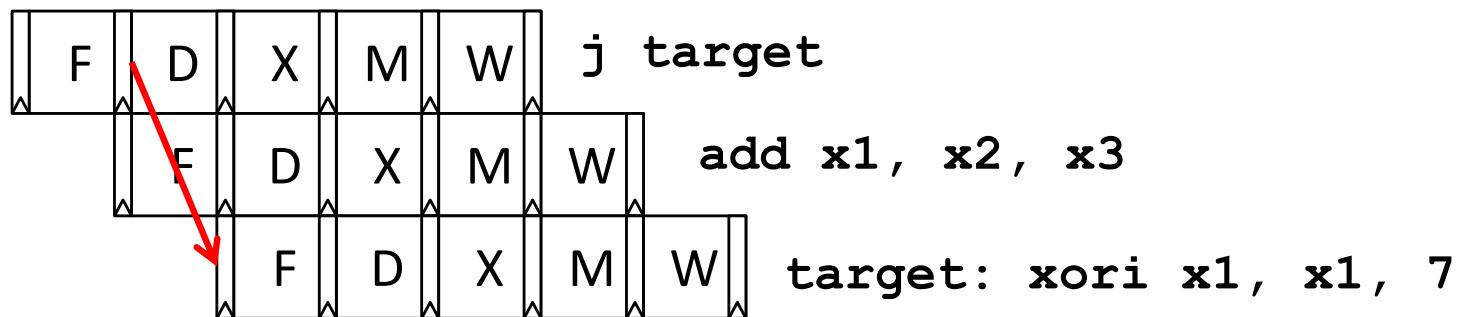
```
0x100 j target
```

```
0x104 add x1, x2, x3 // Executed before target
```

...

```
0x205 target: xori x1, x1, 7
```

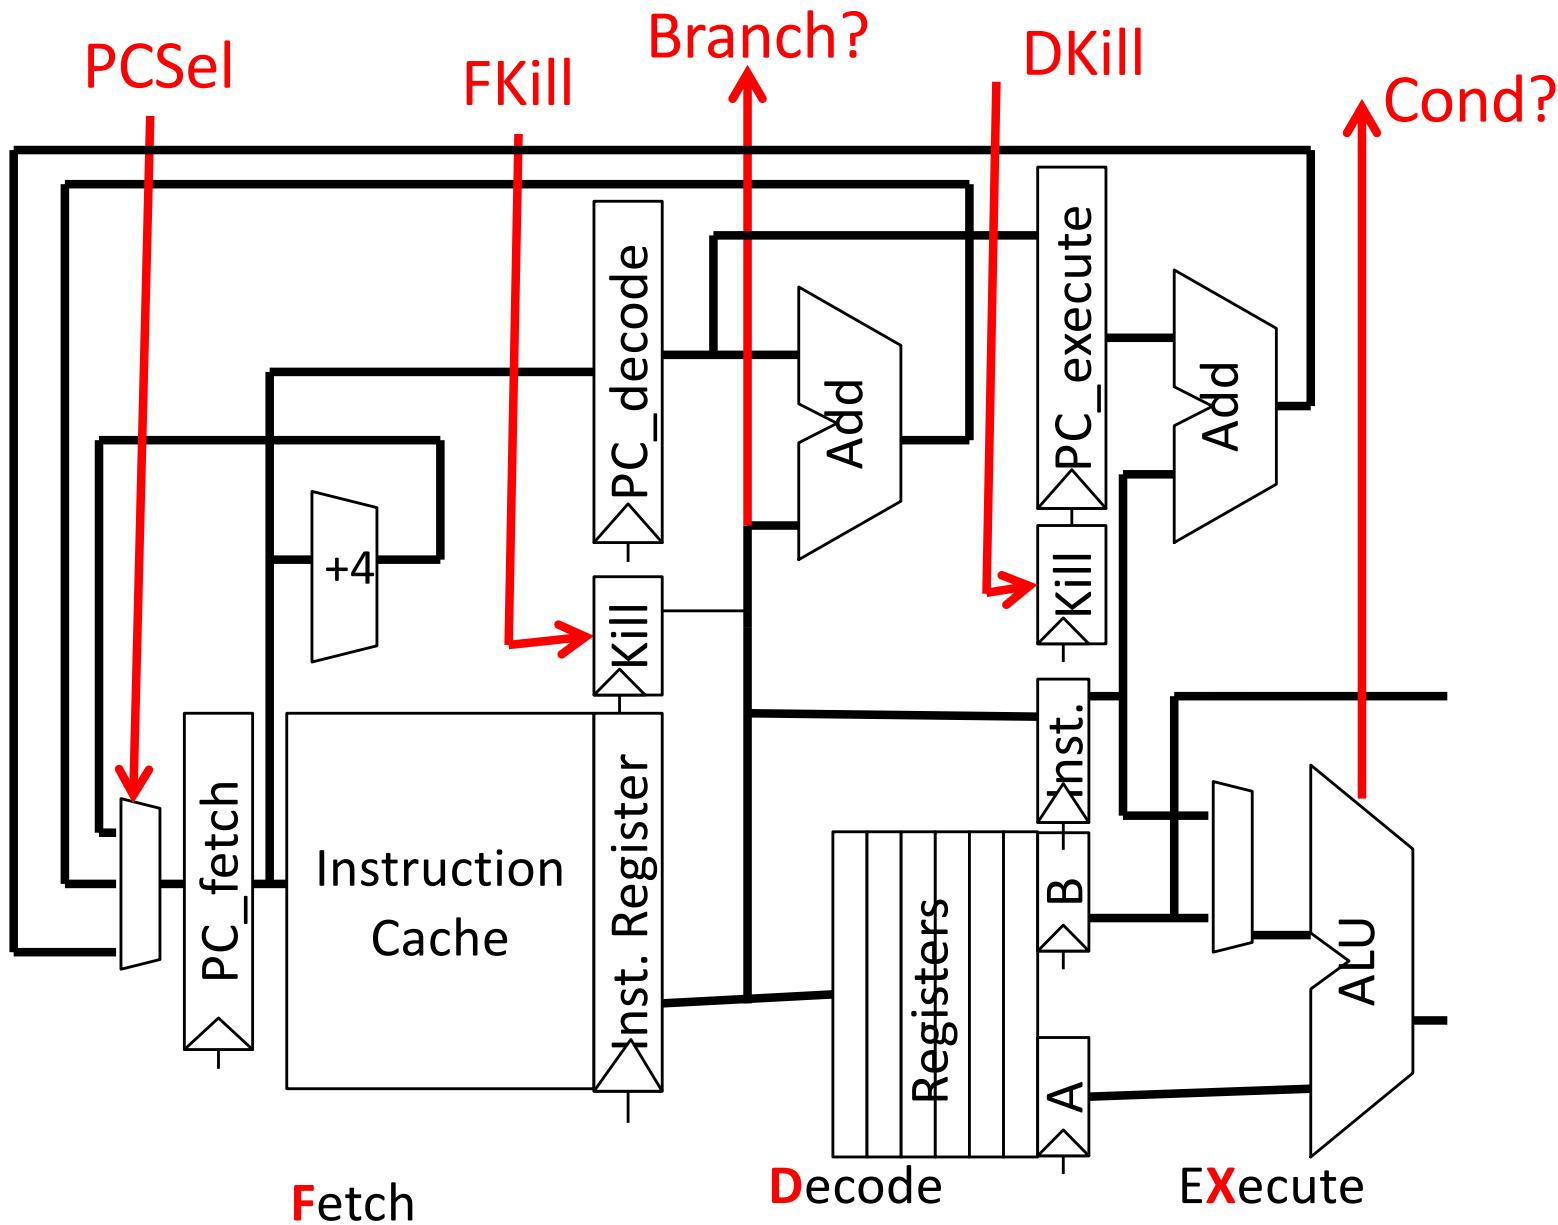
- Software has to fill delay slot with useful work, or fill with explicit NOP instruction



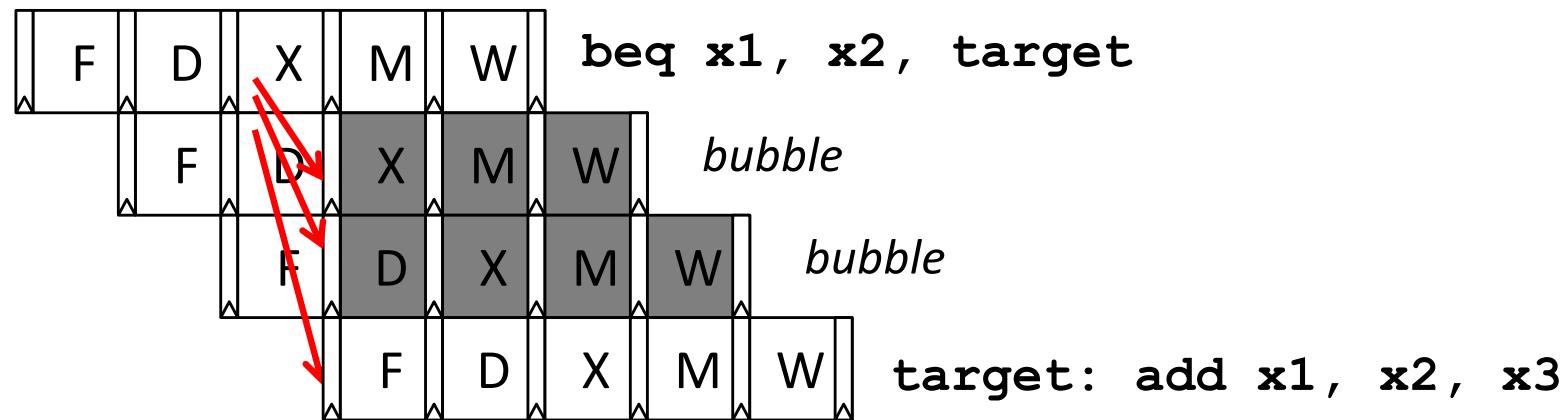
# **Post-1990 RISC ISAs don't have delay slots**

- Encodes microarchitectural detail into ISA
  - c.f. IBM 650 drum layout
- Performance issues
  - Increased I-cache misses from NOPs in unused delay slots
  - I-cache miss on delay slot causes machine to wait, even if delay slot is a NOP
- Complicates more advanced microarchitectures
  - Consider 30-stage pipeline with four-instruction-per-cycle issue
- Better branch prediction reduced need
  - Branch prediction in later lecture

# RISC-V Conditional Branches

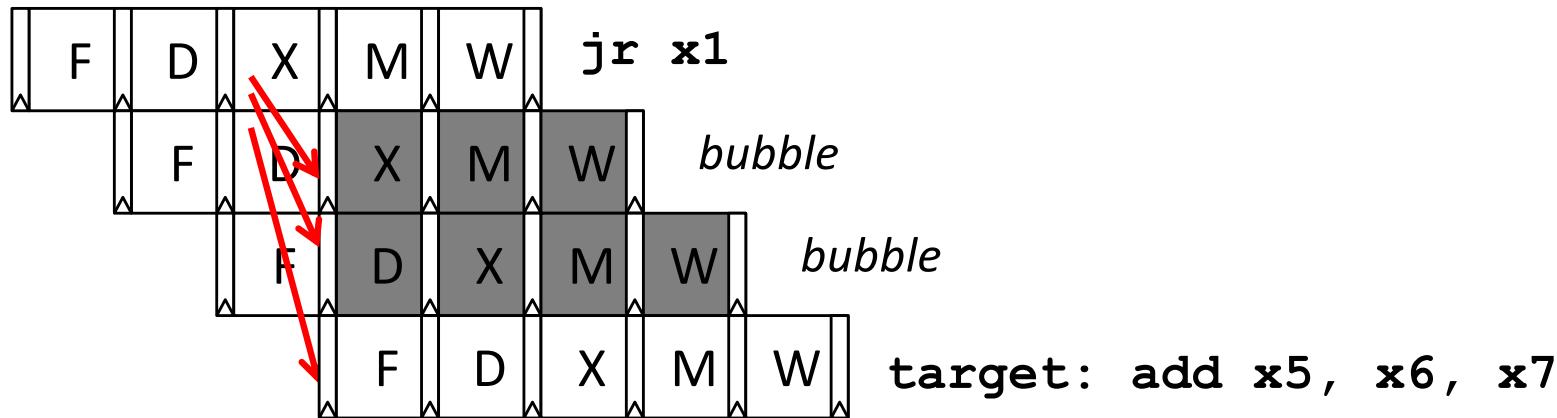


# Pipelining for Conditional Branches



# Pipelining for Jump Register

- Register value obtained in execute stage



# Why instruction may not be dispatched every cycle in classic 5-stage pipeline (CPI>1)

- Full bypassing may be too expensive to implement
  - typically all frequently used paths are provided
  - some infrequently used bypass paths may increase cycle time and counteract the benefit of reducing CPI
- Loads have two-cycle latency
  - Instruction after load cannot use load result
  - MIPS-I ISA defined *load delay slots*, a software-visible pipeline hazard (compiler schedules independent instruction or inserts NOP to avoid hazard). Removed in MIPS-II (pipeline interlocks added in hardware)
    - MIPS: “Microprocessor without Interlocked Pipeline Stages”
- Jumps/Conditional branches may cause bubbles
  - kill following instruction(s) if no delay slots

*Machines with software-visible delay slots may execute significant number of NOP instructions inserted by the compiler.  
NOPs reduce CPI, but increase instructions/program!*

# Traps and Interrupts

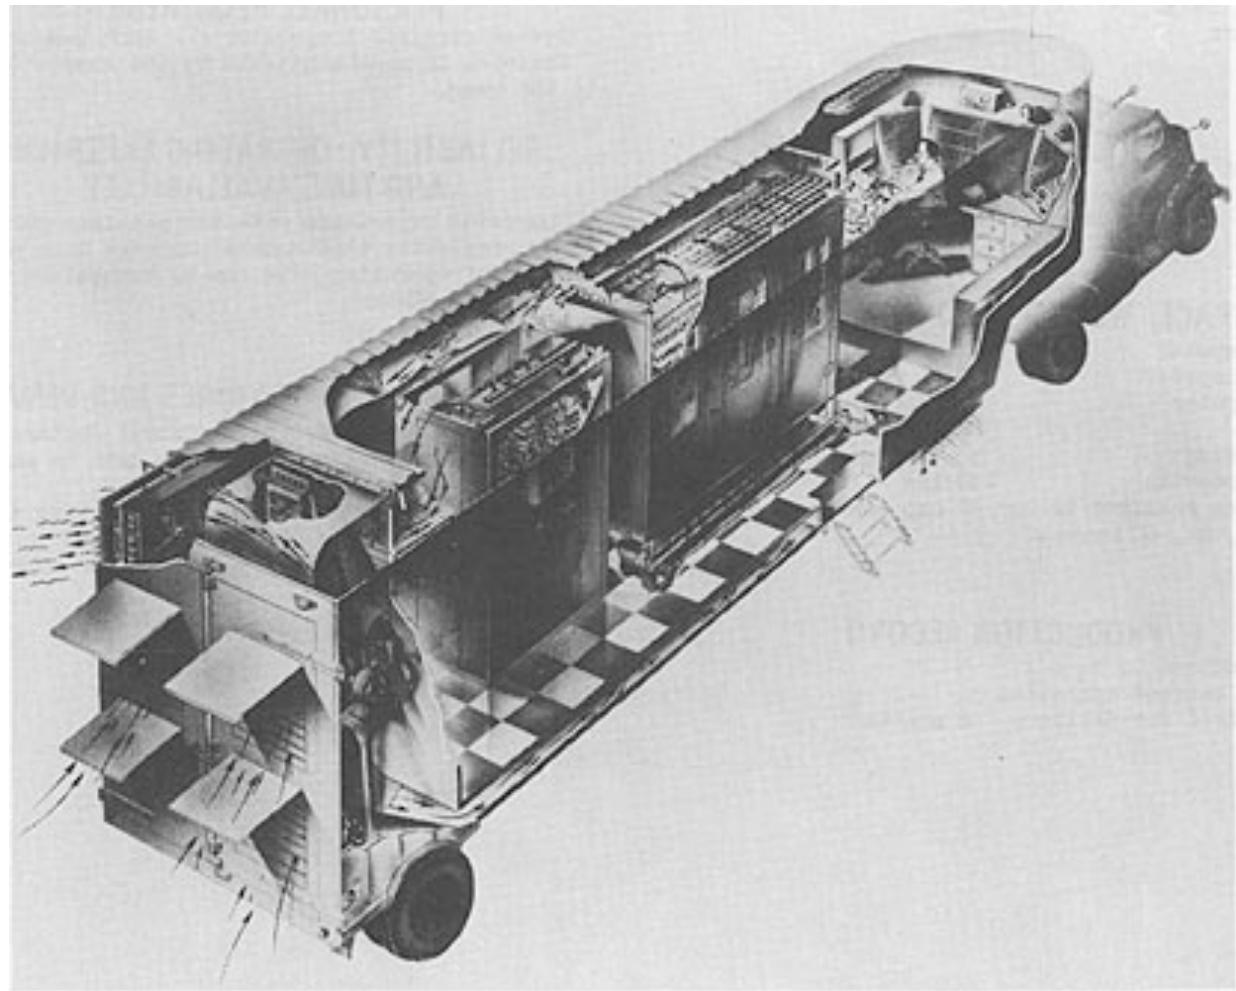
In class, we'll use following terminology

- ***Exception***: An unusual internal event caused by program during execution
  - E.g., page fault, arithmetic underflow
- ***Interrupt***: An external event outside of running program
- ***Trap***: Forced transfer of control to supervisor caused by exception or interrupt
  - Not all exceptions cause traps (c.f. IEEE 754 floating-point standard)

# History of Exception Handling

- Analytical Engine had overflow exceptions
- First system with traps was Univac-I, 1951
  - Arithmetic overflow would either
    - 1. trigger the execution a two-instruction fix-up routine at address 0, or
    - 2. at the programmer's option, cause the computer to stop
  - Later Univac 1103, 1955, modified to add external interrupts
    - Used to gather real-time wind tunnel data
- First system with I/O interrupts was DYSEAC, 1954
  - Had two program counters, and I/O signal caused switch between two PCs
  - Also, first system with DMA (**Direct Memory Access by I/O device**)
  - And, first mobile computer!

# DYSEAC, first mobile computer!



- Carried in two tractor trailers, 12 tons + 8 tons
- Built for US Army Signal Corps

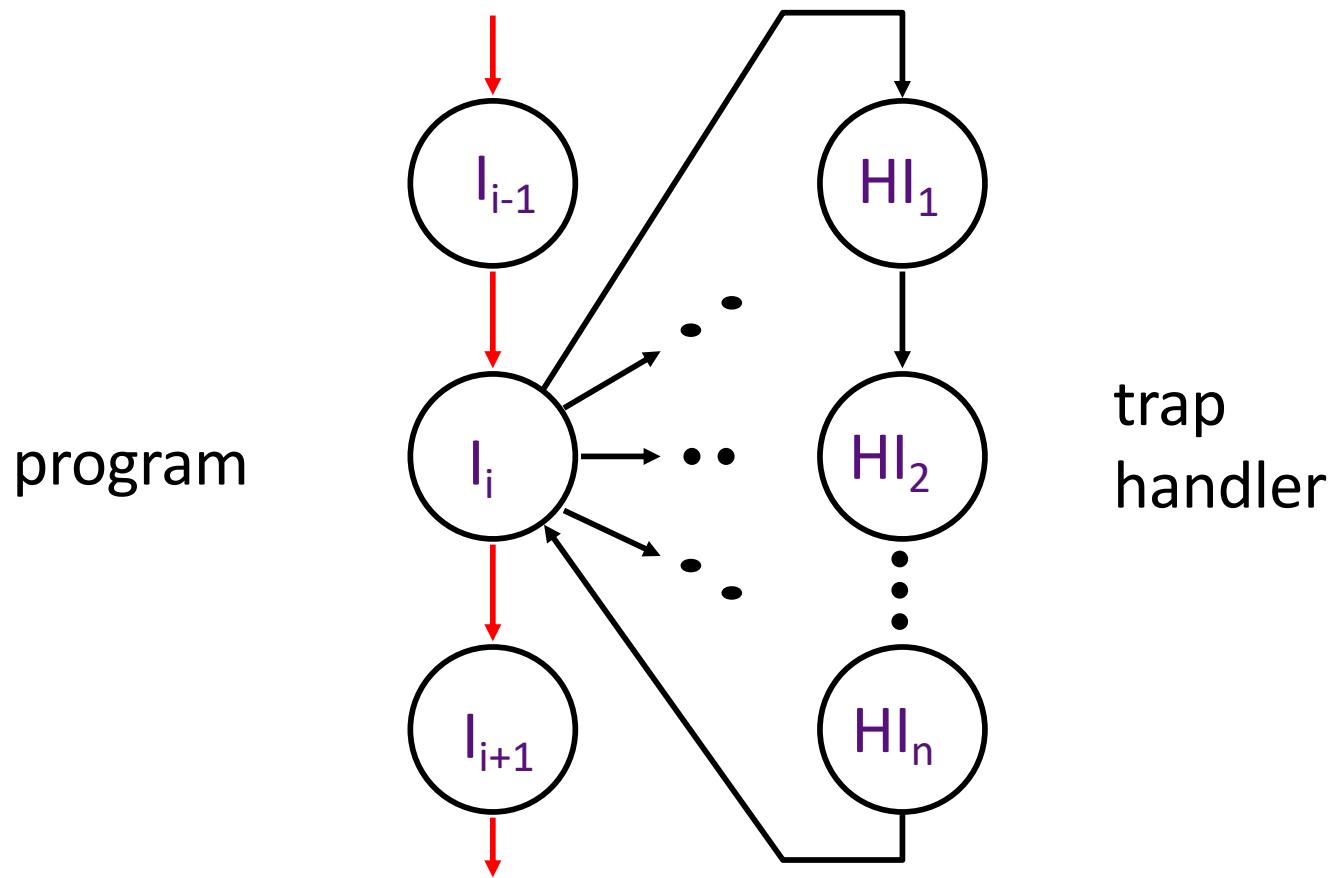
*[Courtesy Mark Smotherman]*

# Asynchronous Interrupts

- An I/O device requests attention by asserting one of the *prioritized interrupt request lines*
- When the processor decides to process the interrupt
  - It stops the current program at instruction  $I_i$ , completing all the instructions up to  $I_{i-1}$  (*precise interrupt*)
  - It saves the PC of instruction  $I_i$  in a special register (EPC)
  - It disables interrupts and transfers control to a designated interrupt handler running in supervisor mode

# Trap:

## altering the normal flow of control



An *external or internal event* that needs to be processed by another (system) program. The event is usually unexpected or rare from program's point of view.

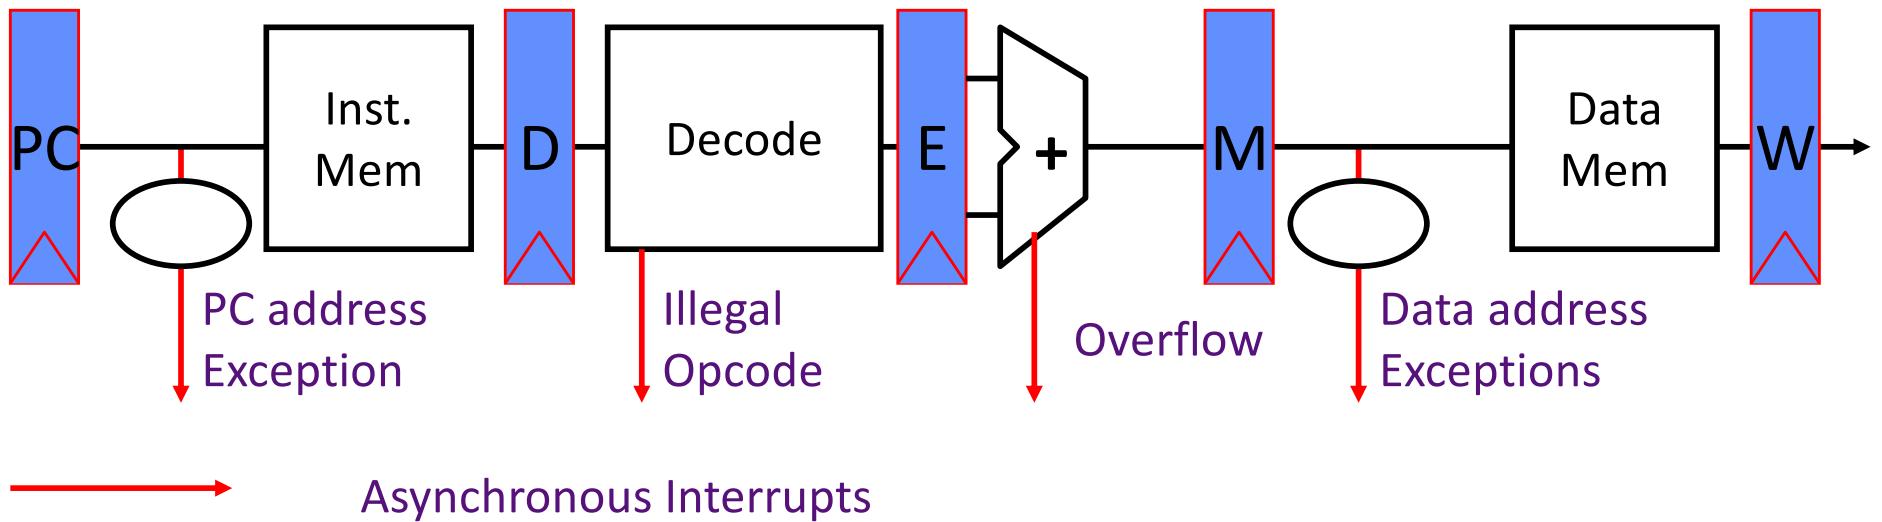
# Trap Handler

- Saves *EPC* before enabling interrupts to allow nested interrupts ⇒
  - need an instruction to move EPC into GPRs
  - need a way to mask further interrupts at least until EPC can be saved
- Needs to read a *status register* that indicates the *cause* of the trap
- Uses a special indirect jump instruction ERET (*return-from-environment*) which
  - enables interrupts
  - restores the processor to the user mode
  - restores hardware status and control state

# Synchronous Trap

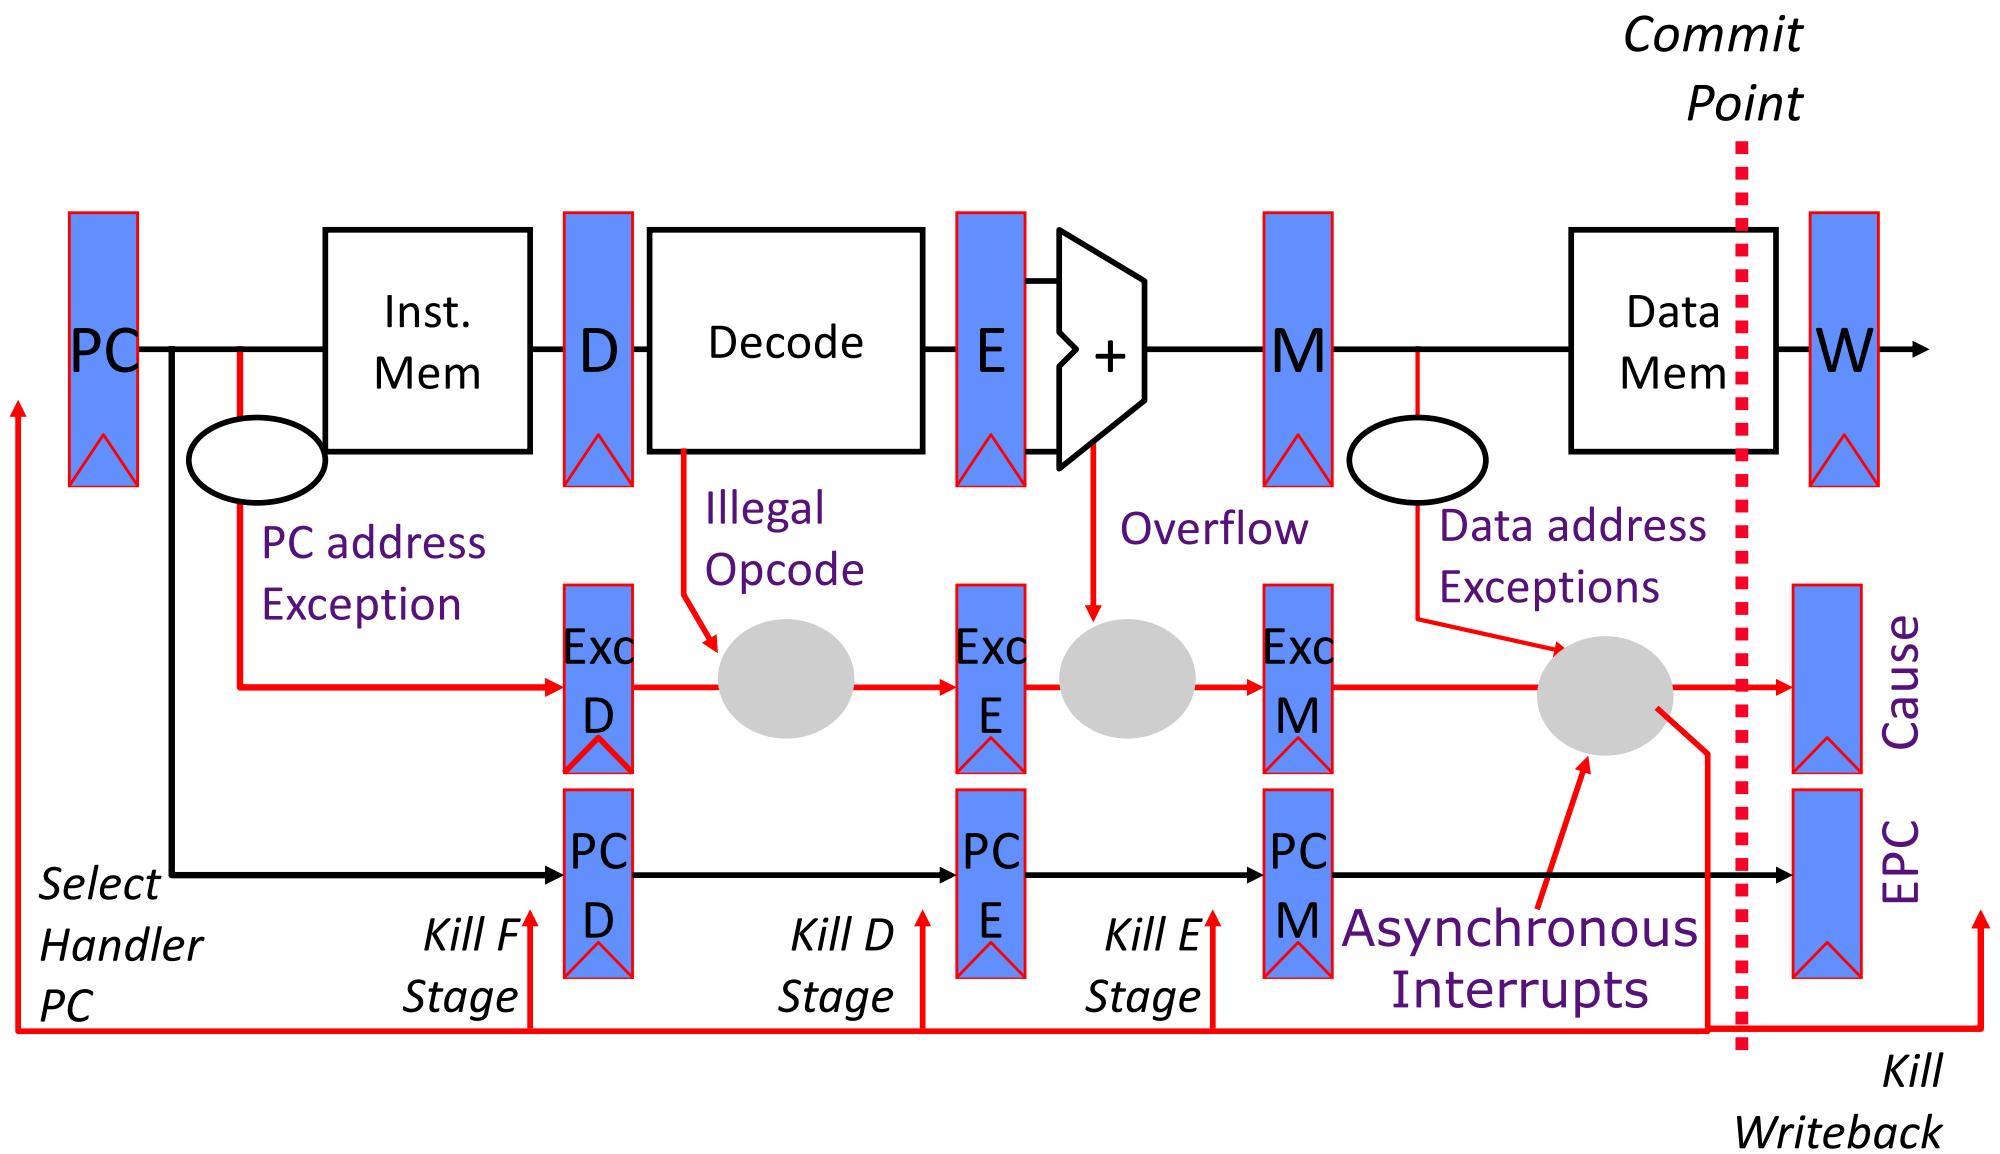
- A synchronous trap is caused by an exception on a *particular instruction*
- In general, the instruction cannot be completed and needs to be *restarted* after the exception has been handled
  - requires undoing the effect of one or more partially executed instructions
- In the case of a system call trap, the instruction is considered to have been completed
  - a special jump instruction involving a change to a privileged mode

# Exception Handling 5-Stage Pipeline



- How to handle multiple simultaneous exceptions in different pipeline stages?
- How and where to handle external asynchronous interrupts?

# Exception Handling 5-Stage Pipeline



## Exception Handling 5-Stage Pipeline

- Hold exception flags in pipeline until commit point (M stage)
- Exceptions in earlier pipe stages override later exceptions *for a given instruction*
- Inject external interrupts at commit point (override others)
- If trap at commit: update Cause and EPC registers, kill all stages, inject handler PC into fetch stage

# Speculating on Exceptions

- Prediction mechanism
  - Exceptions are rare, so simply predicting no exceptions is very accurate!
- Check prediction mechanism
  - Exceptions detected at end of instruction execution pipeline, special hardware for various exception types
- Recovery mechanism
  - Only write architectural state at commit point, so can throw away partially executed instructions after exception
  - Launch exception handler after flushing pipeline
- Bypassing allows use of uncommitted instruction results by following instructions

# Acknowledgements

- This course is partly inspired by previous MIT 6.823 and Berkeley CS252 computer architecture courses created by my collaborators and colleagues:
  - Arvind (MIT)
  - Joel Emer (Intel/MIT)
  - James Hoe (CMU)
  - John Kubiatowicz (UCB)
  - David Patterson (UCB)