

# **CS 152 Computer Architecture and Engineering**

## **CS252 Graduate Computer Architecture**

### **Lecture 18 Cache Coherence**

Krste Asanovic

Electrical Engineering and Computer Sciences  
University of California at Berkeley

**`http://www.eecs.berkeley.edu/~krste`**  
**`http://inst.eecs.berkeley.edu/~cs152`**

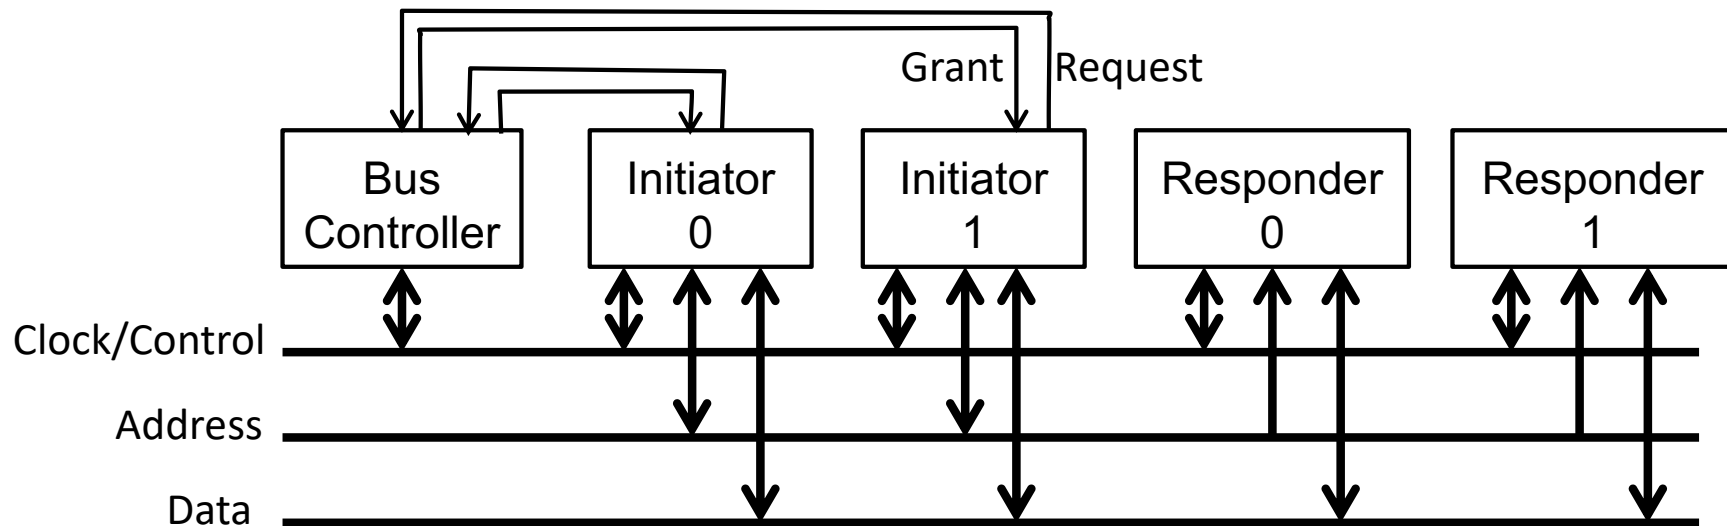
## Last Time in Lecture 17

- RISC-V Standard Vectors

# Adopting Alternative Terminology

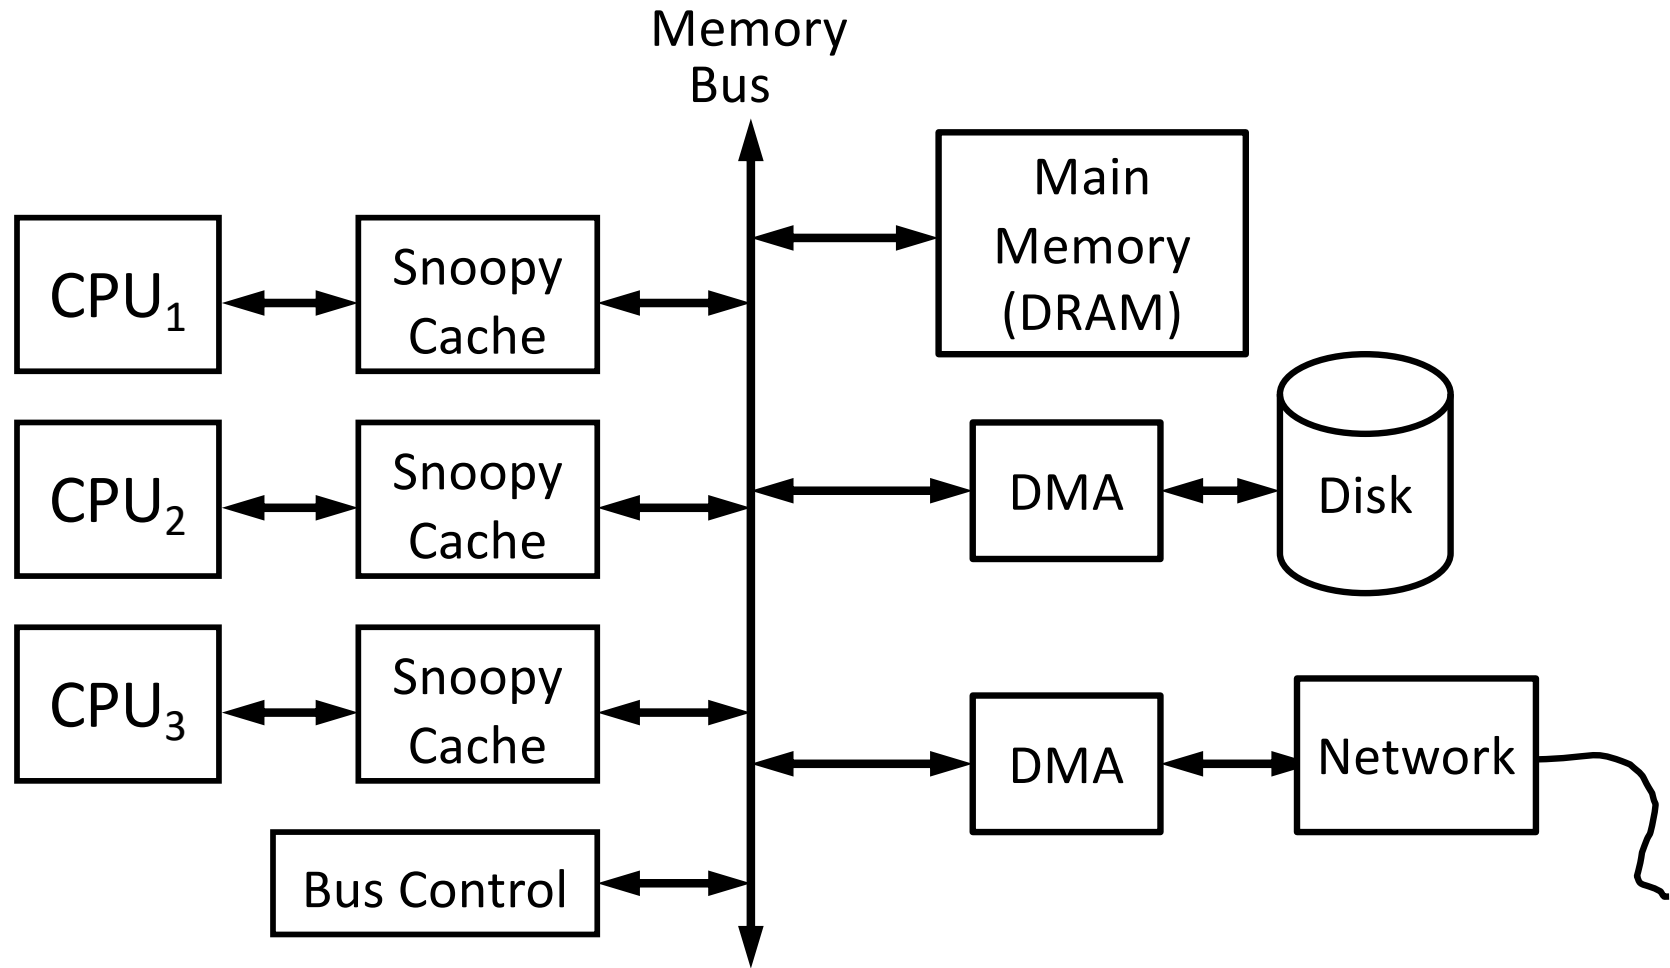
- Traditional terms “master” and “slave” replaced with “initiator” and “responder”
- Most extant documentation will refer to “bus masters” and “bus slaves”

# Bus Management



- A “bus” is a collection of shared wires
  - Newer “busses” use point-point links
- At any instant, only one “initiator” can initiate a transaction by driving wires
  - Initiators arbitrate for access with requests to bus “controller”
  - Some busses only allow one initiator (in which case, it’s also the controller)
- Multiple “responders” can observe and conditionally respond to the transaction on the wires
  - responders decode address on bus to see if they should respond (memory is most common responder)
  - some initiators can also act as responders

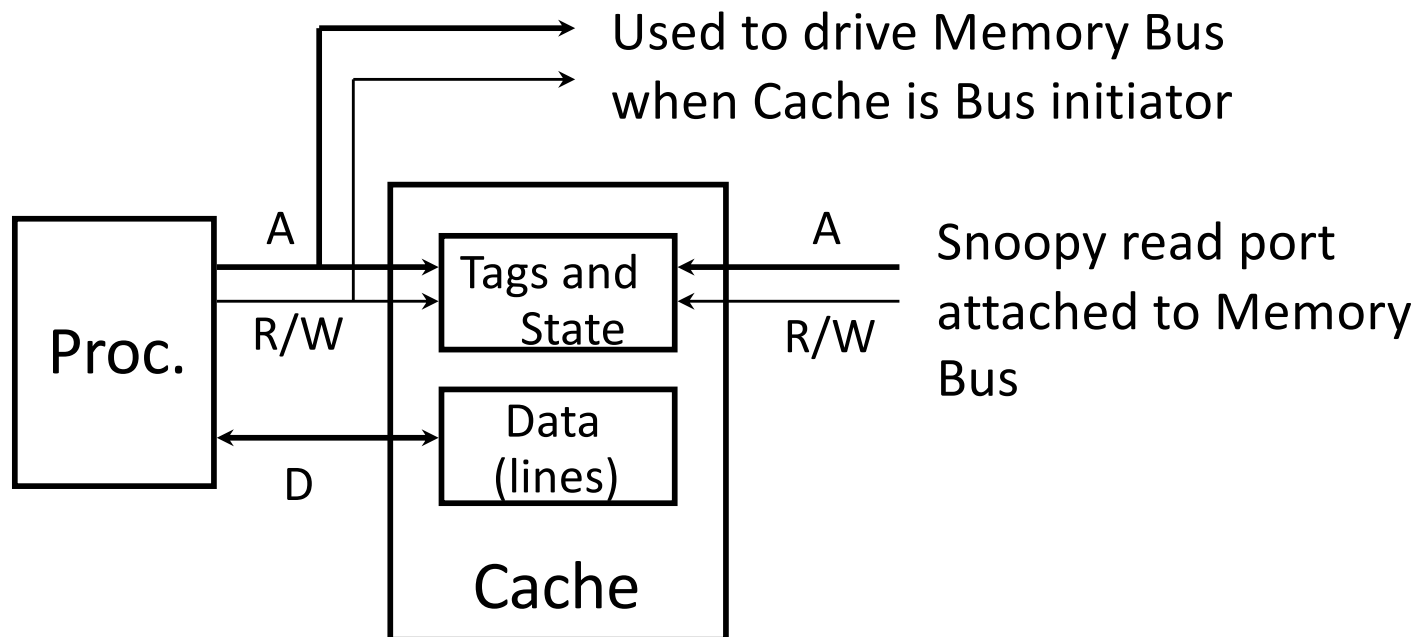
# Shared-Memory Multiprocessor



Use snoop mechanism to keep all processors' view of memory coherent

# Snoopy Cache, *Goodman 1983*

- Idea: Have cache watch (or snoop upon) other memory transactions, and then “do the right thing”
- Snoopy cache tags are dual-ported



# Snoopy Cache-Coherence Protocols

- Write miss:
  - the address is invalidated in all other caches before the write is performed
- Read miss:
  - if a dirty copy is found in some cache, a write-back is performed before the memory is read

# Cache State-Transition Diagram

*The MSI protocol*

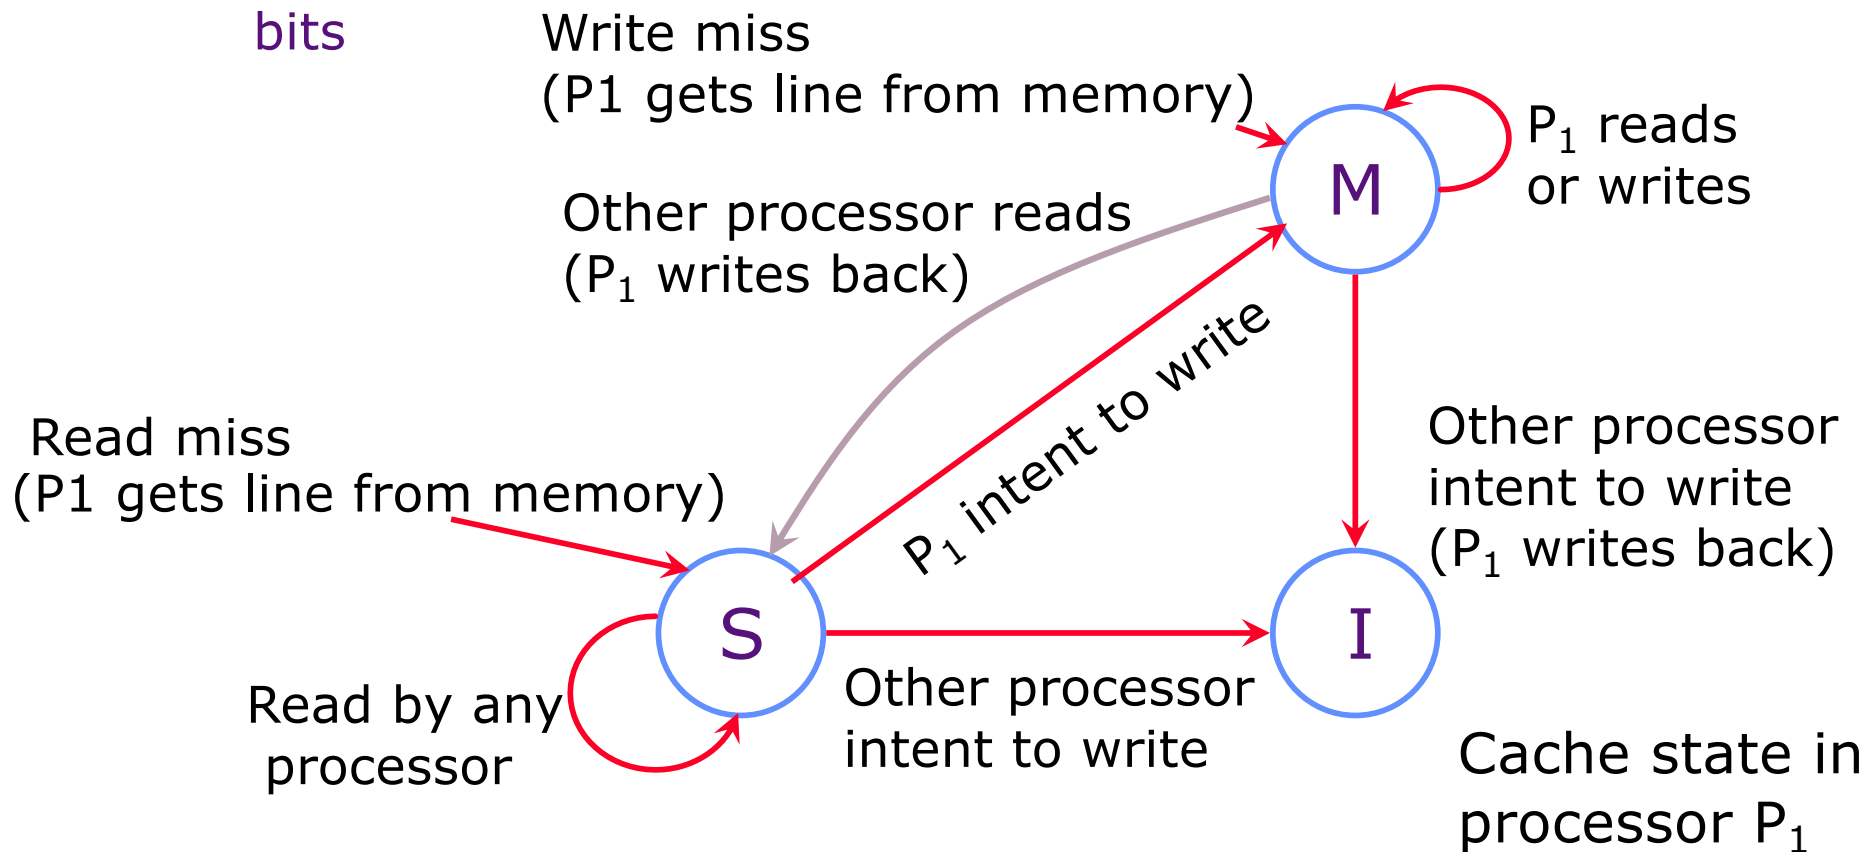
Each cache line has state bits



M: Modified

S: Shared

I: Invalid

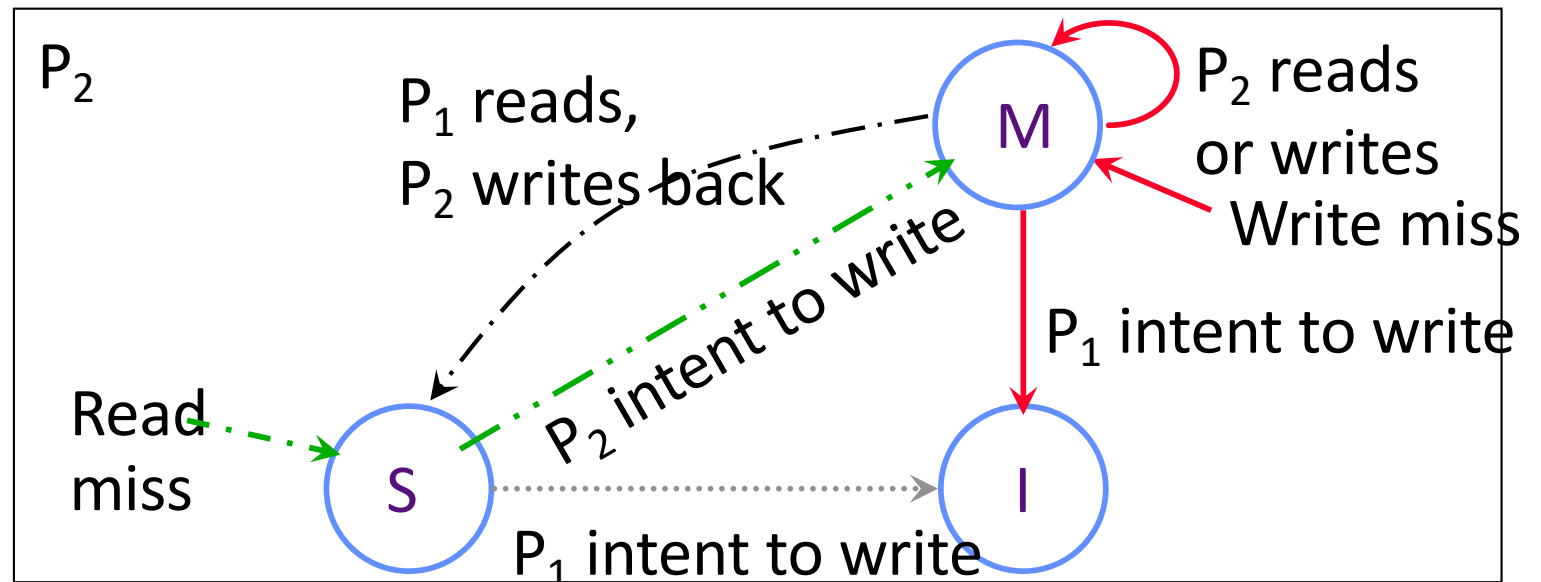
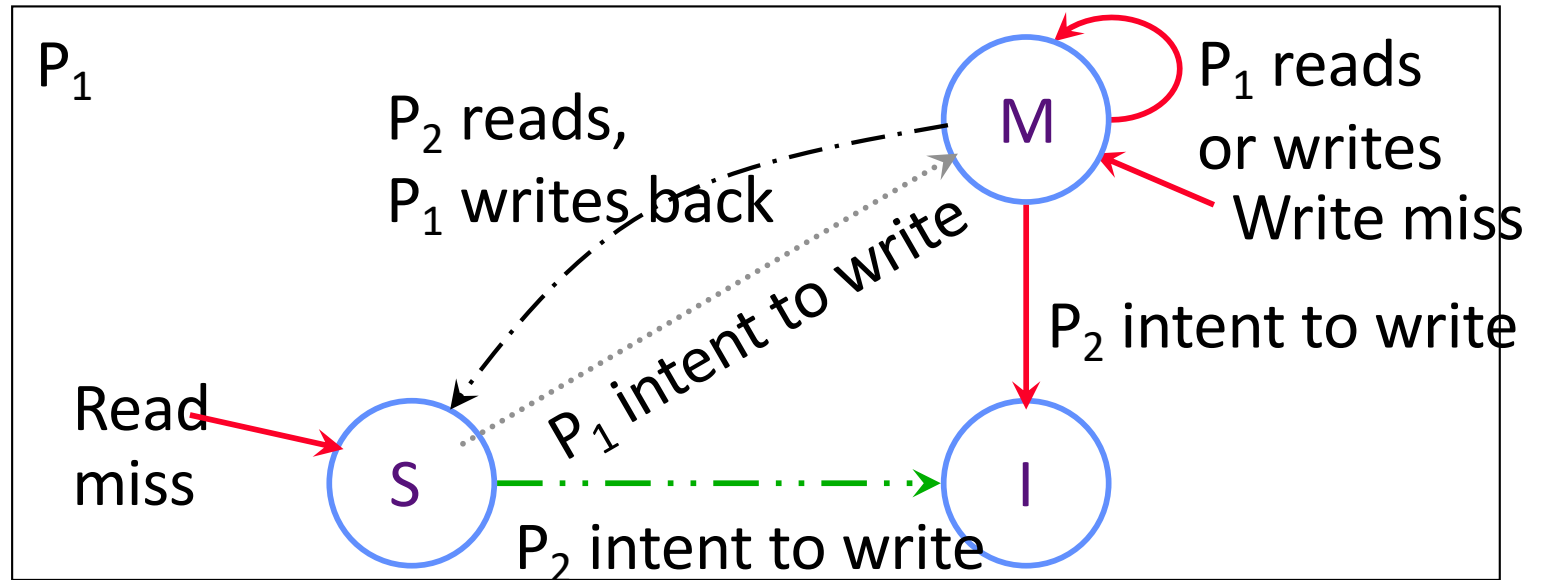




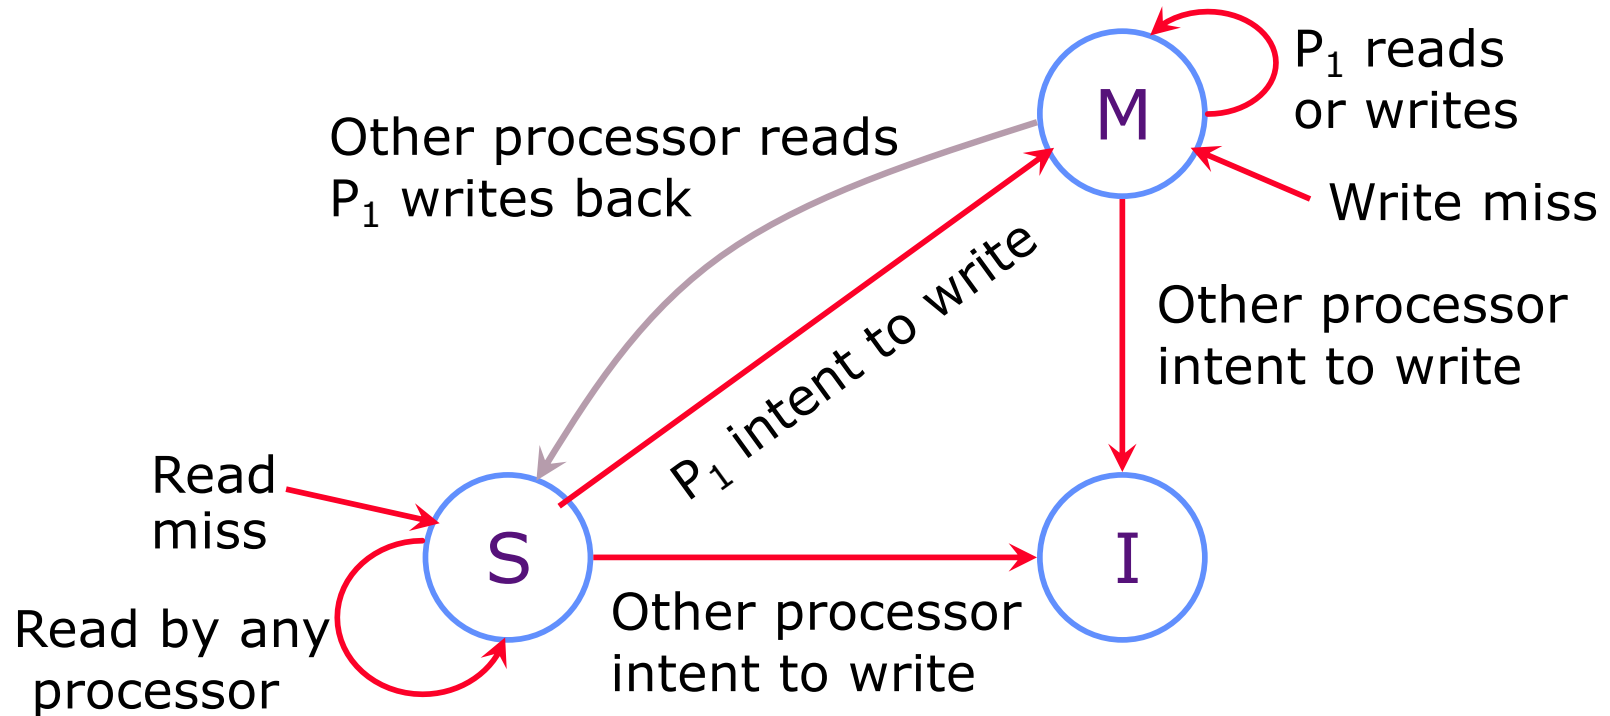
# Two-Processor Example

## (Reading and writing the same cache line)

$P_1$  reads  
 $P_1$  writes  
 $P_2$  reads  
 $P_2$  writes  
 $P_1$  reads  
 $P_1$  writes  
 $P_2$  writes  
 $P_1$  writes



# Observation

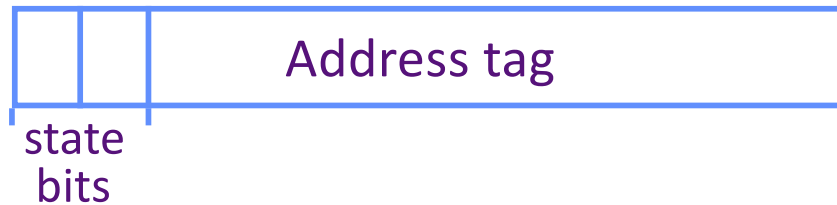


- If a line is in the **M** state then no other cache can have a copy of the line!
- Memory stays coherent, multiple differing copies cannot exist

# MESI: An Enhanced MSI protocol

increased performance for private data

Each cache line has a tag

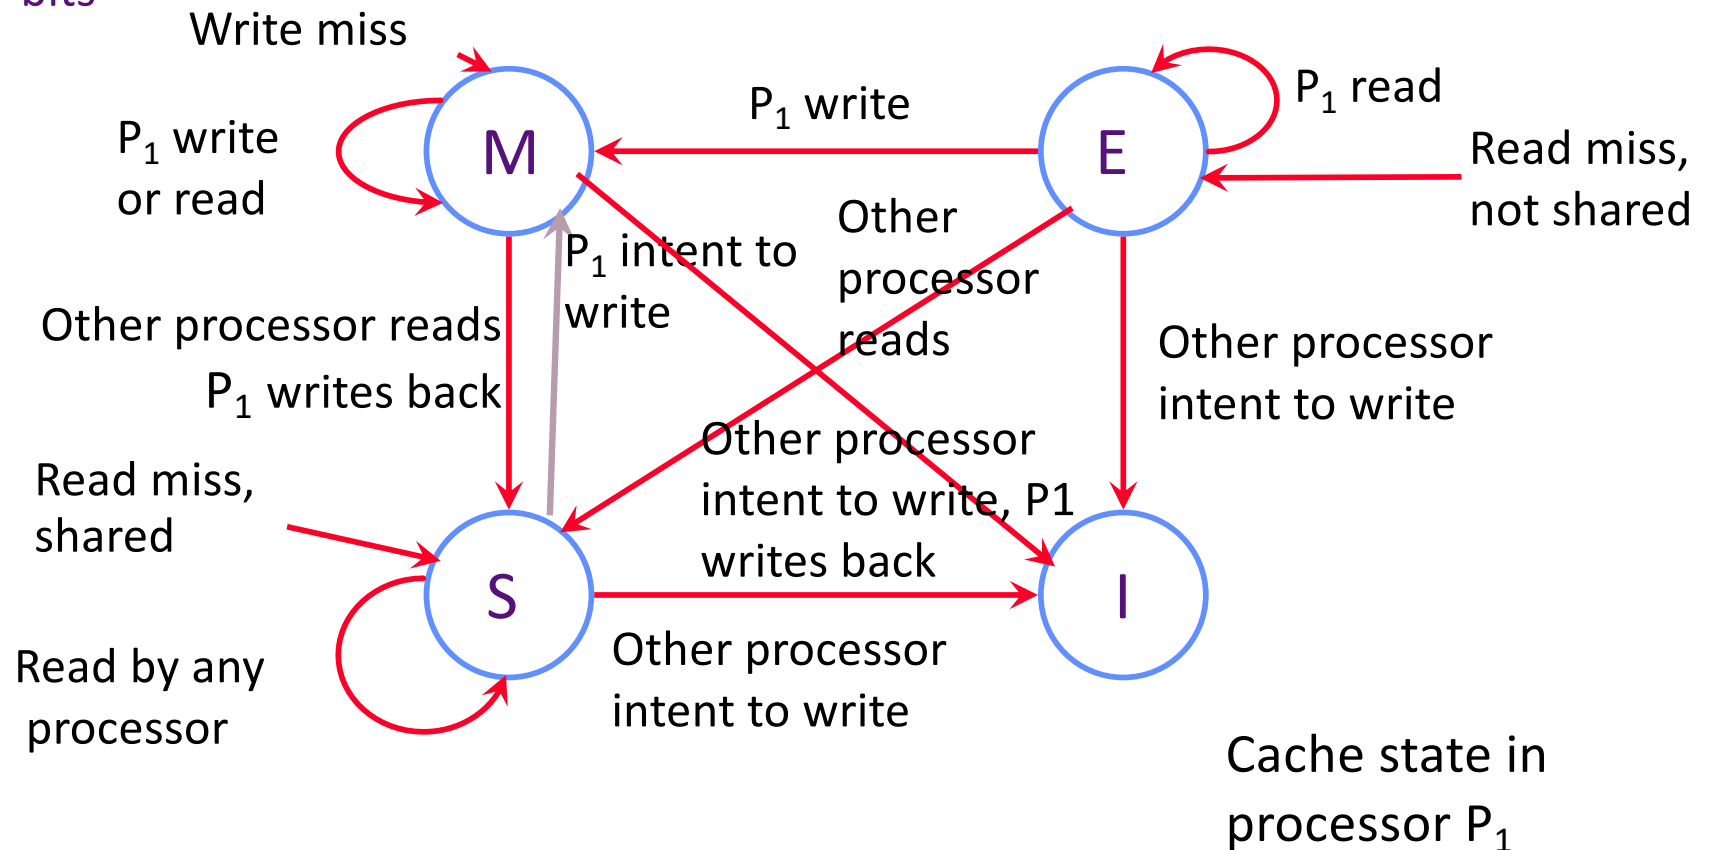


M: Modified Exclusive


E: Exclusive but unmodified

S: Shared

I: Invalid

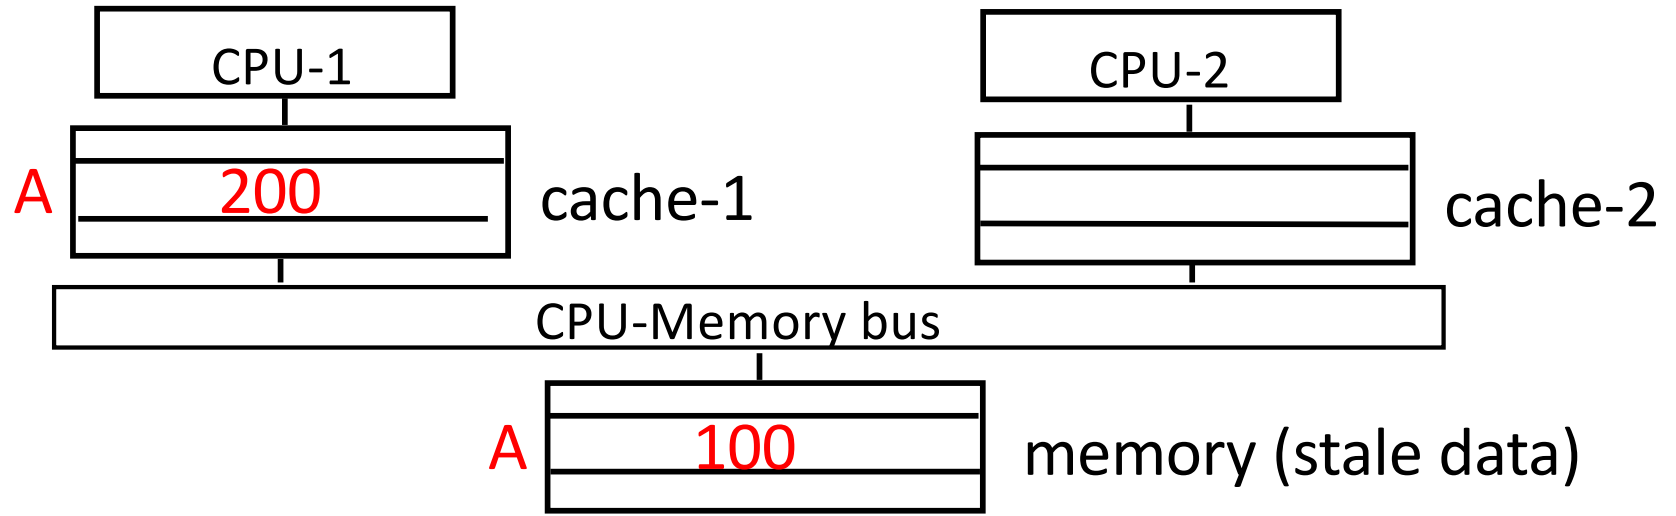


# Optimized Snoop with Level-2 Caches



- Processors often have two-level caches
  - small L1, large L2 (usually both on chip now)
- Inclusion property: entries in L1 must be in L2
  - Miss in L2  $\Rightarrow$  Not present in L1
  - Only if invalidation hits in L2  $\Rightarrow$  probe and invalidate in L1
- Snooping on L2 does not affect CPU-L1 bandwidth

# Intervention



When a read-miss for **A** occurs in cache-2,  
Cache-2 initiates a read request for **A** on the bus

- Cache-1 needs to supply & change its state to shared
- The memory may **respond** to the request also!

*Does memory know it has stale data?*

Cache-1 needs to *intervene* through memory controller to supply correct data to cache-2

# False Sharing

state	line addr	data0	data1	...	dataN
-------	-----------	-------	-------	-----	-------

A cache line contains more than one word

Cache-coherence is done at the line-level and not word-level

Suppose  $M_1$  writes  $\text{word}_i$  and  $M_2$  writes  $\text{word}_k$  and  $i \neq k$  but both words have the same line address.

*What can happen?*

# Performance of Symmetric Multiprocessors (SMPs)

Cache performance is combination of:

- Uniprocessor cache miss traffic
- Traffic caused by communication
  - Results in invalidations and subsequent cache misses
- Coherence misses
  - Sometimes called a Communication miss
  - 4th C of cache misses along with Compulsory, Capacity, & Conflict

# Coherency Misses

- True sharing misses arise from the communication of data through the cache coherence mechanism
  - Invalidates due to 1st write to shared line
  - Reads by another CPU of modified line in different cache
  - Miss would still occur if line size were 1 word
- False sharing misses when a line is invalidated because some word in the line, other than the one being read, is written into
  - Invalidation does not cause a new value to be communicated, but only causes an extra cache miss
  - Line is shared, but no word in line is actually shared  
⇒ miss would not occur if line size were 1 word




## Example: True v. False Sharing v. Hit?

- MSI protocol
- Assume x1 and x2 in same cache line.  
P1 and P2 both read x1 and x2 before.

Time	P1	P2	True, False, Hit? Why?
1	Write x1		True miss; invalidate x1 in P2
2		Read x2	False miss; x1 irrelevant to P2
3	Write x1		False miss; x1 irrelevant to P2
4		Write x2	True miss; x2 not writeable
5	Read x2		True miss; x2 invalid in P1


# MP Performance 4-Processor Commercial Workload: OLTP, Decision Support (Database), Search Engine

- Uniprocessor cache misses improve with cache size increase (Instruction, Capacity/Conflict, Compulsory)
- True sharing and false sharing unchanged going from 1 MiB to 8 MiB (L3 cache)



# MP Performance 2MiB Cache Commercial Workload: OLTP, Decision Support (Database), Search Engine

- True sharing, false sharing increase going from 1 to 8 CPUs



# CS152 Administrivia

- PS4 and Lab 3 due today
- Lab 4 due Monday April 19
- Midterm 2 Wednesday April 14
  - covers lectures 10-17, plus associated problem sets, labs, and readings
  - updated Zoom proctoring procedure

# CS252 Administrivia

- This week's readings Cray-1 and VLIW machines
- Thursday April 15<sup>th</sup> Project Checkpoint
  - Schedule 10-minute individual group zoom calls during discussion period

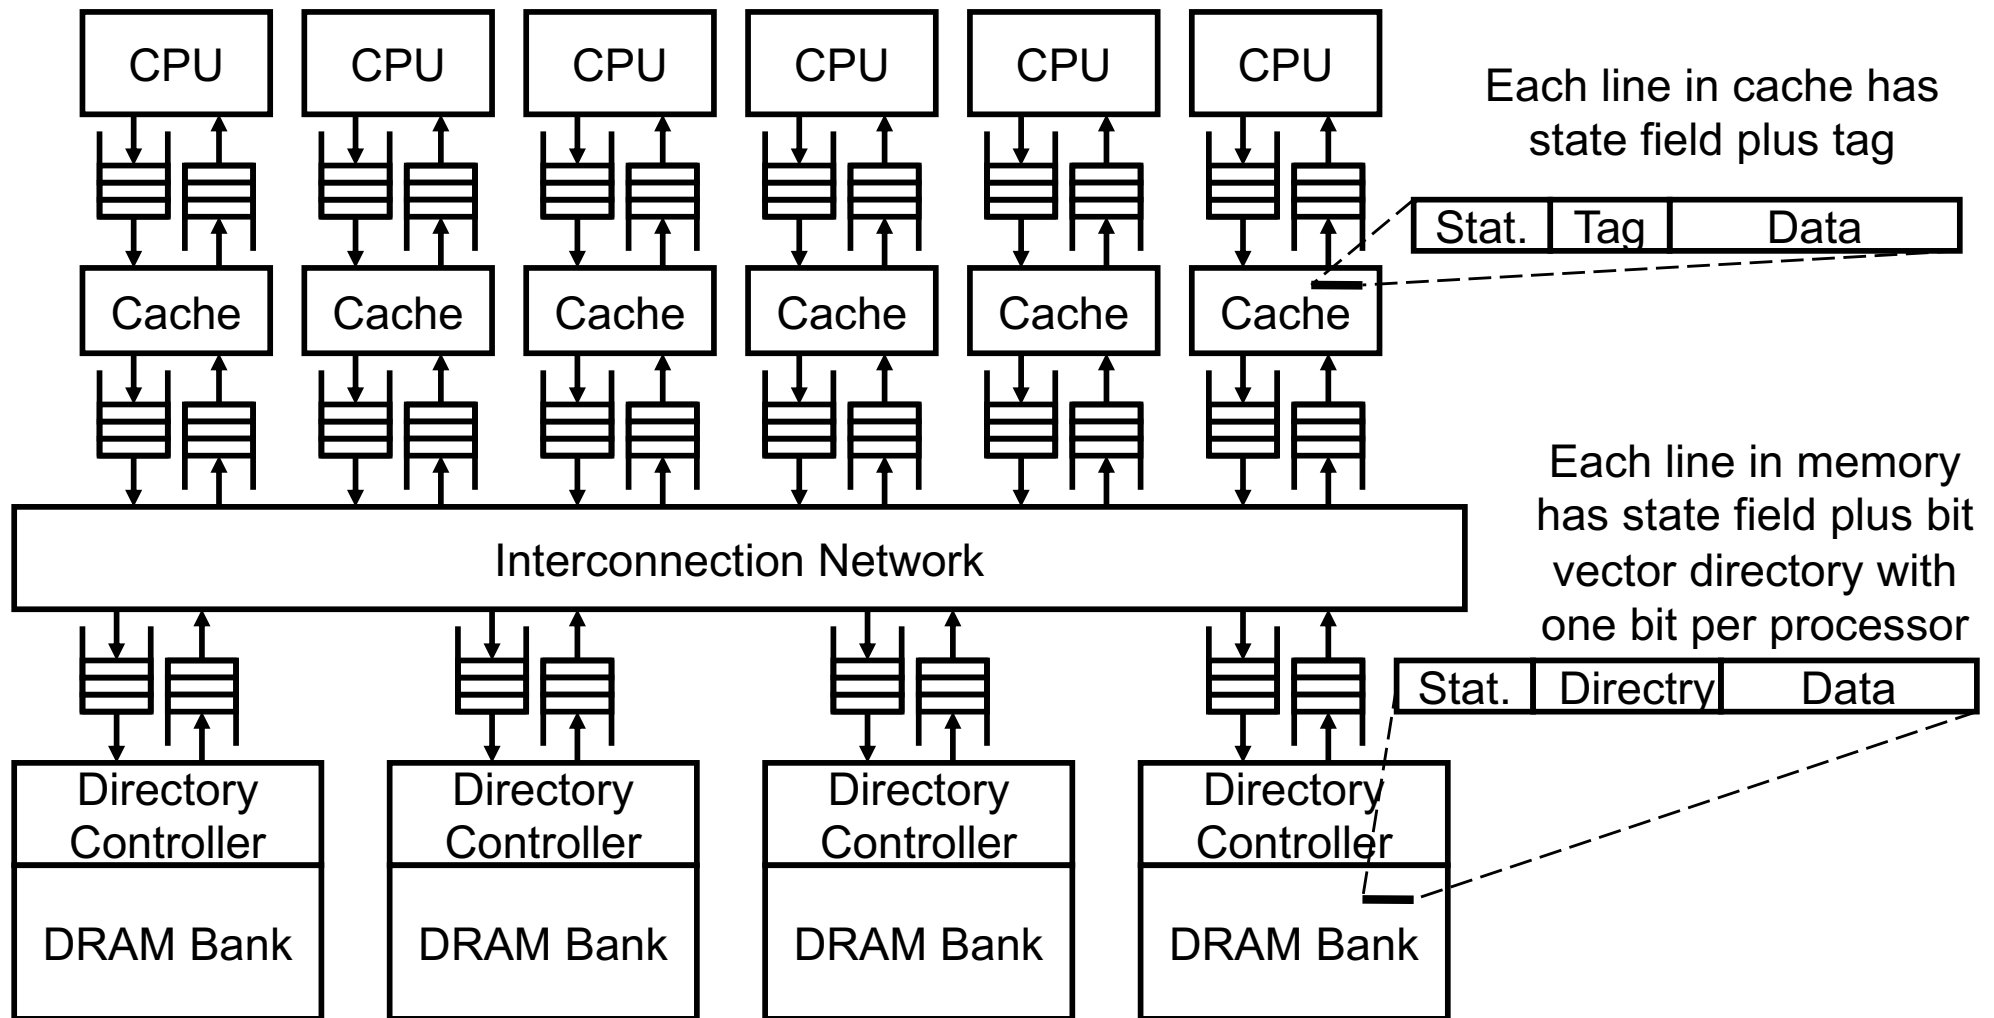
# Scaling Snoopy/Broadcast Coherence

- When any processor gets a miss, must probe every other cache
- Scaling up to more processors limited by:
  - Communication bandwidth over bus
  - Snoop bandwidth into tags
- Can improve bandwidth by using multiple interleaved buses with interleaved tag banks
  - E.g, two bits of address pick which of four buses and four tag banks to use
    - (e.g., bits 7:6 of address pick bus/tag bank, bits 5:0 pick byte in 64-byte line)
- Buses don't scale to large number of connections, so can use point-to-point network for larger number of nodes, but then limited by tag bandwidth when broadcasting snoop requests.
- **Insight:** Most snoops fail to find a match!

# Scalable Approach: Directories

- Every memory line has associated directory information
  - keeps track of copies of cached lines and their states
  - on a miss, find directory entry, look it up, and communicate only with the nodes that have copies if necessary
  - in scalable networks, communication with directory and copies is through network transactions
- Many alternatives for organizing directory information

# Directory Cache Protocol



- Assumptions: Reliable network, FIFO message delivery between any given source-destination pair



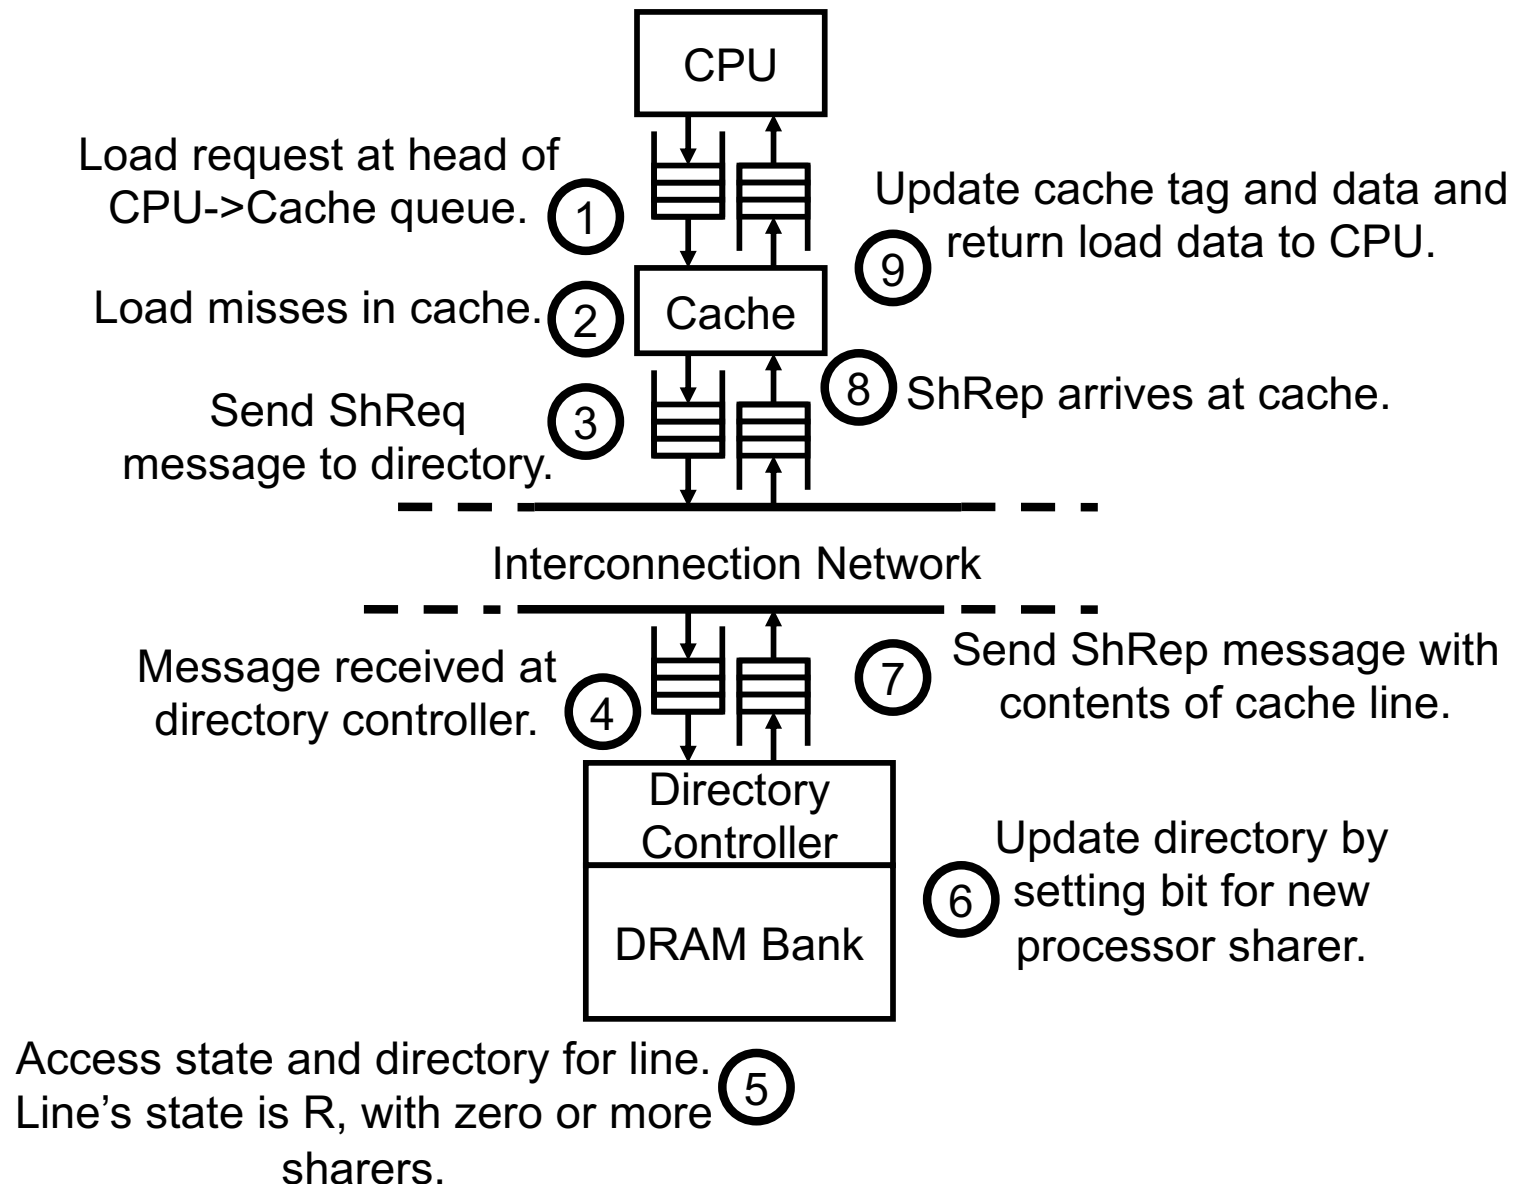
# Cache States

- For each cache line, there are 4 possible states:
  - **C-invalid** (= Nothing): The accessed data is not resident in the cache.
  - **C-shared** (= Sh): The accessed data is resident in the cache, and possibly also cached at other sites. The data in memory is valid.
  - **C-modified** (= Ex): The accessed data is exclusively resident in this cache, and has been modified. Memory does not have the most up-to-date data.
  - **C-transient** (= Pending): The accessed data is in a transient state (for example, the site has just issued a protocol request, but has not received the corresponding protocol reply).

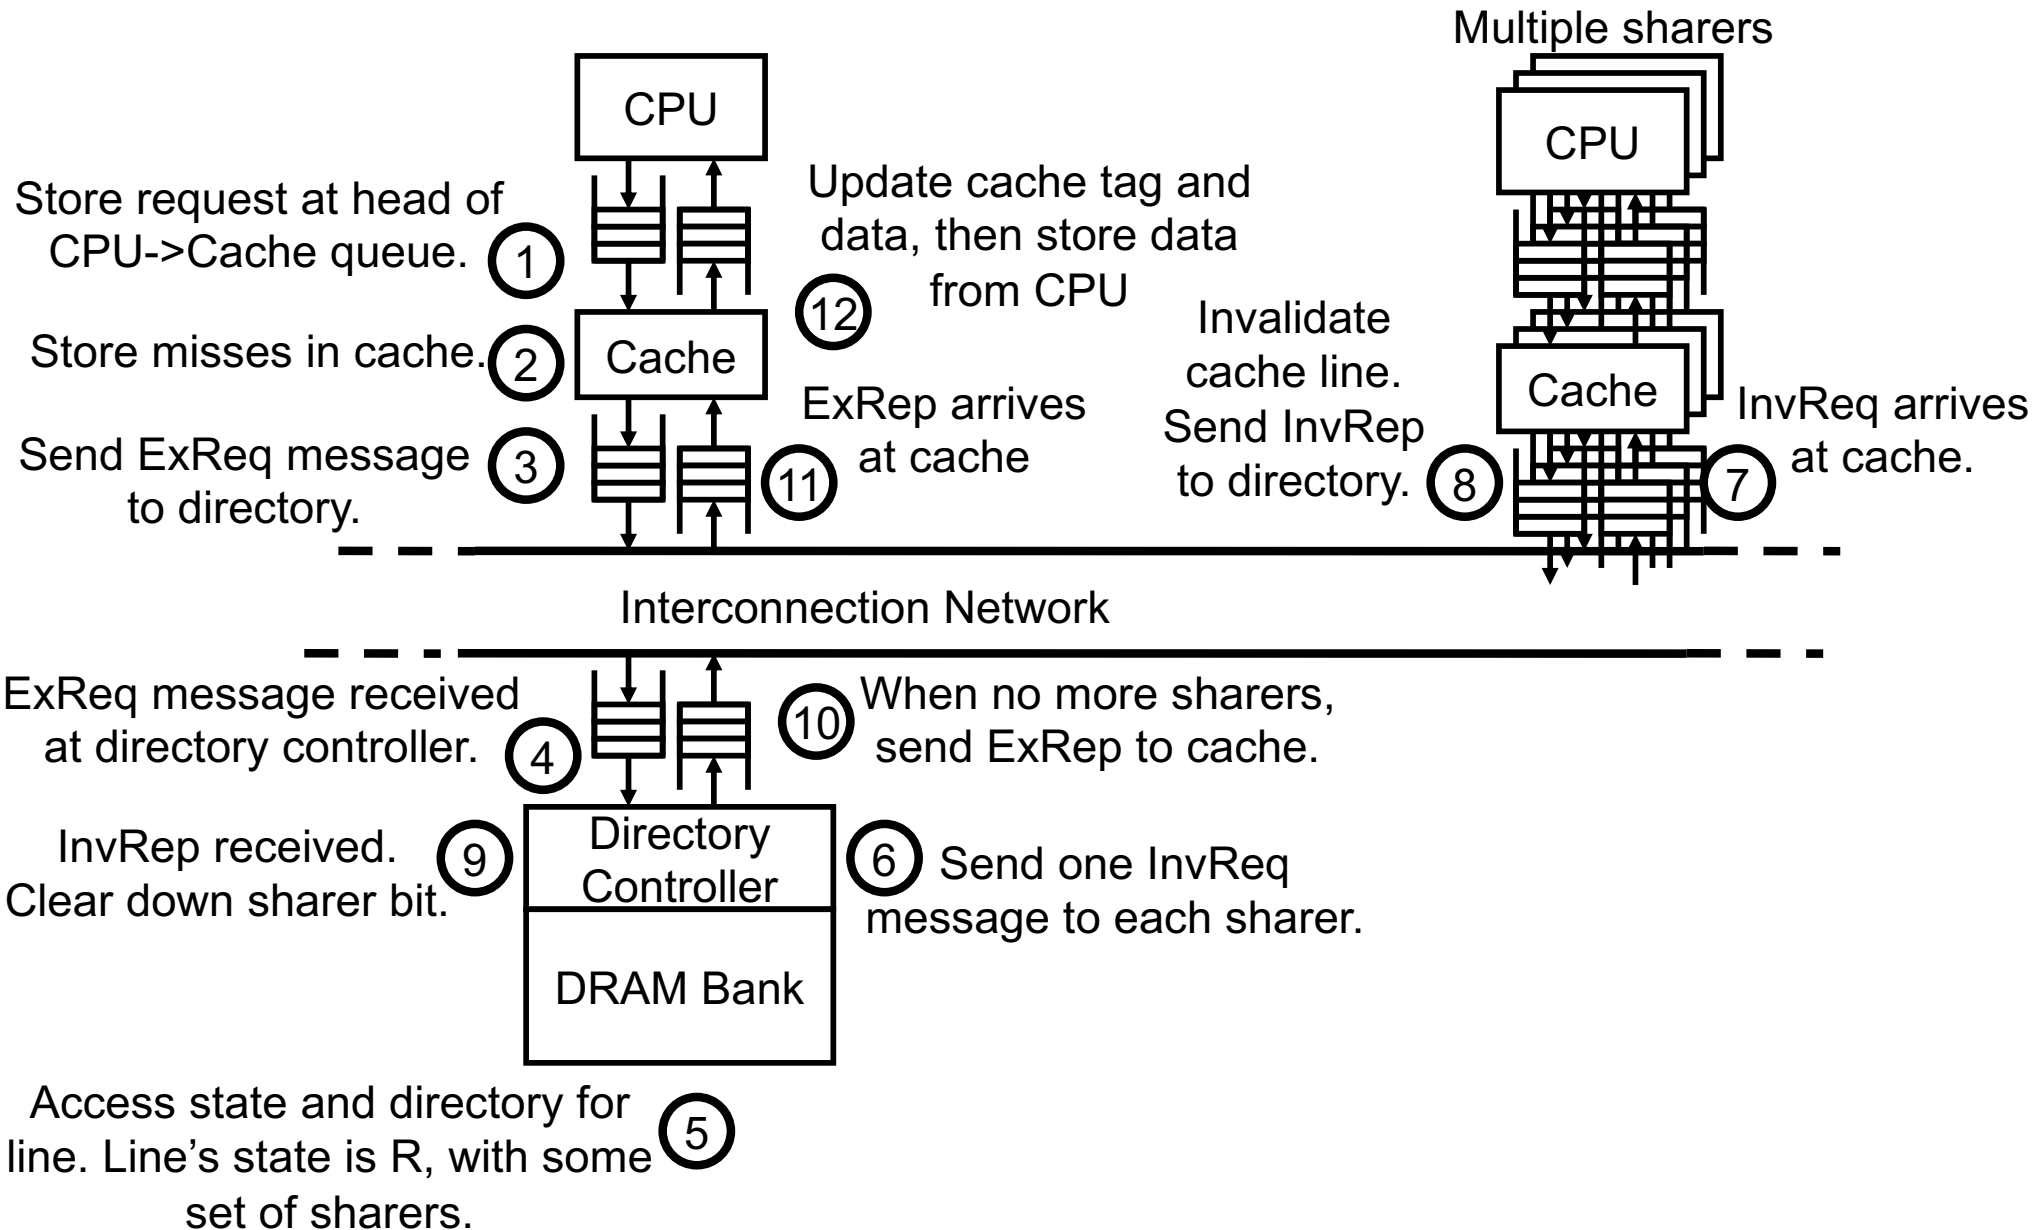
## Home directory states

- For each memory line, there are 4 possible states:
  - **R(dir)**: The memory line is shared by the sites specified in dir (dir is a set of sites). The data in memory is valid in this state. If dir is empty (i.e.,  $\text{dir} = \varepsilon$ ), the memory line is not cached by any site.
  - **W(id)**: The memory line is exclusively cached at site id, and has been modified at that site. Memory does not have the most up-to-date data.
  - **TR(dir)**: The memory line is in a transient state waiting for the acknowledgements to the invalidation requests that the home site has issued.
  - **TW(id)**: The memory line is in a transient state waiting for a line exclusively cached at site id (i.e., in C-modified state) to make the memory line at the home site up-to-date.

# Read miss, to uncached or shared line



# Write miss, to read shared line



# Concurrency Management

- Protocol would be easy to design if only one transaction in flight across entire system
- But, want greater throughput and don't want to have to coordinate across entire system
- Great complexity in managing multiple outstanding concurrent transactions to cache lines
  - Can have multiple requests in flight to same cache line!

# Acknowledgements

- This course is partly inspired by previous MIT 6.823 and Berkeley CS252 computer architecture courses created by my collaborators and colleagues:
  - Arvind (MIT)
  - Joel Emer (Intel/MIT)
  - James Hoe (CMU)
  - John Kubiatowicz (UCB)
  - David Patterson (UCB)