



CS 152/252A Computer Architecture and Engineering



Sophia Shao

Lecture 17: Vectors

Nvidia in Talks to Use Intel as a Foundry to Manufacture Chips

As inconceivable as it may sound, your next Nvidia GPU could be manufactured by Intel. Nvidia CEO Jensen Huang held a question and answer session with the press today, and the topic quickly turned to Intel's Foundry Services (IFS) initiative that will see Intel making chips for other companies as part of its IDM 2.0 initiative. Surprisingly, Huang confirmed that his company is considering using Intel's foundry to possibly make some of its chips. Intel is now a direct competitor with Nvidia on both the CPU and GPU fronts, but Huang also explained that Intel and AMD have known Nvidia's secret roadmaps for years, so he isn't paranoid about sharing more information.

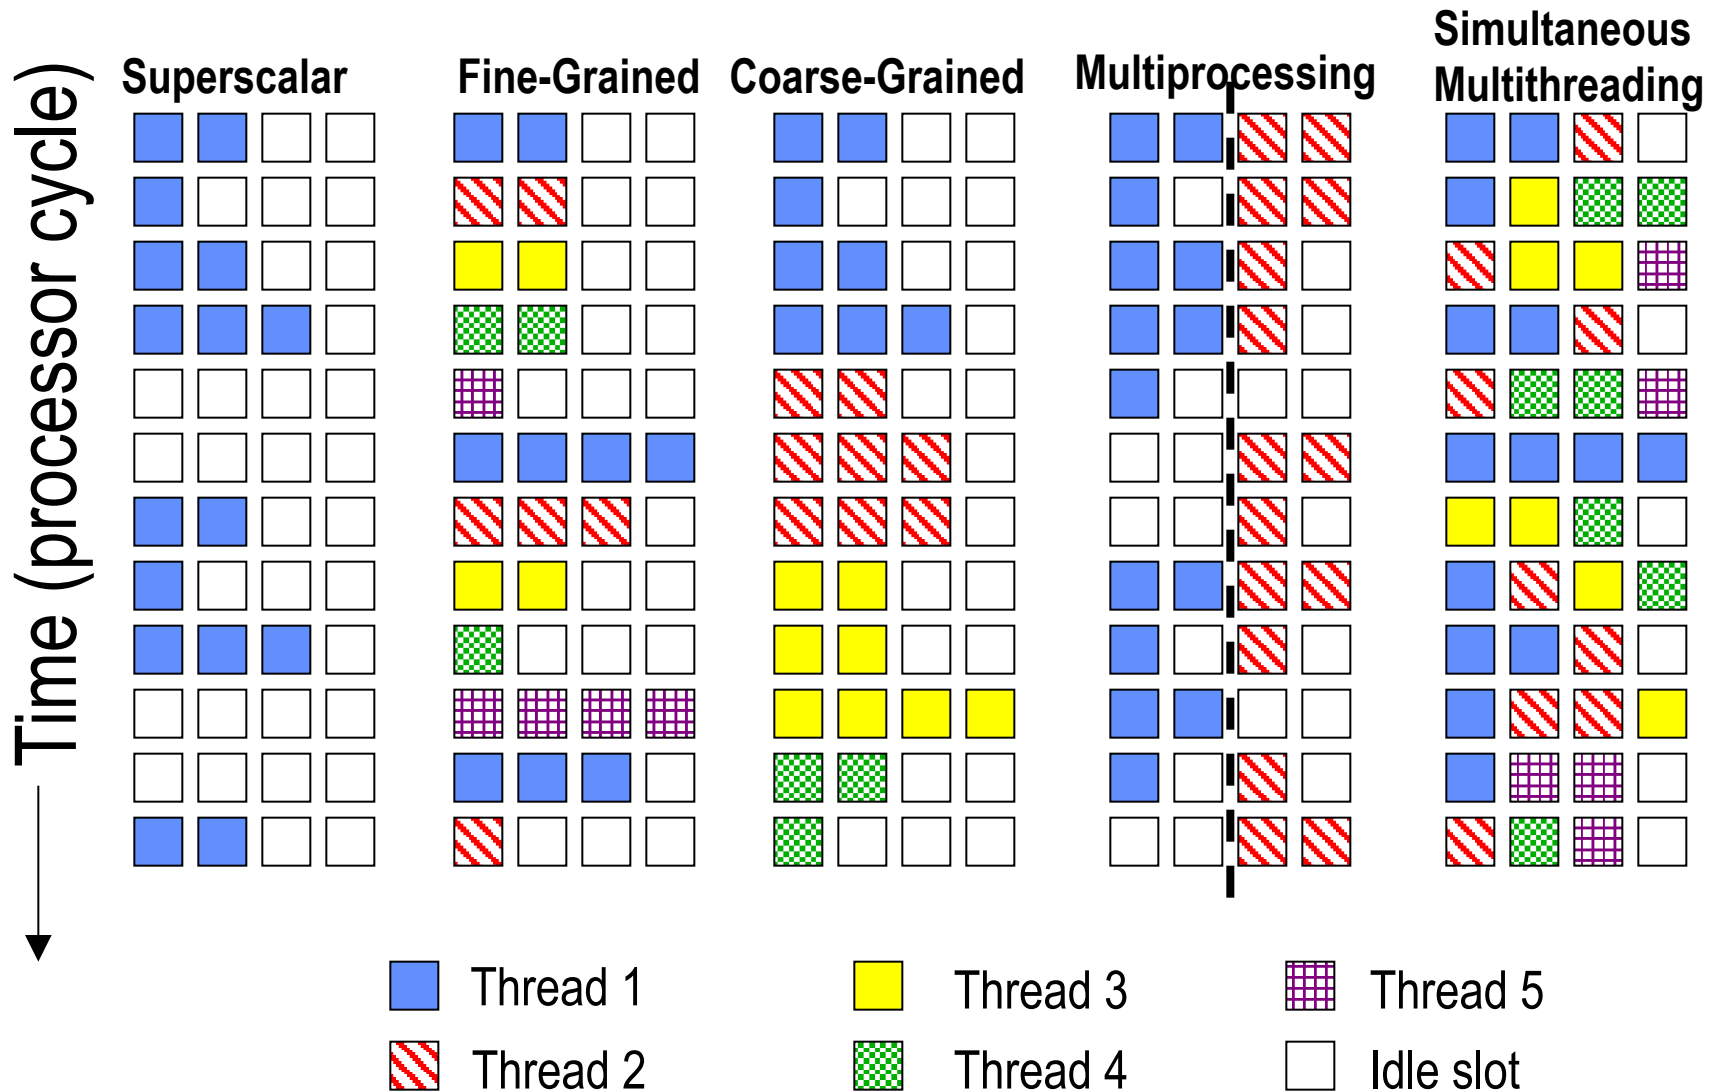
<https://www.tomshardware.com/news/nvidia-in-talks-with-intel-foundry-intel-and-amd-know-all-our-secrets>

SilTerra										
X-FAB										
Dongbu HiTek										
ADI	ADI									
Atmel	Atmel									
Rohm	Rohm									
Sanyo	Sanyo									
Mitsubishi	Mitsubishi									
ON	ON									
Hitachi	Hitachi									
Cypress	Cypress	Cypress								
SkyWater	SkyWater	SkyWater								
Sony	Sony	Sony								
Infineon	Infineon	Infineon								
Sharp	Sharp	Sharp								
Freescale	Freescale	Freescale								
Renesas (NEC)	Renesas	Renesas	Renesas	Renesas						
Toshiba	Toshiba	Toshiba	Toshiba	Toshiba						
Fujitsu	Fujitsu	Fujitsu	Fujitsu	Fujitsu						
TI	TI	TI	TI	TI						
Panasonic	Panasonic	Panasonic	Panasonic	Panasonic	Panasonic					
STMicroelectronics	STM	STM	STM	STM	STM					
HLMC	HLMC		HLMC	HLMC	HLMC					
IBM	IBM	IBM	IBM	IBM	IBM	IBM				
UMC	UMC	UMC	UMC	UMC	UMC		UMC			
SMIC	SMIC	SMIC	SMIC	SMIC	SMIC		SMIC			
AMD	AMD	AMD	GlobalFoundries	GF	GF	GF	GF			
Samsung	Samsung	Samsung	Samsung	Samsung	Samsung	Samsung	Samsung	Samsung	Samsung	Samsung
TSMC	TSMC	TSMC	TSMC	TSMC	TSMC	TSMC	TSMC	TSMC	TSMC	TSMC
Intel	Intel	Intel	Intel	Intel	Intel	Intel	Intel	Intel	Intel	Intel
180 nm	130 nm	90 nm	65 nm	45 nm/40 nm	32 nm/28 nm	22 nm/20 nm	16 nm/14 nm	10 nm	7 nm	5 nm

https://en.wikichip.org/wiki/technology_node



Last Time Lecture: Multithreading



Supercomputer Applications

- Typical application areas
 - Military research (nuclear weapons, cryptography)
 - Scientific research
 - Weather forecasting
 - Oil exploration
 - Industrial design (car crash simulation)
 - Bioinformatics
 - Cryptography
- All involve huge computations on large data set
- Supercomputers: CDC6600, CDC7600, Cray-1, ...
- In 70s-80s, Supercomputer \equiv Vector Machine

Vector Supercomputers




- Epitomized by Cray-1, 1976:
- Scalar Unit
 - Load/Store Architecture
- Vector Extension
 - Vector Registers
 - Vector Instructions
- Implementation
 - Hardwired Control
 - Highly Pipelined Functional Units
 - Interleaved Memory System
 - No Data Caches
 - No Virtual Memory


[©Cray Research, 1976]

Vector Programming Model

Scalar Registers



Vector Registers




Vector Length Register




Vector Arithmetic Instructions

vadd v3,v1,v2



Vector Load and Store Instructions

vls v1, (x1), x2



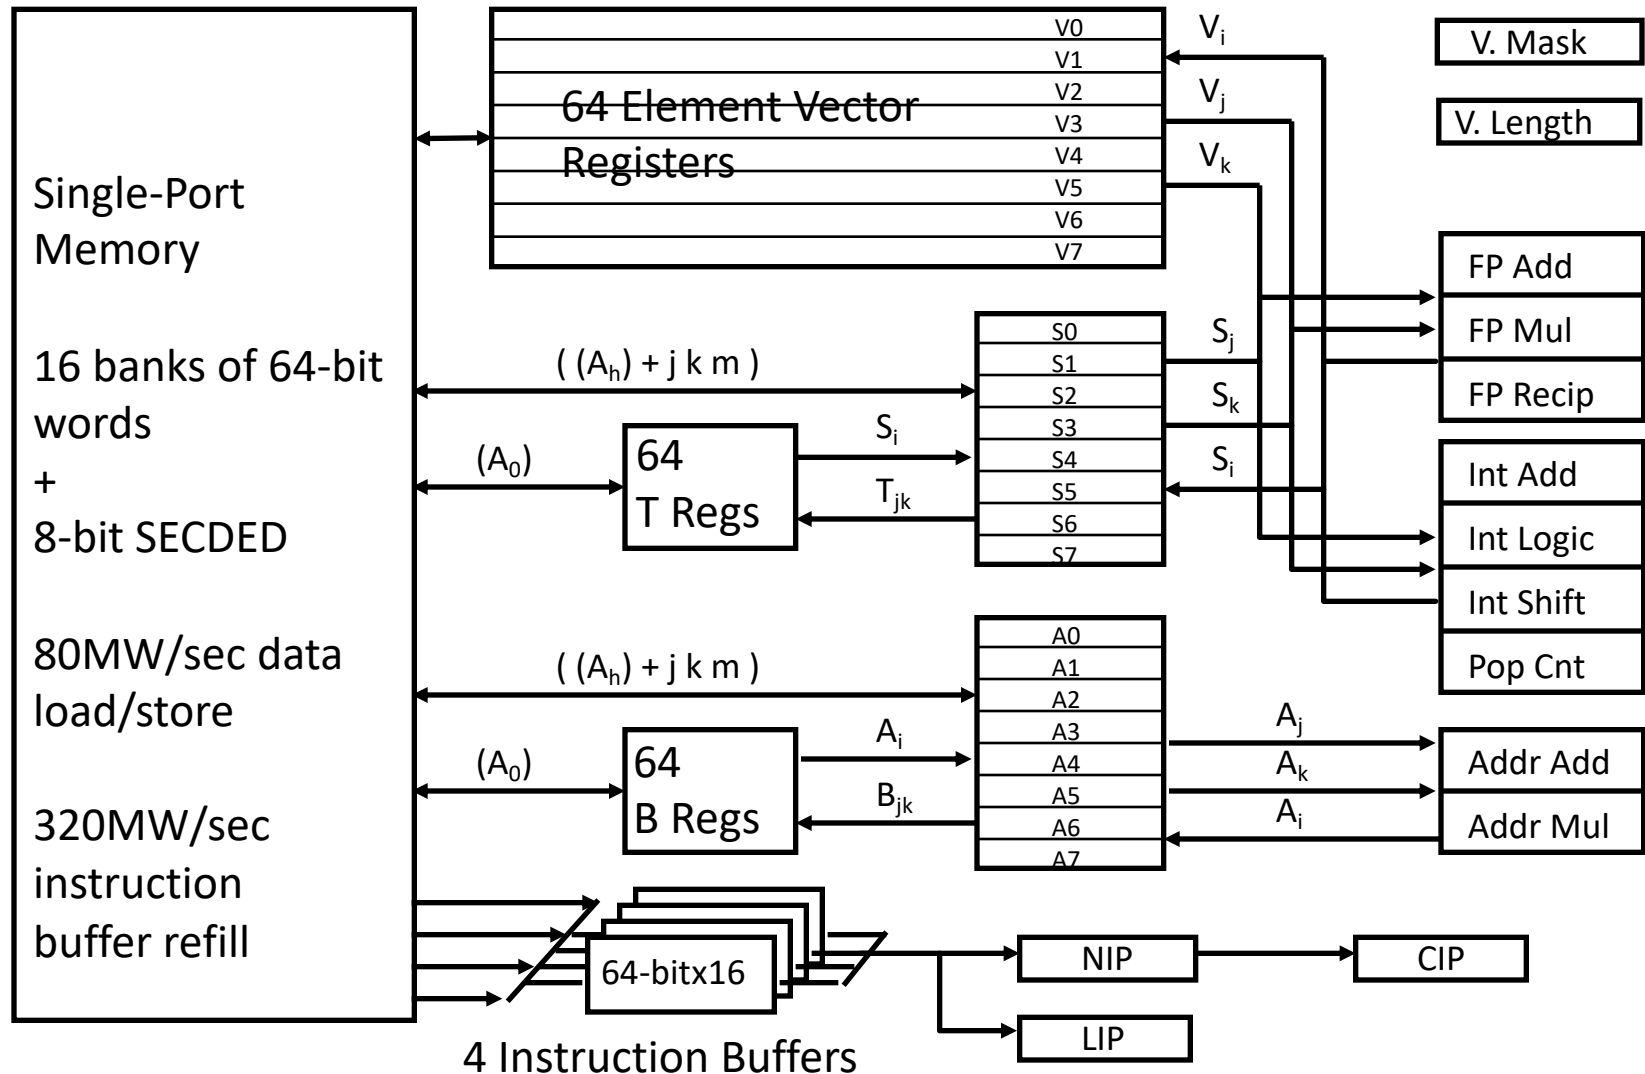
Vector Code Example

```
# C code
for (i=0; i<64; i++)
    C[i] = A[i] + B[i];
```

```
# Scalar Code
    li x4, 64
loop:
    fld f1, 0(x1)
    fld f2, 0(x2)
    fadd.d f3, f1, f2
    fsd f3, 0(x3)
    addi x1, x1, 8
    addi x2, x2, 8
    addi x3, x3, 8
    subi x4, x4, 1
    bnez x4, loop
```

```
# Vector Code
    li x4, 64
    vsetv1 x4
    vld v1, (x1)
    vld v2, (x2)
    vadd v3, v1, v2
    vst v3, (x3)
```

Cray-1 (1976)



memory bank cycle 50 ns processor cycle 12.5 ns (80MHz)

Vector Instruction Set Advantages

- Compact

- one short instruction encodes N operations

- Expressive, tells hardware that these N operations:

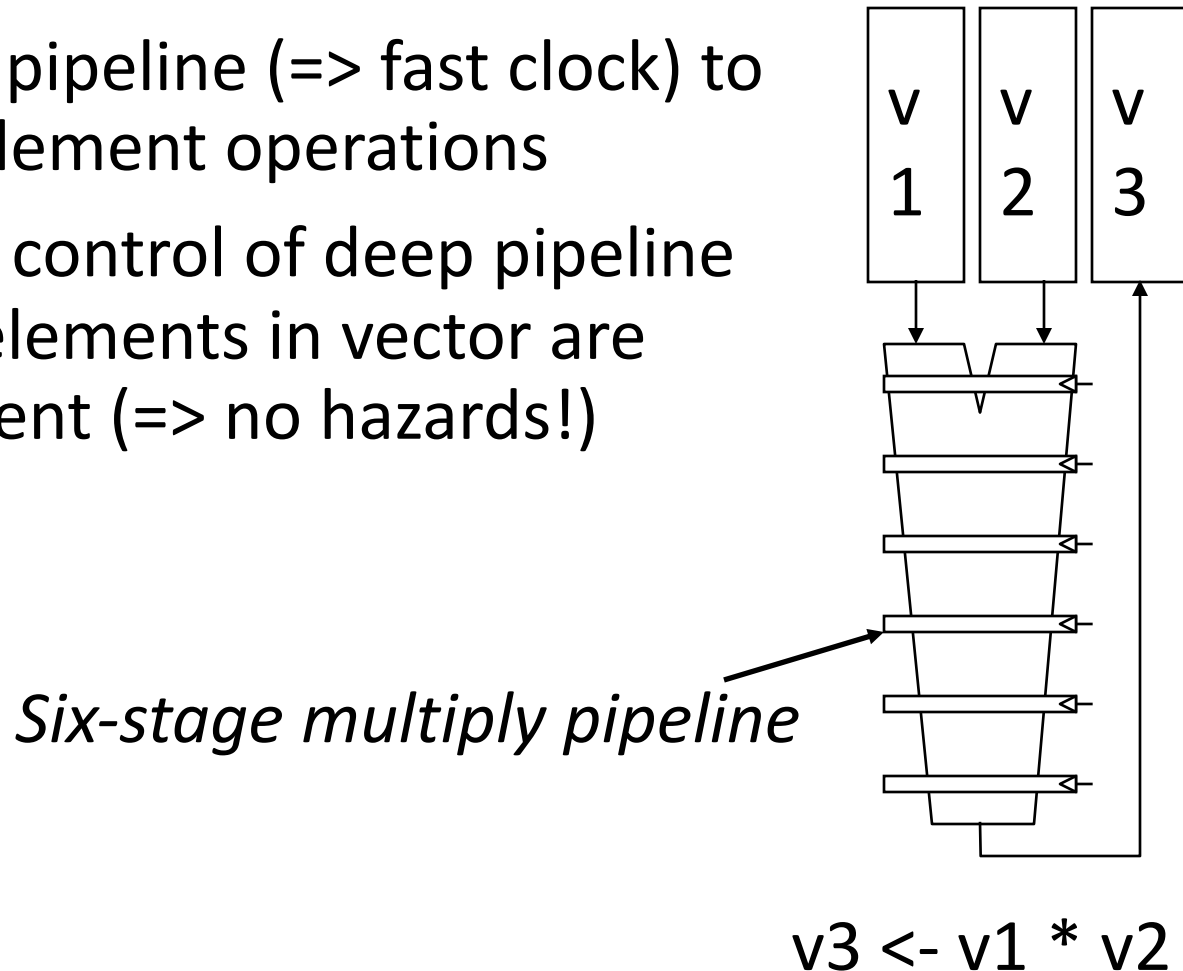
- are independent
- use the same functional unit
- access disjoint registers
- access registers in same pattern as previous instructions
- access a contiguous block of memory (unit-stride load/store)
- access memory in a known pattern (strided load/store)

- Scalable

- can run same code on more parallel pipelines (lanes)

Vector Arithmetic Execution

- Use deep pipeline (\Rightarrow fast clock) to execute element operations
- Simplifies control of deep pipeline because elements in vector are independent (\Rightarrow no hazards!)




Vector Instruction Execution

vadd vc, va, vb





*Execution using
one pipelined
functional unit*

A[6] B[6]
A[5] B[5]
A[4] B[4]
A[3] B[3]




*Execution using
four pipelined
functional units*

A[24] B[24] A[25] B[25] A[26] B[26] A[27] B[27]
A[20] B[20] A[21] B[21] A[22] B[22] A[23] B[23]
A[16] B[16] A[17] B[17] A[18] B[18] A[19] B[19]
A[12] B[12] A[13] B[13] A[14] B[14] A[15] B[15]




Interleaved Vector Memory System

- Bank busy time: Time before bank ready to accept next request
- Cray-1, 16 banks, 4 cycle bank busy time, 12 cycle latency




Vector Unit Structure



T0 Vector Microprocessor (UCB/ICSI, 1995)

*Vector register
elements striped
over lanes*

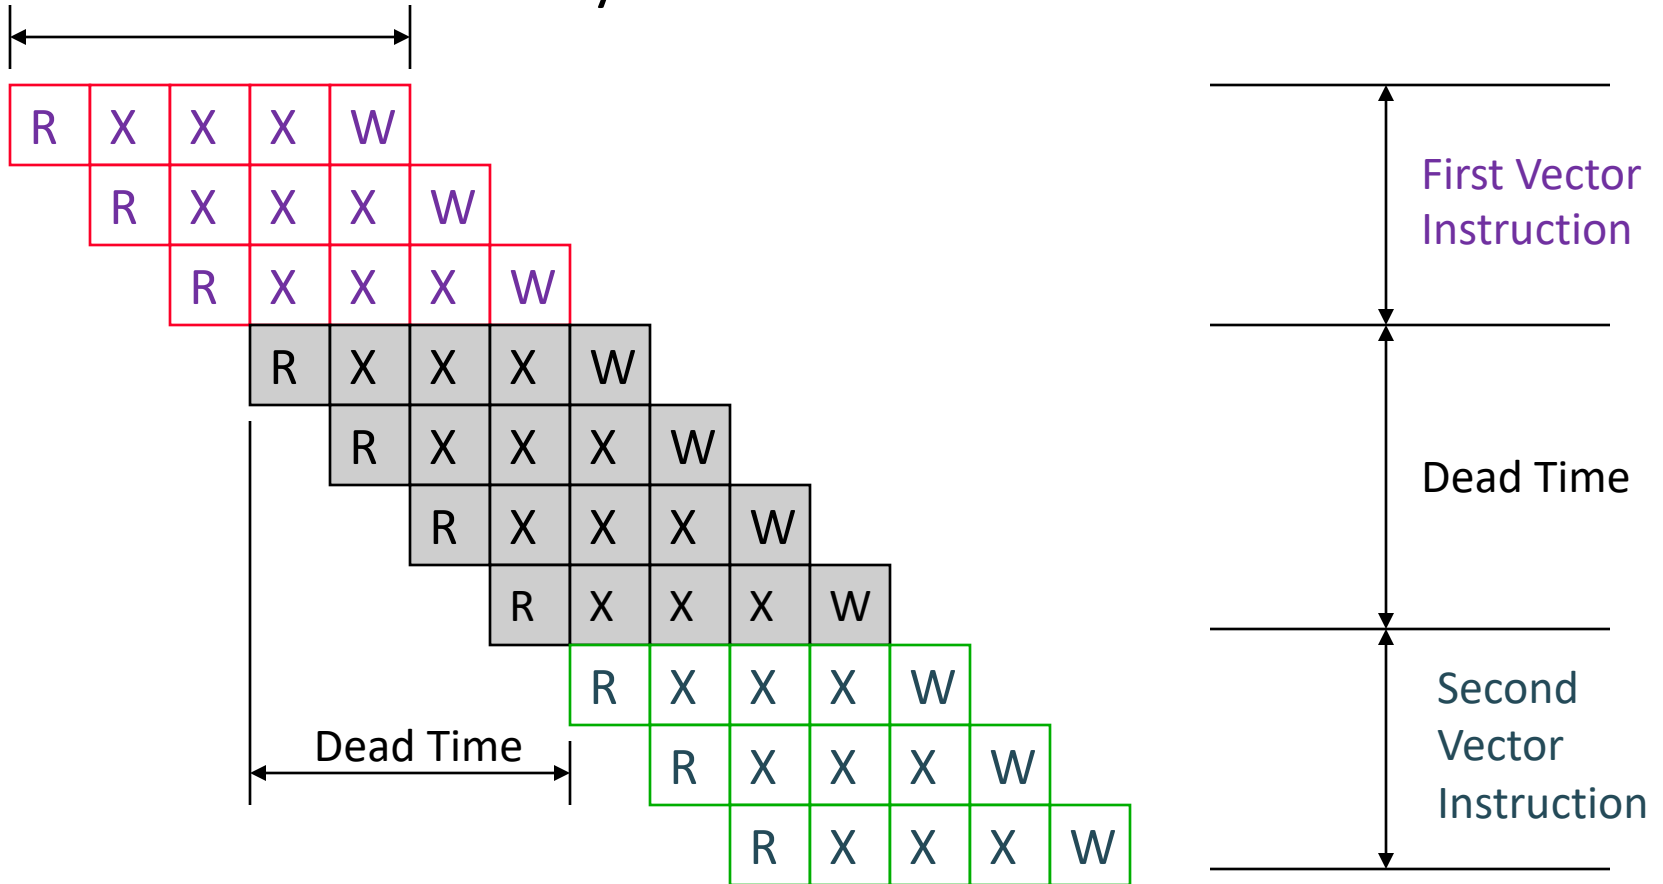


Lane

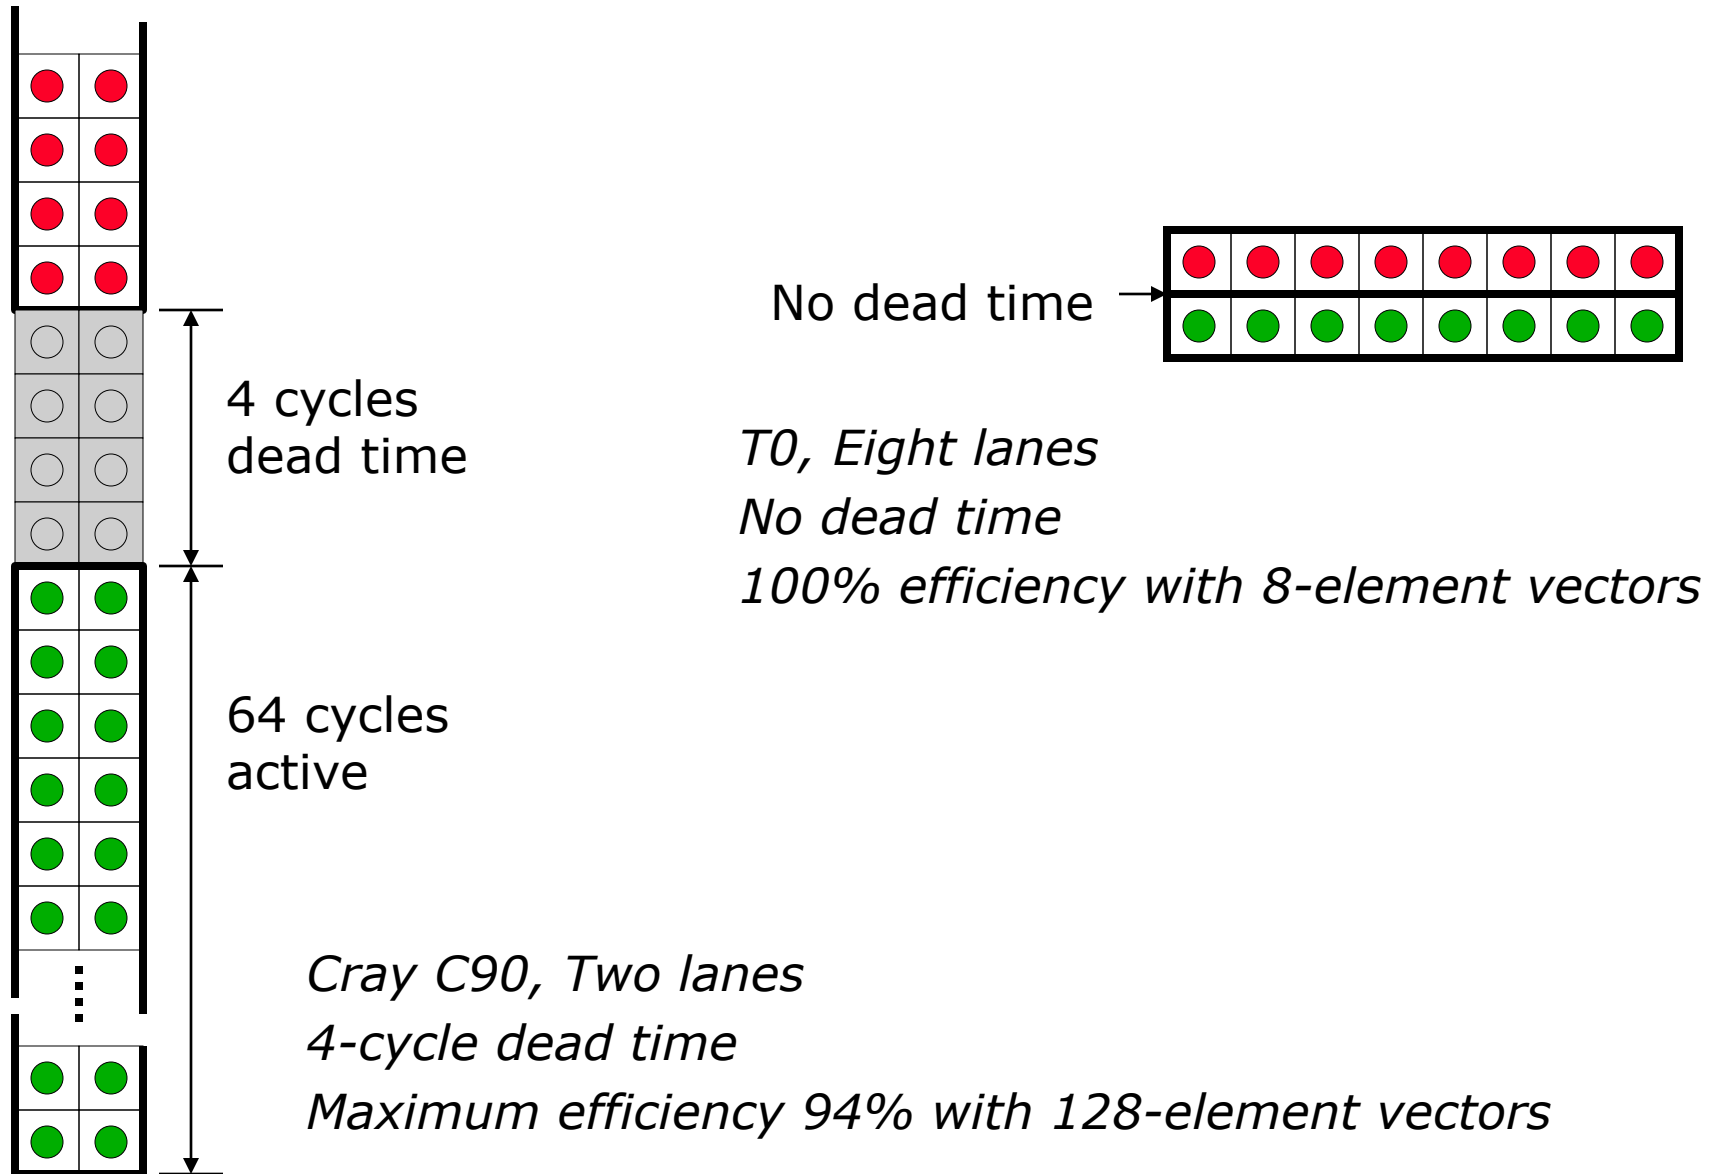
Vector Startup

- Two components of vector startup penalty
 - functional unit latency (time through pipeline)
 - dead time or recovery time (time before another vector instruction can start down pipeline)

Functional Unit Latency




Dead Time and Short Vectors



Vector Chaining


- Vector version of register bypassing
 - introduced with Cray-1

```
vld  v1  
vmul v3, v1, v2  
vadd v5, v3, v4
```



Vector Chaining Advantage

- Without chaining, must wait for last element of result to be written before starting dependent instruction

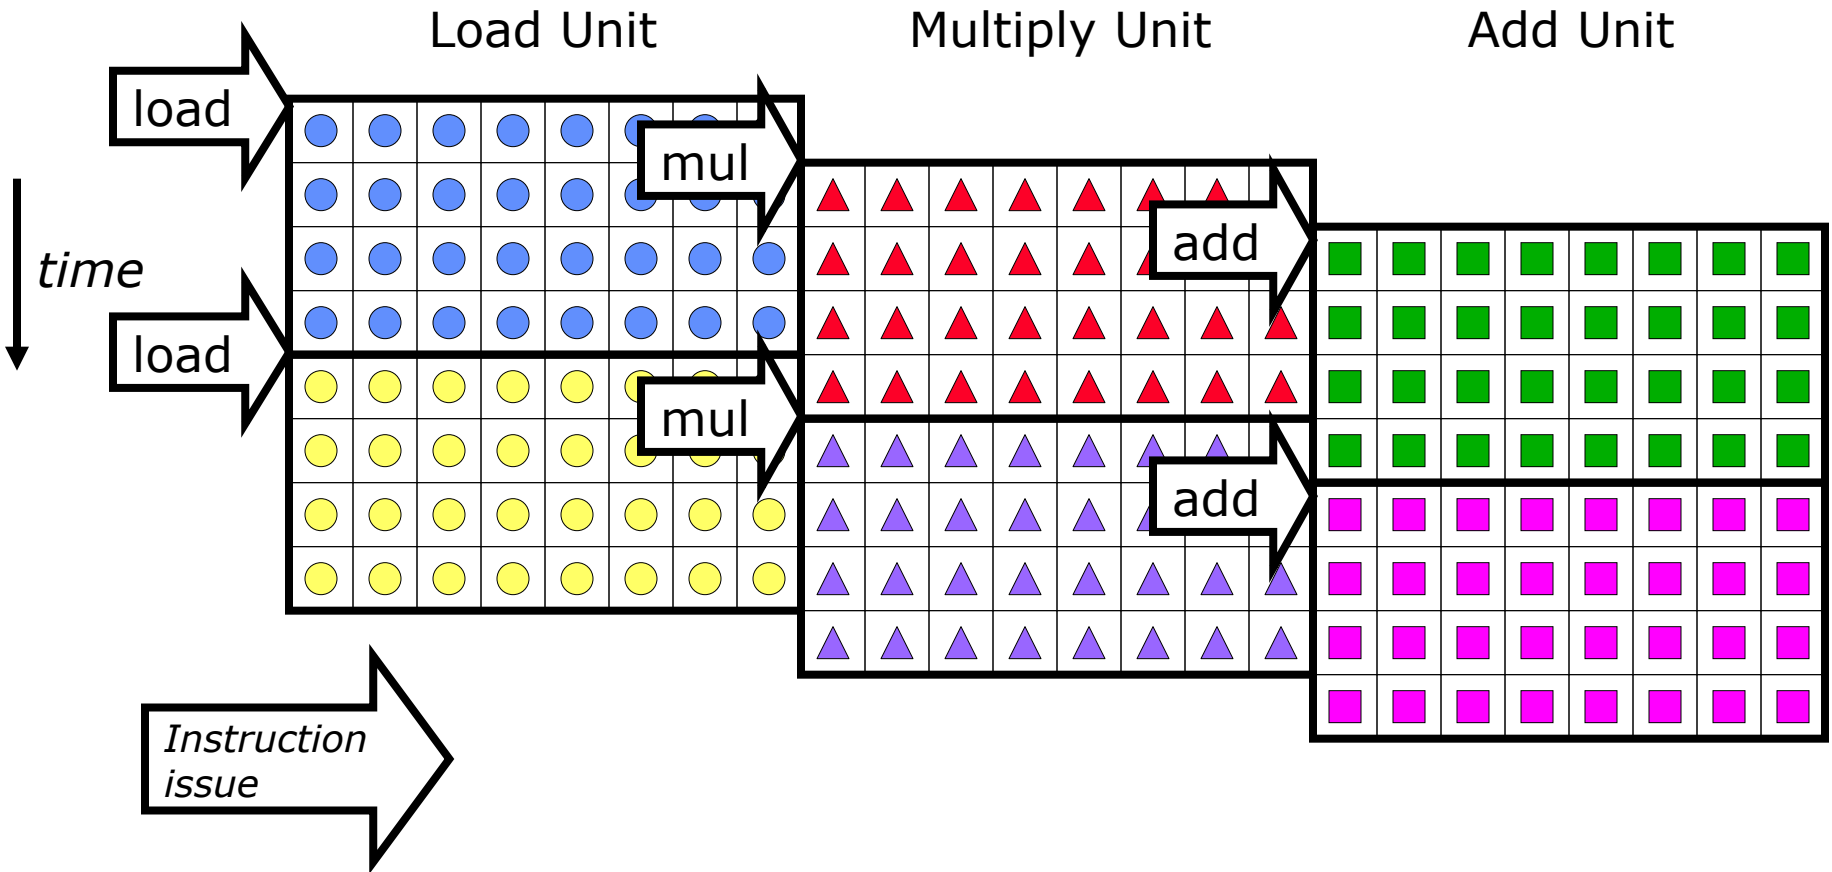


- With chaining, can start dependent instruction as soon as first result appears



Vector Instruction Parallelism

- Can overlap execution of multiple vector instructions
 - example machine has 32 elements per vector register and 8 lanes



Complete 24 operations/cycle while issuing 1 short instruction/cycle

Vector Memory-Memory versus Vector Register Machines

- Vector memory-memory instructions hold all vector operands in main memory
- The first vector machines, CDC Star-100 ('73) and TI ASC ('71), were memory-memory machines
- Cray-1 ('76) was first vector register machine

Example Source Code

```
for (i=0; i<N; i++)  
{  
    C[i] = A[i] + B[i];  
    D[i] = A[i] - B[i];  
}
```

Vector Memory-Memory Code

```
vadd (C) , (A) , (B)  
vsub (D) , (A) , (B)
```

Vector Register Code

```
vld V1, (A)  
vld V2, (B)  
vadd V3, V1, V2  
vst V3, (C)  
vsub V4, V1, V2  
vst V4, (D)
```

Vector Memory-Memory vs. Vector Register Machines

- Vector memory-memory architectures (VMMA) require greater main memory bandwidth, why?
 - All operands must be read in and out of memory
- VMMA make it difficult to overlap execution of multiple vector operations, why?
 - Must check dependencies on memory addresses
- VMMA incur greater startup latency
 - Scalar code was faster on CDC Star-100 for vectors < 100 elements
 - For Cray-1, vector/scalar breakeven point was around 2-4 elements
- Apart from CDC follow-ons (Cyber-205, ETA-10) all major vector machines since Cray-1 have had vector register architectures
- (we ignore vector memory-memory from now on)

CS152 Administrivia

- HW 4 out this week
- Lab 3 due 3/23
- Midterm feedback

CS252 Administrivia


- Readings on OoO architectures this Wednesday

Automatic Code Vectorization

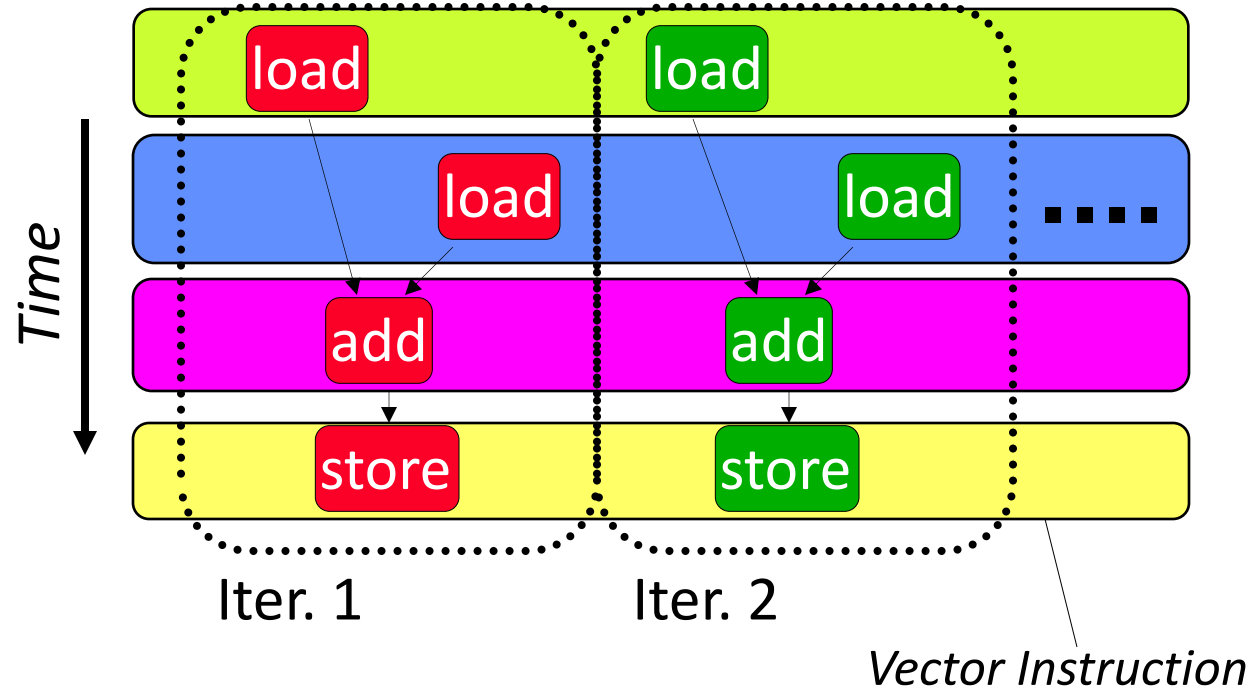
```
for (i=0; i < N; i++)
```

```
  C[i] = A[i] + B[i];
```

Scalar Sequential Code



Vectorized Code



Vectorization is a massive compile-time reordering of operation sequencing

⇒ requires extensive loop-dependence analysis


Vector Stripmining

Problem: Vector registers have finite length

Solution: Break loops into pieces that fit in registers, “*Stripmining*”

```
andi x1, xN, 63 # N mod 64
vsetvl x1      # Do remainder
loop:
vld v1, (xA)
slli x2, x1, 3 # Multiply by 8
add xA, xA, x2 # Bump pointer
vld v2, (xB)
add xB, xB, x2
vadd v3, v1, v2
vst v3, (xC)
add xC, xC, x2
sub xN, xN, x1 # Subtract elements
li x1, 64
vsetvl x1      # Reset full length
bgtz xN, loop  # Any more to do?
```

```
for (i=0; i<N; i++)
    C[i] = A[i]+B[i];
```



Vector Conditional Execution

Problem: Want to vectorize loops with conditional code:

```
for (i=0; i<N; i++)  
    if (A[i]>0) then  
        A[i] = B[i];
```

Solution: Add vector *mask* (or *flag*) registers

- vector version of predicate registers, 1 bit per element

...and *maskable* vector instructions

- vector operation becomes bubble (“NOP”) at elements where mask bit is clear

Code example:

```
vld vA, (xA)      # Load entire A vector  
vgt vA, f0        # Set bits in mask register where A>0  
vld vA, (xB)      # Load B vector into A under mask  
vst vA, (xA)      # Store A back to memory under mask
```

Masked Vector Instructions

Simple Implementation

- execute all N operations, turn off result writeback according to mask


M[7]=1 A[7] B[7]

M[6]=0 A[6] B[6]

M[5]=1 A[5] B[5]

M[4]=1 A[4] B[4]

M[3]=0 A[3] B[3]



M[2]=0

M[1]=1

M[0]=0

C[0]

Density-Time Implementation

- scan mask vector and only execute elements with non-zero masks

M[7]=1

M[6]=0

M[5]=1

M[4]=1

M[3]=0

M[2]=0

M[1]=1

M[0]=0

A[7] B[7]

C[5]

C[4]

C[3]

C[2]

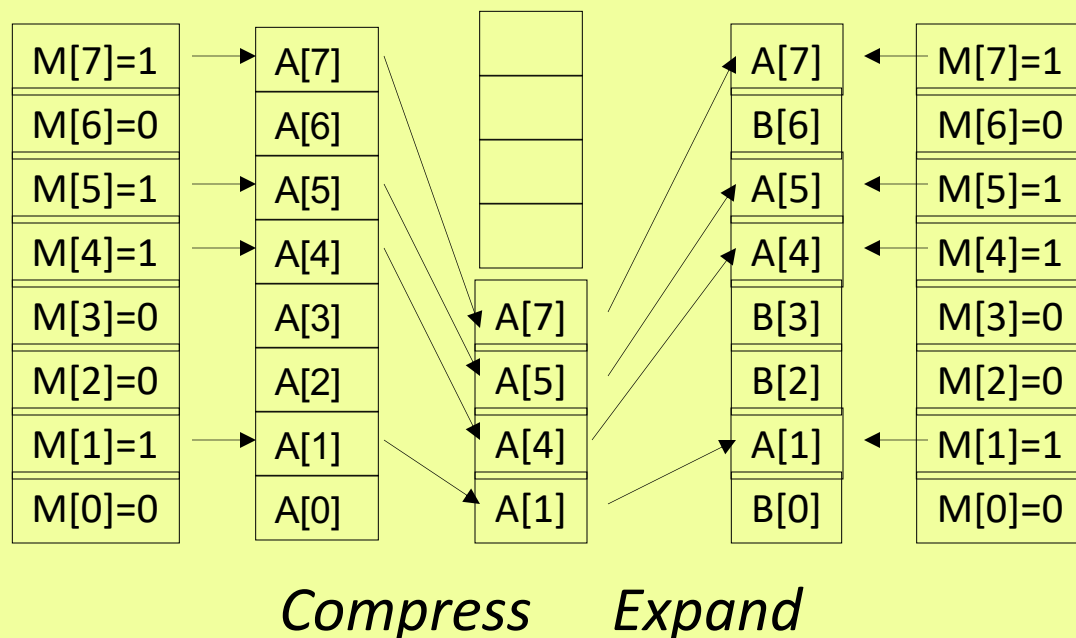
C[1]

C[0]

Write data port

Compress/Expand Operations

- Compress packs non-masked elements from one vector register contiguously at start of destination vector register
 - population count of mask vector gives packed vector length
- Expand performs inverse operation



Used for density-time conditionals and also for general selection operations

Vector Reductions

Problem: Loop-carried dependence on reduction variables

```
sum = 0;
for (i=0; i<N; i++)
    sum += A[i];  # Loop-carried dependence on sum
```

Solution: Re-associate operations if possible, use binary tree to perform reduction

```
# Rearrange as:
sum[0:VL-1] = 0          # Vector of VL partial sums
for(i=0; i<N; i+=VL)    # Stripmine VL-sized chunks
    sum[0:VL-1] += A[i:i+VL-1]; # Vector sum
# Now have VL partial sums in one vector register
do {
    VL = VL/2;           # Halve vector length
    sum[0:VL-1] += sum[VL:2*VL-1] # Halve no. partials
} while (VL>1)
```

Vector Scatter/Gather

Want to vectorize loops with indirect accesses:

```
for (i=0; i<N; i++)  
    A[i] = B[i] + C[D[i]]
```

Indexed load instruction (*Gather*)

```
vld vD, (xD)           # Load indices in D vector  
vlx vC, (xC), vD        # Load indexed from xC base  
vld vB, (xB)           # Load B vector  
vadd vA, vB, vC         # Do add  
vst vA, (xA)           # Store result
```

Histogram with Scatter/Gather


Histogram example:

```
for (i=0; i<N; i++)  
    A[B[i]]++;
```

Is following a correct translation?


```
vld vB, (xB)          # Load indices in B vector  
vlx vA, (xA), vB      # Gather initial A values  
vadd vA, vA, 1        # Increment  
vsx vA, (xA), vB      # Scatter incremented values
```

Packed SIMD Extensions



- Very short vectors added to existing ISAs for microprocessors
- Use existing 64-bit registers split into 2x32b or 4x16b or 8x8b
 - Lincoln Labs TX-2 from 1957 had 36b datapath split into 2x18b or 4x9b
 - Newer designs have wider registers
 - 128b for PowerPC AltiVec, Intel SSE2/3/4
 - 256b/512b for Intel AVX

- Single instruction operates on all elements within register



Packed SIMD versus Vectors

- Limited instruction set:
 - no vector length control
 - no strided load/store or scatter/gather
 - unit-stride loads must be aligned to 64/128-bit boundary
- Limited vector register length:
 - requires superscalar dispatch to keep multiply/add/load units busy
 - loop unrolling to hide latencies increases register pressure
- Trend towards fuller vector support in microprocessors
 - Better support for misaligned memory accesses
 - Support of double-precision (64-bit floating-point)
 - New Intel AVX spec (announced April 2008), 256b vector registers (expandable up to 1024b), gather added, scatter to follow
 - ARM Scalable Vector Extensions (SVE)

Acknowledgements

- This course is partly inspired by previous MIT 6.823 and Berkeley CS252 computer architecture courses created by my collaborators and colleagues:
 - Arvind (MIT)
 - Krste Asanovic (MIT/UCB)
 - Joel Emer (Intel/MIT)
 - James Hoe (CMU)
 - John Kubiatowicz (UCB)
 - David Patterson (UCB)