

CS 152 Computer Architecture and Engineering

CS252 Graduate Computer Architecture

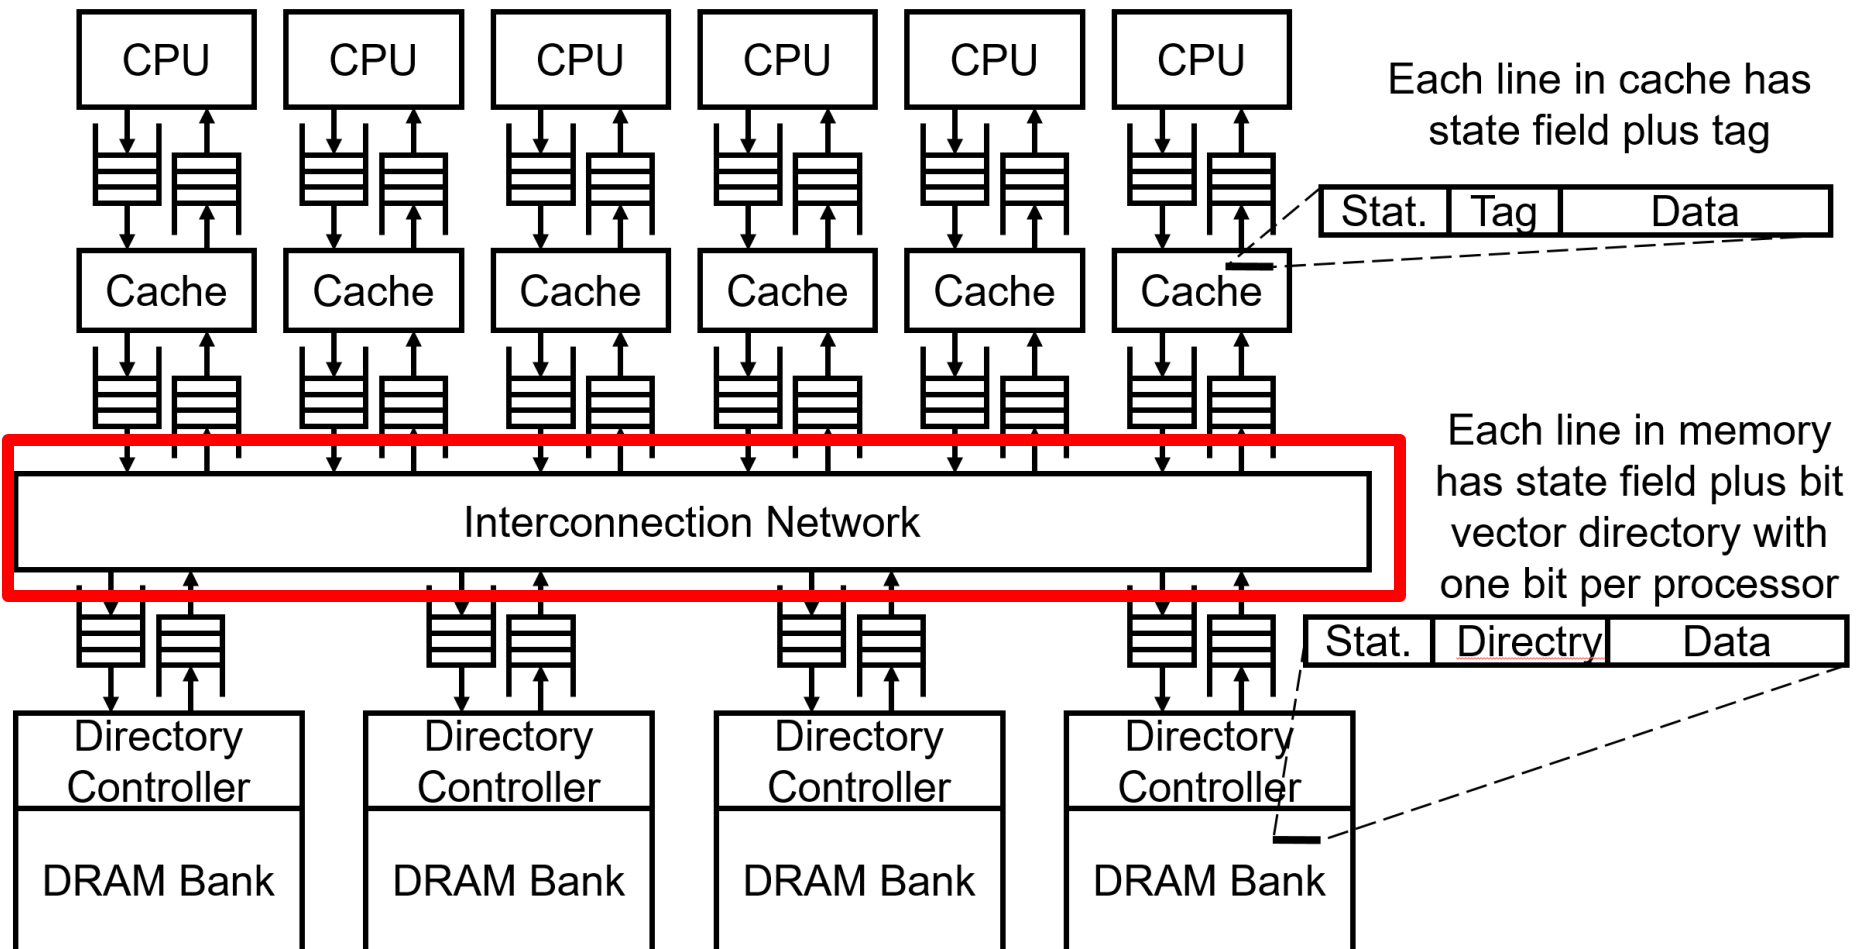
Lecture 25 - Interconnects

Jerry Zhao

Electrical Engineering and Computer Sciences
University of California at Berkeley

jzh@berkeley.edu

Directory Cache Protocol




This Lecture

- Interconnects are not traditionally covered in 152
 - More in-depth in EECS251B
- This Lecture:
 - Topology
 - Flow-control
 - Routing
 - Micro-architecture
 - Notable interconnect implementations


Interconnect Scale



Rack-to-rack
(datacenter)



Chip-to-chip



On-chip

- This lecture will focus on **on-chip interconnects for shared-memory systems**
- Same principles apply when scaling out

Interconnect Agents

Agents on a **shared-memory interconnect** can be:


- Cores/private-caches
- Shared-caches/memory banks
- Accelerators
- I/O devices (ethernet, PCI-E, HDMI)
- Off-chip memory channels (DDR)

Alternative interpretations:

- FPGA routing between CLBs
- Datapath between PEs in an accelerator

Architectural View



- Interconnect agents speak a common memory protocol
- AMBA/AXI/ACE/TileLink/Wishbone
- Interconnect implementation routes traffic to correct destination
 - Client-to-manager traffic: route by address
 - Manager-to-client traffic: route by client identifier (ID)
- Across homogeneous channels/banks, address ranges can be striped or hashed for load-balancing





Interconnect Components

- **Topology** – physical layout of routers (nodes) and channels (links)
- **Flow control** – mechanics of how messages travel from source-to-destination
- **Routing** – which paths messages take through the network
- **Micro-architecture** – implementation of routers and channels

Topologies



- Shared-bus
 - Advantages: cheap & easy to implement, broadcast, serialized messages
 - Disadvantages: low bandwidth, tri-state logic
- Crossbar – all-to-all connection
 - Advantages: high bandwidth due to all-to-all routing, no contention, serialized messages, predictable latency
 - Disadvantages: $O(n^2)$ scaling, scales poorly past 4x4 networks



- **1D torus/ring (unidirectional/bidirectional)**
 - Advantages: simple to implement, well-behaved
 - Disadvantages: low bisection bandwidth, high-hop-count
- **2D mesh**
 - Advantages: scalable with good bandwidth/low-latency
 - Disadvantages: complex routing for deadlock-freedom



These are **network-on-chips (NoCs)**,
with proper routers and channels

Topology Properties

- Diameter – longest shortest path between any two nodes
- Bisection bandwidth – minimum bandwidth between any two partitions

Assuming single-direction channel bandwidth is **b**:


- 1D unidirectional torus with **n** nodes
 - Diameter = $n - 1$
 - Bisection bandwidth = $2b$
- 1D bidirectional torus with **n** nodes:
 - Diameter = $n / 2$
 - Bisection bandwidth = $4b$
- **n x n** 2D mesh
 - Diameter = $2n$
 - Bisection bandwidth = $2nb$



Topology Variations


■ Hypercube

- Popular in the 70s/80s, connect multi-board machines
- Low diameter, high bandwidth
- Simpler 1D/2D topologies won for single-chip multi-processors




■ 3D mesh

- Extension of 2D mesh
- 3D integration with TSVs (through-silicon-vias)




■ Butterfly/Clos networks


- Multi-stage crossbar
- Used in big network switches




Topology Spectrum



Mesh



CMesh



Clos



Crossbar

Increasing diameter

Easy to design
Hard to program

Hard to design
Easy to program


Increasing radix

Radix – Number of inputs and outputs of each switching node
Diameter – largest minimal hop count over all node pairs

Flow Control – Packets/Flits

Packet: unit of routing/flow-control, a single message in the network

- Size can be from ~40b (just address) to ~600b (addr + cache line) to 1KB (streaming data)



Packets are subdivided into equal-length **flits**,

- **Head/tail flits** mark the boundaries of a packet
- All flits follow the same path in order
- Variable packet size => variable flits/packet

Flow Control – Packet/Flit Tradeoff


- Small packets/flits => better resource utilization
- Large packets/flits => reduce overheads

Flow Control – Two Types

- **Bufferless** : no buffers between ingress and egress, packets move straight from channel-to-channel
 - Like a train system/railroad
- **Buffered**: packets can stall in a buffer in the network while they wait for a channel
 - Like a road system with intersections
 - Lanes are buffers

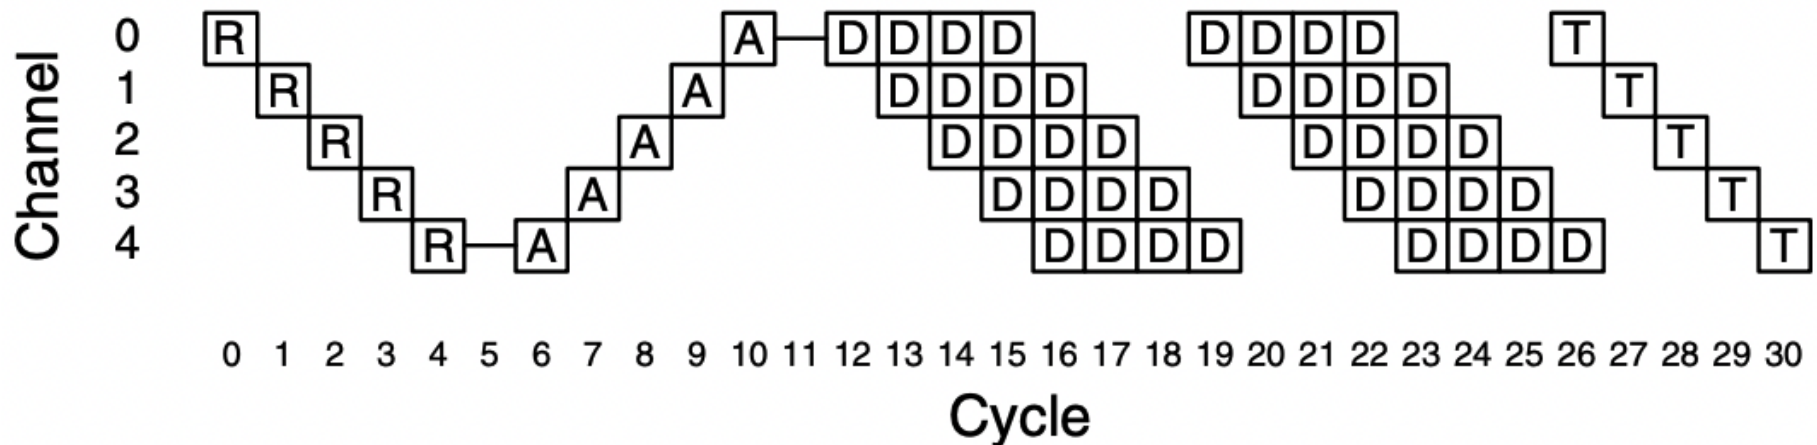
Flow Control - Dropping

- **Dropping flow control:** nack packets if they fail to allocate a channel, sender should retry
 - Sender must buffer all inflight requests
 - Low throughput under load, requires buffers



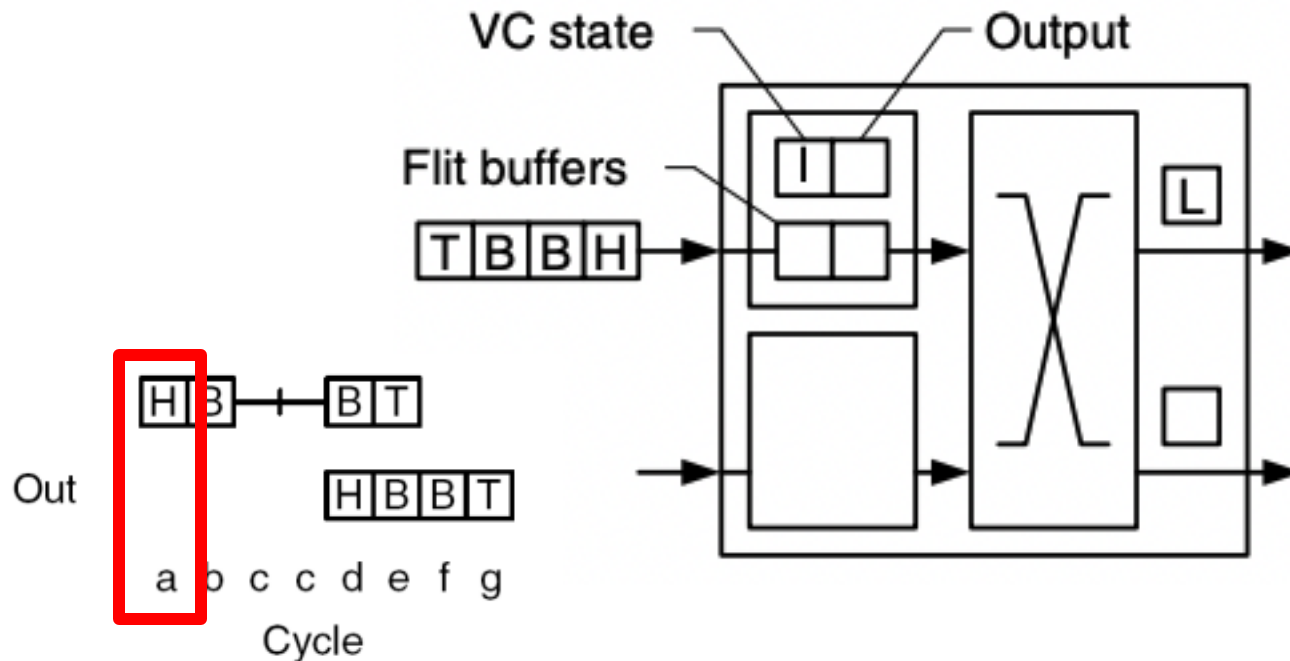
Flow Control – Circuit-switched

- **Circuit-switched flow control:** try to acquire the entire path from ingress to egress, only send packet if acquired
 - Low throughput/high latency
 - Telephone network works like this



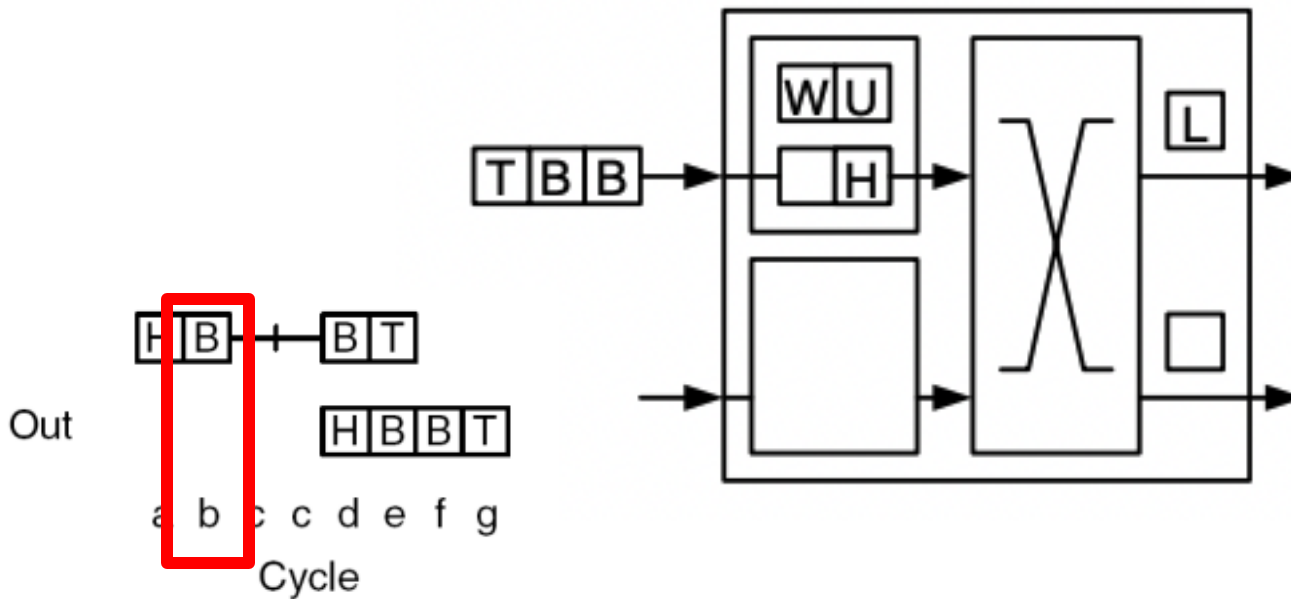
Flow Control - Buffered

- **Packet-buffered:** Buffer the entire packet, then send over the link
 - High buffer capacity requirements
- **Flit-buffered “wormhole” flow-control:** Buffer flits, not packets
 - Flits travel as a “worm” from head to tail



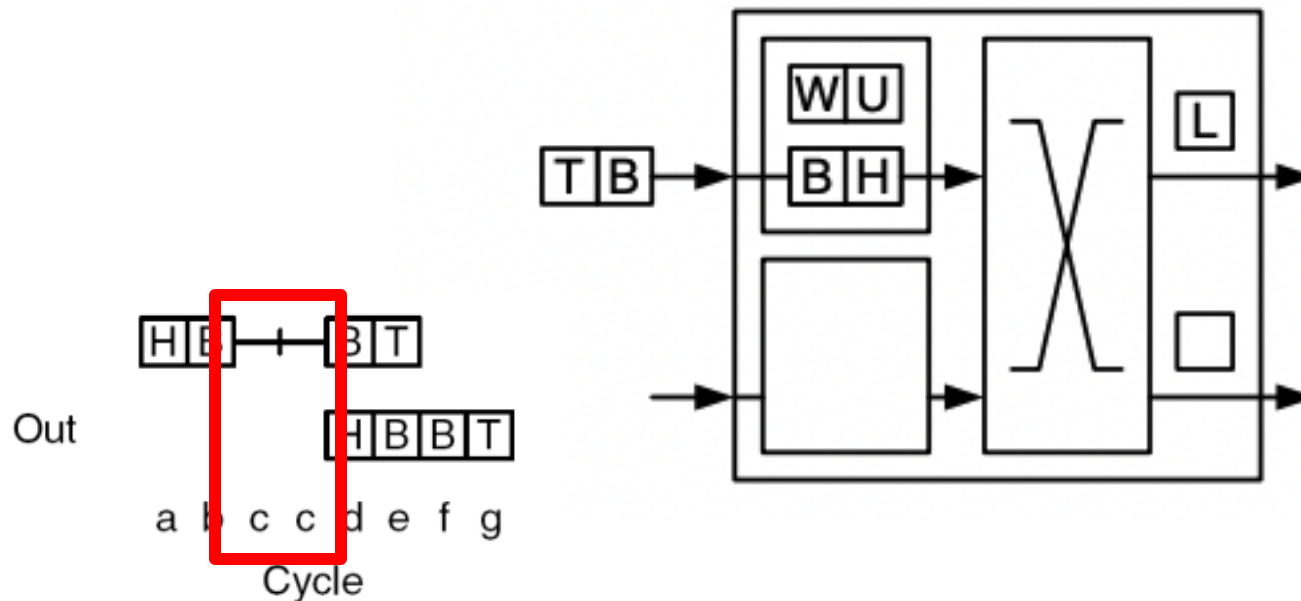
Flow Control - Buffered

- **Packet-buffered:** Buffer the entire packet, then send over the link
 - High buffer capacity requirements
- **Flit-buffered “wormhole” flow-control:** Buffer flits, not packets
 - Flits travel as a “worm” from head to tail



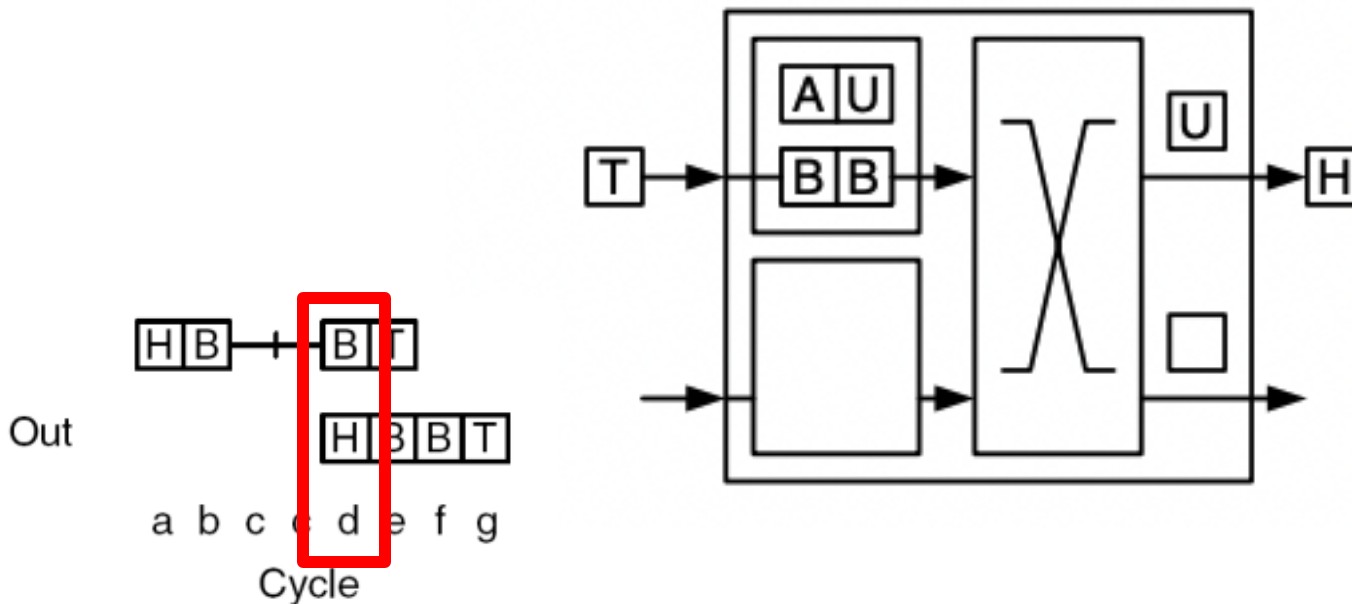
Flow Control - Buffered

- **Packet-buffered:** Buffer the entire packet, then send over the link
 - High buffer capacity requirements
- **Flit-buffered “wormhole” flow-control:** Buffer flits, not packets
 - Flits travel as a “worm” from head to tail




Flow Control - Buffered

- **Packet-buffered:** Buffer the entire packet, then send over the link
 - High buffer capacity requirements
- **Flit-buffered “wormhole” flow-control:** Buffer flits, not packets
 - Flits travel as a “worm” from head to tail




Flow Control - Buffered

- **Packet-buffered:** Buffer the entire packet, then send over the link
 - High buffer capacity requirements
- **Flit-buffered “wormhole” flow-control:** Buffer flits, not packets
 - Flits travel as a “worm” from head to tail




Flow Control - Buffered

- **Packet-buffered:** Buffer the entire packet, then send over the link
 - High buffer capacity requirements
- **Flit-buffered “wormhole” flow-control:** Buffer flits, not packets
 - Flits travel as a “worm” from head to tail



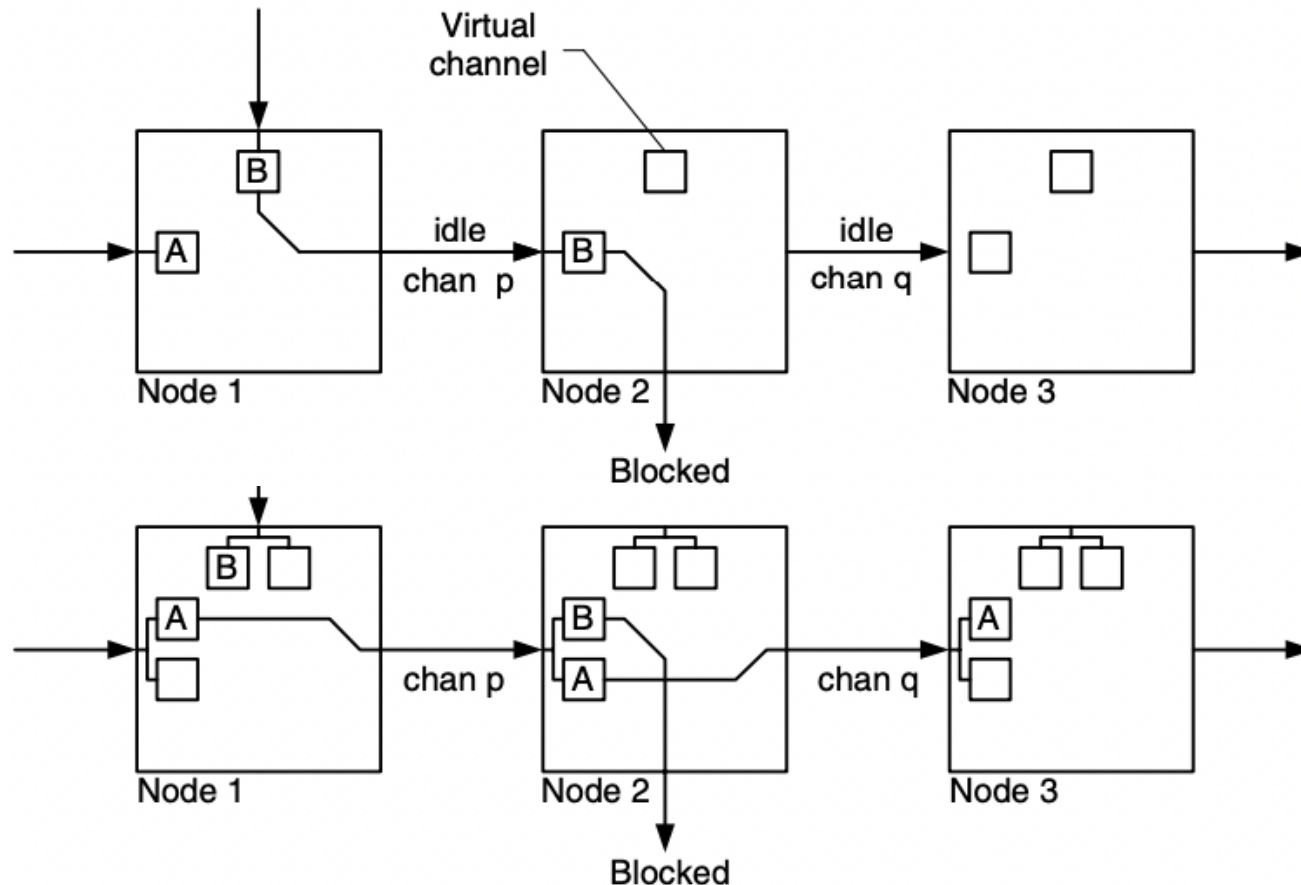
Flow Control - Buffered

- **Packet-buffered:** Buffer the entire packet, then send over the link
 - High buffer capacity requirements
- **Flit-buffered “wormhole” flow-control:** Buffer flits, not packets
 - Flits travel as a “worm” from head to tail



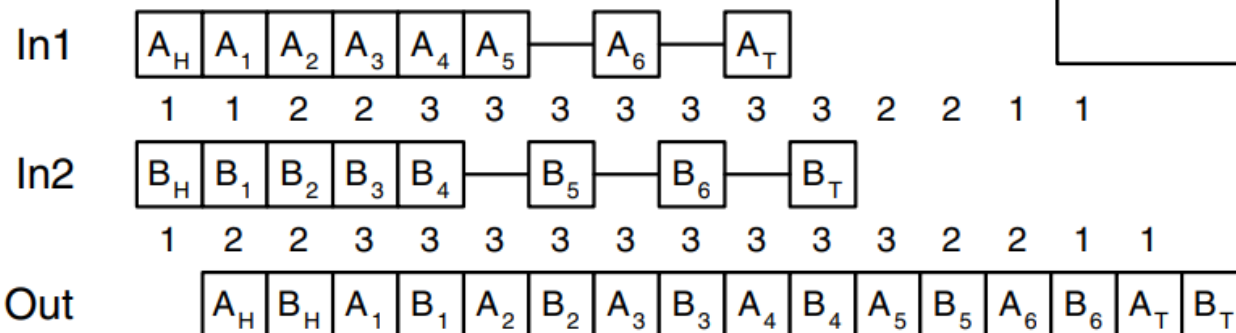

Flow Control – Virtual Channels

- Allow flits from different packets to **time-multiplex** access to a physical channel
- **Flit buffers are "virtual channels" (VCs)**
- Standard approach for NoC architectures




Flow Control – Virtual Channels

- Suppose A/B are contending for a single physical channel
- What happens if fair arbitration is used between **A flits** and **B flits**?



A downstream

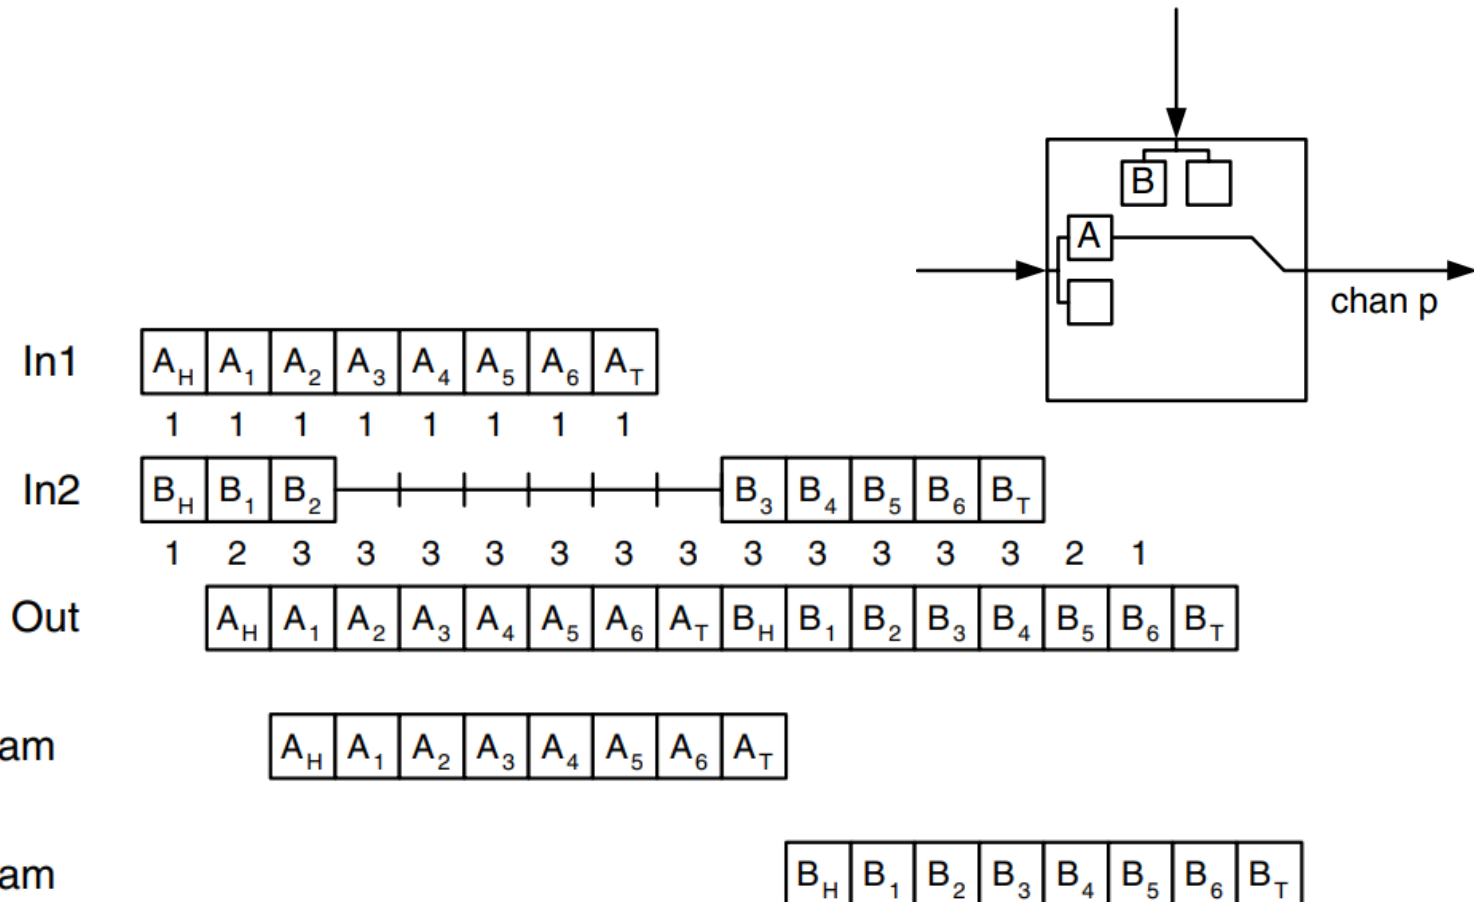


B downstream



Flow Control – Virtual Channels


- Suppose A/B are contending for a single physical channel
- Winner-take-all arbitration improves latency



Routing

- **Deterministic** – no consideration for other traffic
 - Easy to implement, can fully calculate route at the source
 - Poor load-balancing
- **Adaptive** – consider existing traffic in the network
 - Consider network's current state
 - Can't query the entire network before routing, need to compute the route on-the-fly
- **Minimal** – route along minimal paths only


Deadlock-free routing



- Generalizes to more complex topologies – 2D mesh
- **Solution A:** break circular dependencies in channel dependency graph
 - Disadvantage: Requires lots of VCs (buffers)
- **Solution B:** disallow certain physical routes
 - Dimension ordered routing: route first in X, then Y
 - Disadvantage: Creates network hotspots
- Both A and B are used in complex routing algorithms

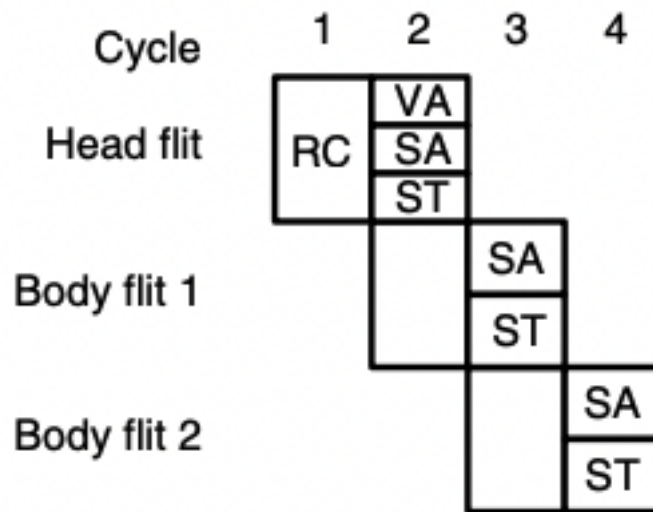

Router Microarchitecture

- Virtual channels (buffers)
- Allocators/arbiters schedule access to channels, buffers, switch
- Small crossbar-switch to connect inputs with outputs




Router Pipelining

- Hop latency is critical, interconnect latency might be on the critical path for a L1/L2 miss
 - 1 cycle on a hop, is 2 cycles round trip
- Modern NoCs are 1-2 cycles per hop
 - Speculation allows merging of pipeline stages



Evaluating Interconnects

- Latency vs Offered Traffic
 - Measure average latency as across offered traffic
 - Traffic limit is the network capacity
 - Latency limit is zero-load latency



QoS – Quality of Service

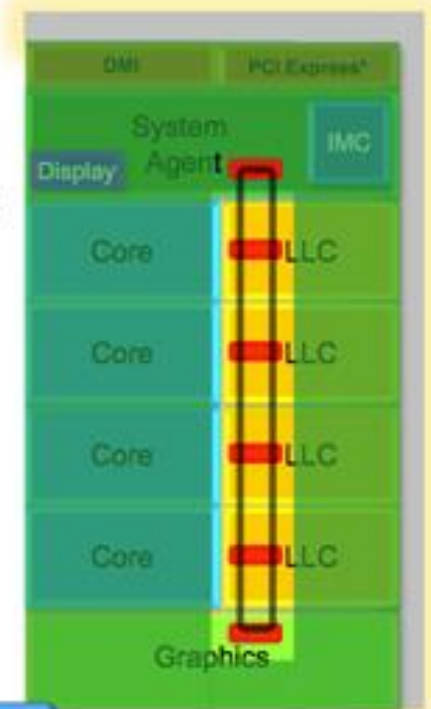
- **Service classes** – categorization of traffic by properties
 - **Real time** – camera input
 - **Latency sensitive** – CPU cache misses
 - **Bandwidth sensitive** – display out
 - **Best effort** - ethernet

☞ Typically service classes have some traffic limit. Why?
- **Service guarantees** – achievable bandwidth/latency per service class
- **Fairness** – traffic from the same class should be prioritized equally
- How to implement?
 - Allocator/arbiter prioritization
 - Resource isolation

Intel Sandy Bridge

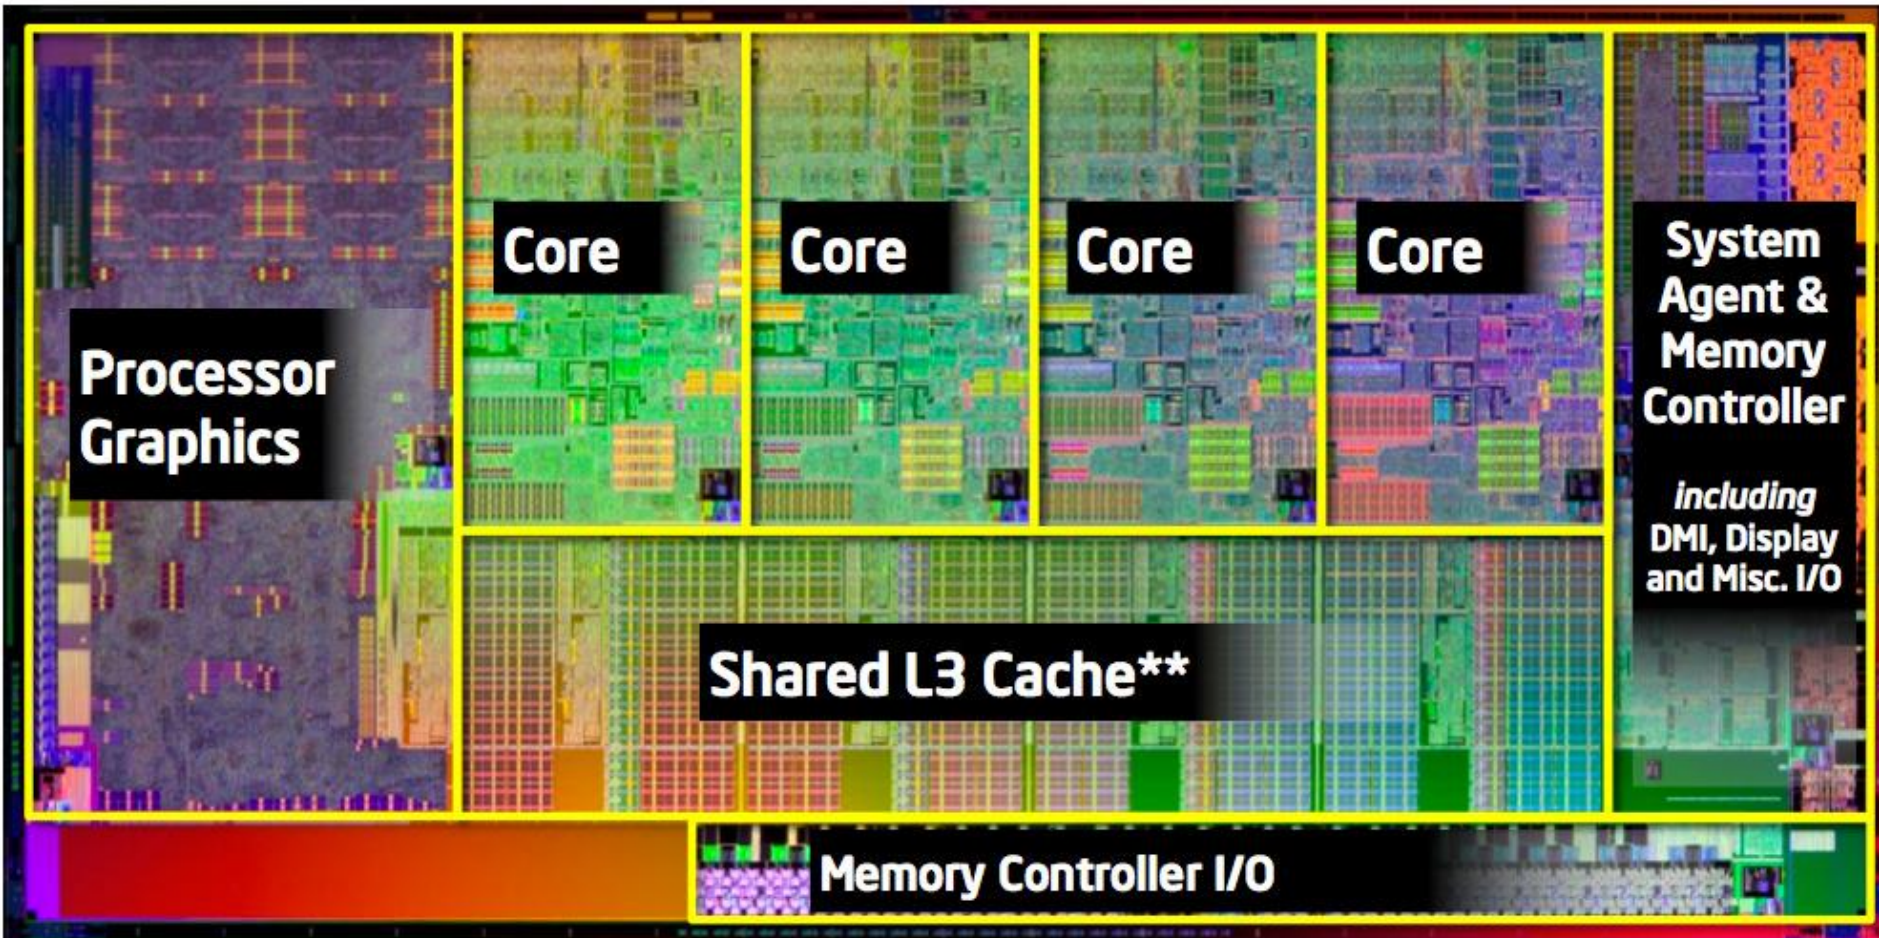
Scalable Ring On-die Interconnect

- **Ring-based** interconnect between Cores, Graphics, Last Level Cache (LLC) and System Agent domain
- Composed of **4 rings**
 - 32 Byte *Data* ring, *Request* ring, *Acknowledge* ring and *Snoop* ring
 - Fully pipelined at **core frequency/voltage**: bandwidth, latency and power scale with cores
- Massive ring **wire routing** runs over the LLC with no area impact
- Access on ring always picks the **shortest path** – minimize latency
- **Distributed arbitration**, sophisticated ring protocol to handle coherency, ordering, and core interface
- **Scalable to servers** with large number of processors



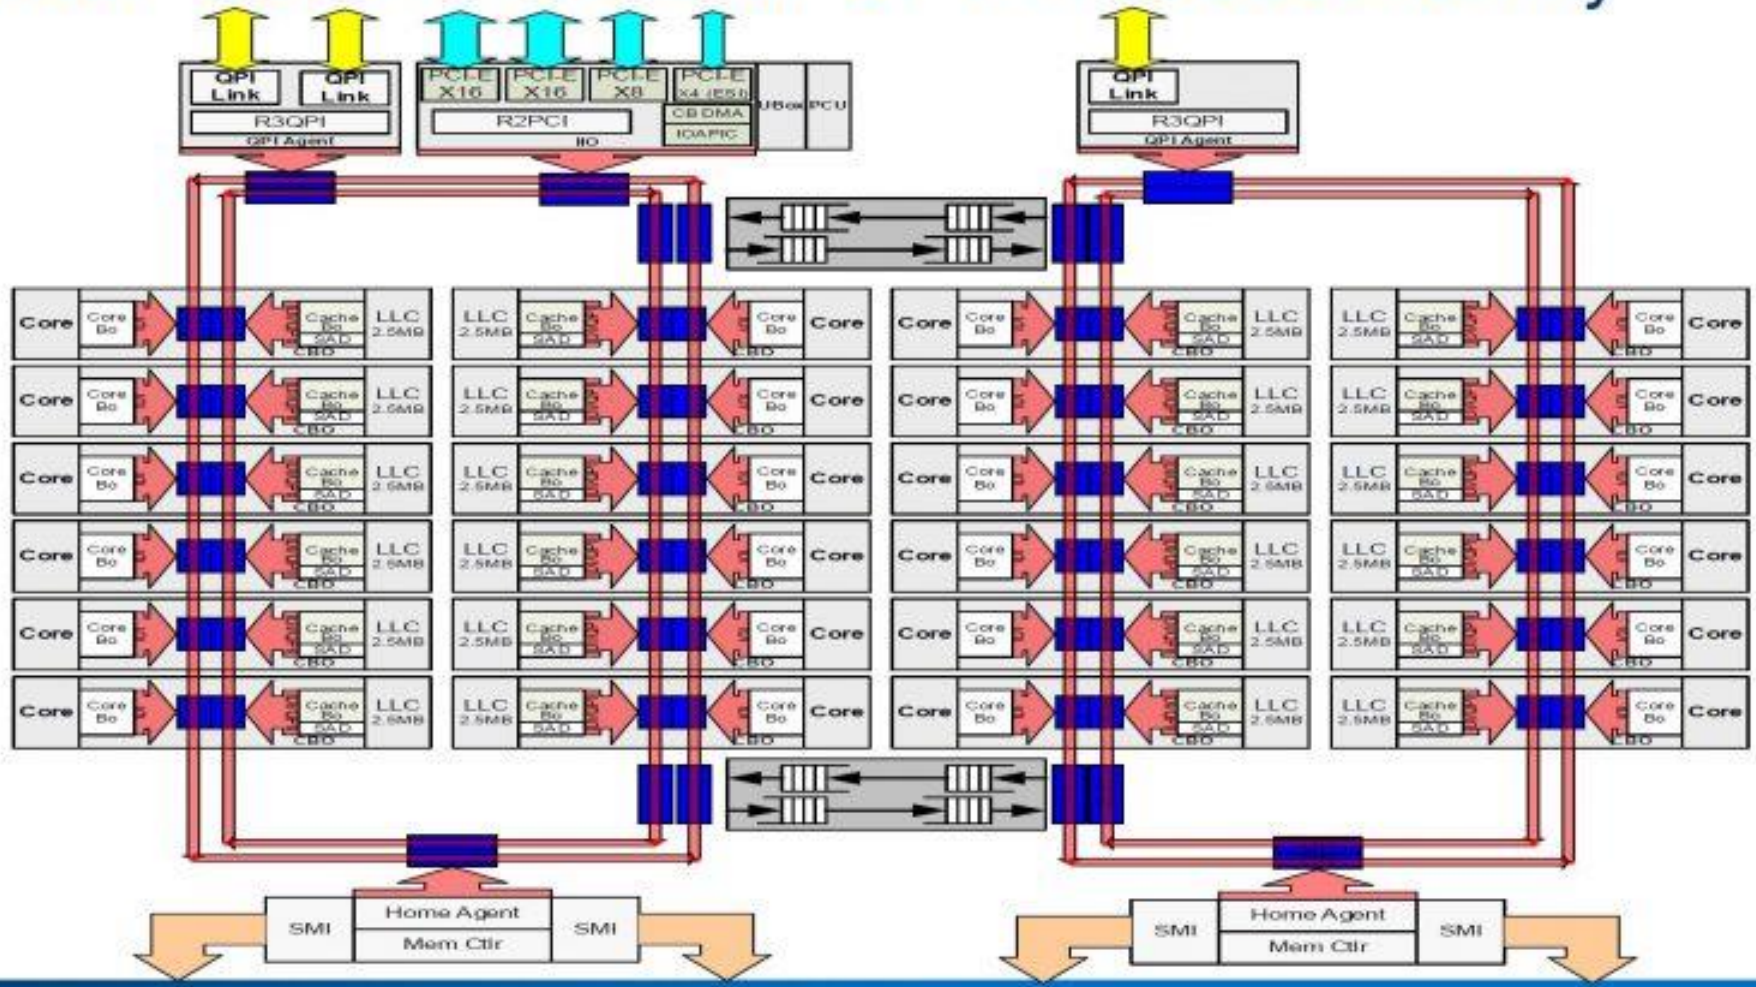
High Bandwidth, Low Latency, Modular

Intel Sandy Bridge

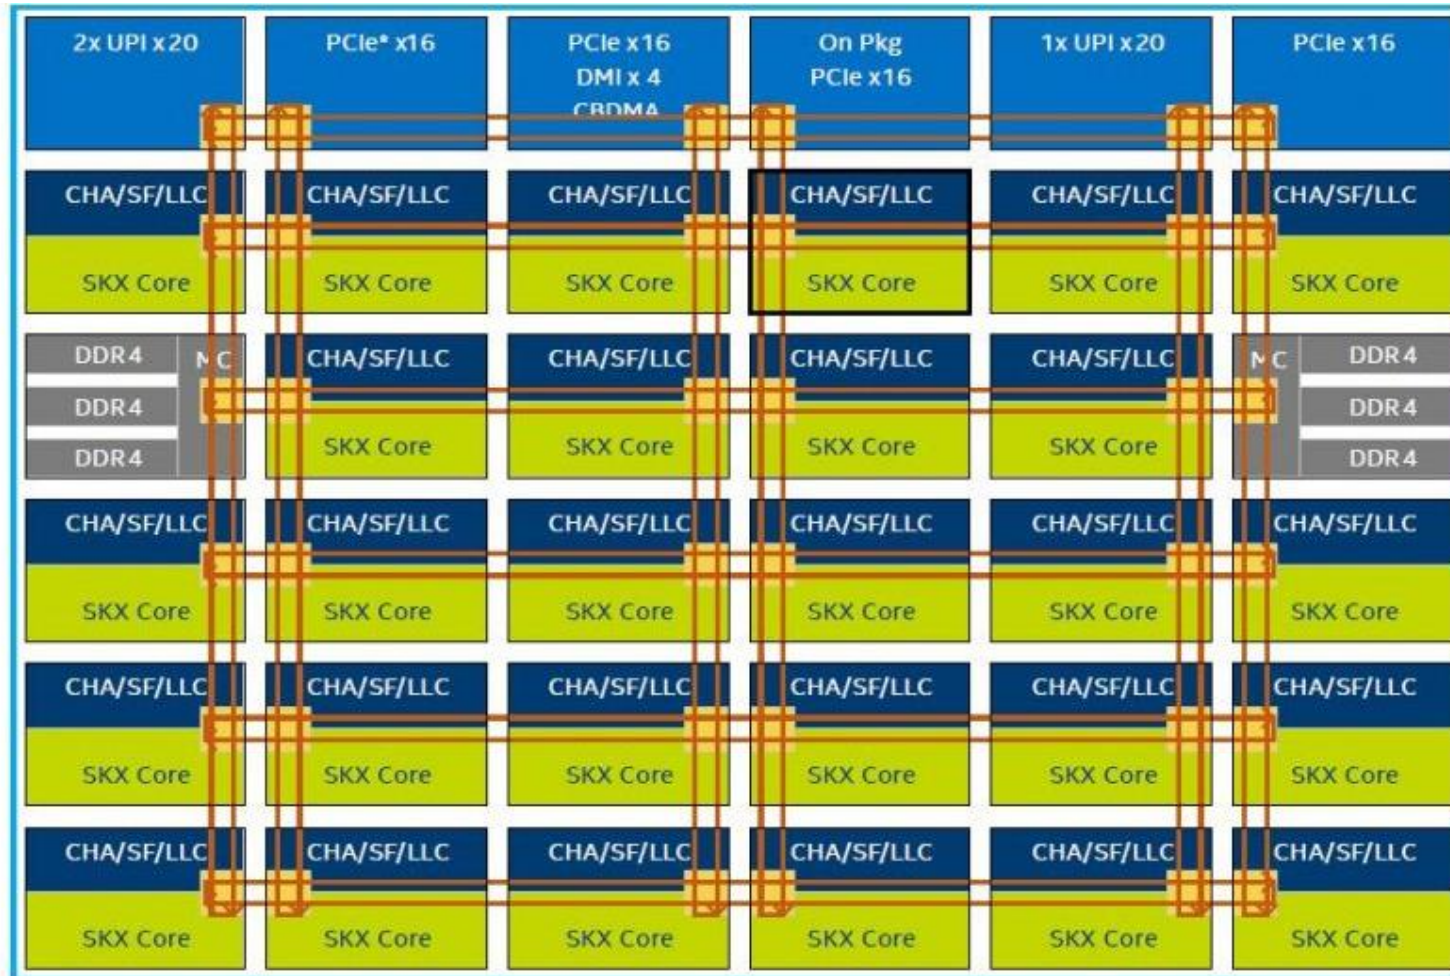


Intel Broadwell

Intel® Xeon® Processor E7 v4 Product Family




Intel Skylake



CHA – Caching and Home Agent ; SF – Snoop Filter; LLC – Last Level Cache;
 SKX Core – Skylake Server Core; UPI – Intel® UltraPath Interconnect

- 256b/cycle per channel (half a cache line)
- Dimension-ordered routing (First X, then Y)


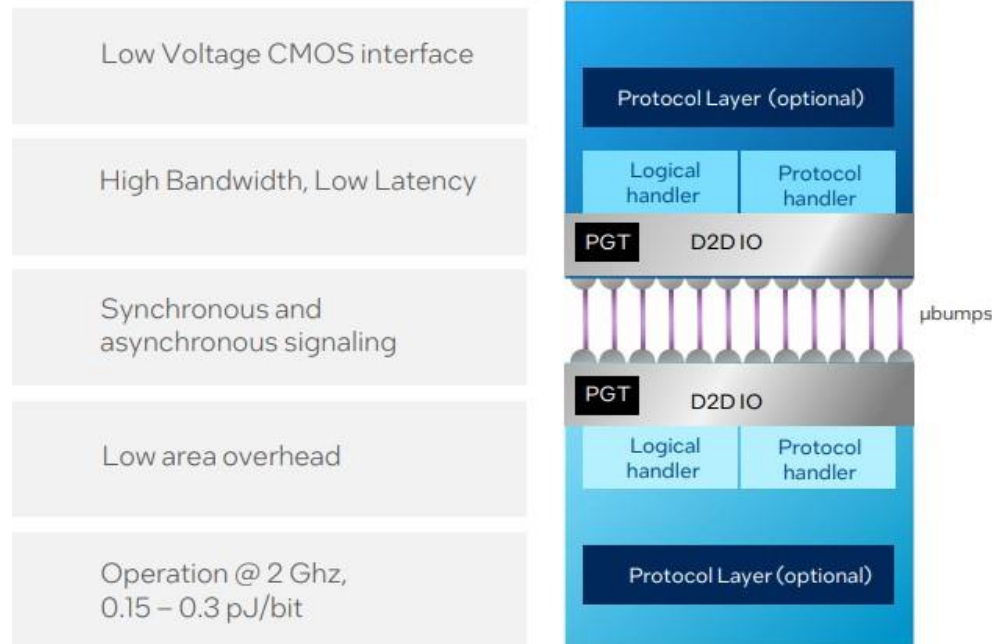
Intel Skylake



Intel Meteor Lake

- 3D integration
- Separate SoC/GPU/CPU/IO Tiles
- “Base Tile” die implements routing

FDI - Foveros Die Interconnect




Cerebras WSE

- 2D mesh across an entire wafer

Largest Chip Ever Built


- 46,225 mm² silicon
- 1.2 trillion transistors
- 400,000 AI optimized cores
- 18 Gigabytes of On-chip Memory
- 9 PByte/s memory bandwidth
- 100 Pbit/s fabric bandwidth
- TSMC 16nm process



Cerebras WSE

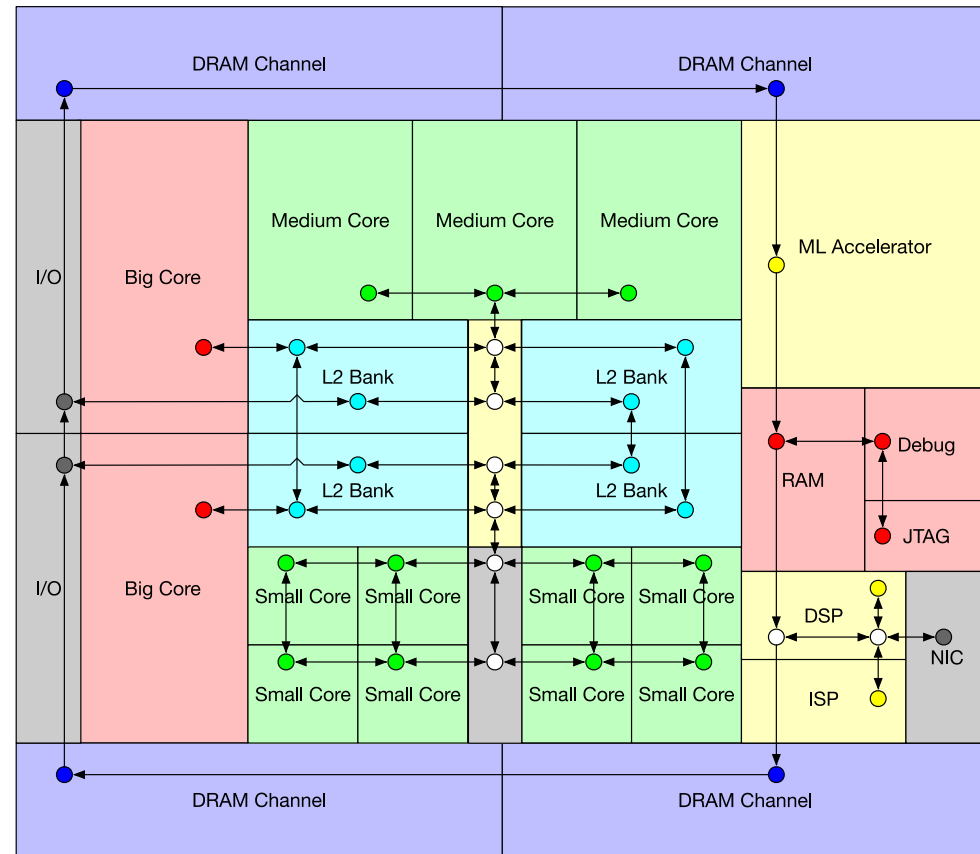
Redundancy is Your Friend

- Uniform small core architecture enables redundancy to address yield at very low cost
- Design includes redundant cores and redundant fabric links
- Redundant cores replace defective cores
- Extra links reconnect fabric to restore logical 2D mesh



Generating Network-on-Chips

- Constellation – a network-on-chip generator
- Generate realistic interconnects for modern SoCs
- Configurable routing/topology/micro-architecture
- If interested, email me – jzh@berkeley.edu



Acknowledgements

- Most diagrams are from the Dally and Towles textbook: *Principles and Practices of Interconnection Networks*
- Die shots, diagrams, slides are from HotChips presentations and AnandTech